

# Bachelor Thesis

Nsak as Framework for Scenario Based Network Security

|                 |                                                    |
|-----------------|----------------------------------------------------|
| Course of study | Bachelor of Science in Computer Science            |
| Author          | Frank Gauss (gausf1) and Lukas von Allmen (vonal3) |
| Advisor         | Wenger Hansjürg                                    |

Version 1.0 of May 24, 2026

# Abstract

As adversaries grow in number, security teams increasingly rely on **red and blue team exercises** to identify vulnerabilities before they can be exploited. Manually engineered scenarios offer flexibility and reproducibility, but require a high level of cybersecurity expertise and lack adaptability to changing conditions. **Large Language Models**, by contrast, can reason over tool output and dynamically adjust their behavior.

This thesis investigates whether **agentic AI** can effectively automate and enhance network security scenarios within the **NSAK** framework — a modular Python-based cybersecurity orchestration platform developed in a predecessor project.

To address this question, the framework is extended with an agentic drill integration built on LangChain. The agent is equipped with tools for CLI interaction, host configuration, and autonomous **purple team reconnaissance**. A deterministic reference scenario — chaining the drills **host-discovery**, **port-scanning**, and **service-enumeration** — serves as the reproducible baseline.

Three models are benchmarked across 10 runs in three environments (*contain-erlab*, *physical lab*, and *BFH Cyber Lab*): a frontier model (Claude Opus 4.7), an institutionally hosted open-weight model (gpt-oss:120b), and a locally-run open-source model (qwen3:30b via Ollama).

results  
/ con-  
clusion  
sentence

# Contents

|                                        |    |
|----------------------------------------|----|
| Abstract                               | ii |
| 1. Introduction                        | 1  |
| 1.1. Motivation                        | 2  |
| 1.2. Predecessor Project: Project 2    | 2  |
| 1.3. Project Goals                     | 2  |
| 1.3.1. Research Goals                  | 3  |
| 1.3.2. Framework                       | 4  |
| 1.3.3. Scenarios                       | 6  |
| 1.3.4. Evaluation Goals                | 8  |
| 1.3.5. Out of Scope / Optional Feature | 9  |
| 2. Related Research                    | 10 |
| 2.1. Framework Comparison              | 10 |
| 2.1.1. Atomic Red Team                 | 11 |
| 2.1.2. Caldera                         | 12 |
| 2.1.3. Metasploit                      | 12 |
| 2.1.4. NSAK                            | 12 |
| 2.2. AI-Research                       | 13 |
| 2.2.1. AI and Language Models          | 13 |
| 2.2.2. Model Context Protocol (MCP)    | 14 |
| 2.2.3. Agentic-AI                      | 14 |
| 3. Concepts                            | 17 |
| 3.1. Framework Resources               | 17 |
| 3.1.1. Device                          | 17 |
| 3.1.2. Environment                     | 17 |
| 3.1.3. Drill                           | 18 |
| 3.1.4. Scenario                        | 18 |
| 3.2. Configuration Management          | 18 |
| 3.2.1. Static Configuration            | 18 |
| 3.2.2. Runtime Configuration           | 19 |
| 3.2.3. Resource Configuration          | 19 |
| 4. Methodology                         | 20 |

|                                                  |    |
|--------------------------------------------------|----|
| 5. Implementation                                | 21 |
| 5.1. Configuration Management Implementation     | 22 |
| 5.1.1. NSAK Configuration Management             | 22 |
| 5.1.2. Device Configuration Management           | 27 |
| 5.1.3. Scenario & Drill Configuration Management | 35 |
| 5.2. Reproducible Test Environment               | 38 |
| 5.2.1. Network Topology                          | 39 |
| 5.2.2. DMZ-Server                                | 40 |
| 5.2.3. Domain Control and File Server            | 41 |
| 5.2.4. Workstations                              | 42 |
| 5.2.5. Printer                                   | 42 |
| 5.2.6. Introduced Vulnerabilities                | 42 |
| 5.3. AI –Scenarios                               | 43 |
| 5.3.1. Artificial Intelligence (AI) Agent        | 43 |
| 5.3.2. LangChain Concepts                        | 45 |
| 5.3.3. Human Interaction Hook                    | 55 |
| 5.3.4. Logging                                   | 55 |
| 5.3.5. E-Mail Alerting                           | 56 |
| 5.3.6. CLI Kill Switch                           | 58 |
| 5.3.7. Scenario: AI Agent                        | 59 |
| 5.3.8. Scenario: AI Reconnaissance               | 61 |
| 5.3.9. Scenario: AI Intrusion Detection          | 64 |
| 5.3.10. LangChain Agent Chat UI                  | 67 |
| 5.4. Red Team: Reconnaissance Scenario           | 68 |
| 5.4.1. Goal                                      | 68 |
| 5.4.2. Discover Hosts Drill                      | 68 |
| 5.4.3. Enumeration                               | 69 |
| 5.4.4. Output Artifacts                          | 69 |
| 5.5. Benchmark Suite                             | 70 |
| 6. Evaluation                                    | 71 |
| 6.1. Evaluation Setup                            | 71 |
| 6.1.1. Criteria                                  | 71 |
| 6.1.2. Environments                              | 72 |
| 6.1.3. Models                                    | 72 |
| 6.2. Todos                                       | 72 |
| 6.3. Virtual Environment (Containerlab)          | 73 |
| 6.4. Physical Lab                                | 75 |
| 6.5. Cyberlab bfh                                | 78 |
| 7. Discussion                                    | 82 |
| 8. Conclusion                                    | 83 |



|                                       |     |
|---------------------------------------|-----|
| Bibliography                          | 84  |
| List of Figures                       | 86  |
| List of Tables                        | 88  |
| List of Listings                      | 89  |
| Glossary                              | 91  |
| A. Project Management                 | 92  |
| A.1. Stakeholders                     | 92  |
| A.2. Agile Management Process         | 92  |
| A.2.1. Issue Lifecycle:               | 93  |
| A.2.2. Definition of Ready (DoR)      | 93  |
| A.2.3. Definition of Done (DoD)       | 94  |
| A.3. Project Structure                | 94  |
| A.3.1. Milestones                     | 94  |
| A.3.2. Epics                          | 98  |
| A.3.3. Issues                         | 98  |
| A.4. Sprints / Iterations             | 98  |
| A.4.1. Planning                       | 98  |
| A.5. Planning and Estimation          | 99  |
| A.6. Project Roadmap                  | 99  |
| A.7. Sprint Ceremonies                | 99  |
| A.7.1. Sprint 1: 19.02 –10.03.2026    | 100 |
| A.7.2. Sprint 2: 05.03 –18.03.2026    | 103 |
| A.7.3. Sprint 3: 26.03 –09.04.2026    | 109 |
| A.7.4. Sprint 4: 10.04 –02.05.2026    | 113 |
| A.7.5. Sprint 5: 03.05 –15.05.2026    | 116 |
| A.7.6. Sprint 6: 15.05 –28.05.2026    | 118 |
| A.7.7. Sprint 7: 29.05 –11.06.2026    | 119 |
| A.7.8. Risk Assessment                | 121 |
| A.8. Meeting Notes                    | 126 |
| A.8.1. Checkpoint Meeting 1: 26.02.26 | 126 |
| A.8.2. Checkpoint Meeting 2: 05.03.26 | 126 |
| A.8.3. Checkpoint Meeting 3: 19.03.26 | 128 |
| A.8.4. Checkpoint Meeting 4: 02.04.26 | 129 |
| A.8.5. Checkpoint Meeting 5: 16.04.26 | 130 |
| A.8.6. Checkpoint Meeting 6: 30.04.26 | 130 |
| A.8.7. Checkpoint Meeting 7: 14.05.26 | 131 |
| A.8.8. Checkpoint Meeting 8: 28.05.26 | 131 |
| A.8.9. Checkpoint Meeting 9: 11.06.26 | 131 |

|                                                                               |     |
|-------------------------------------------------------------------------------|-----|
| A.9. Variant Decision: Thesis Goals . . . . .                                 | 132 |
| A.9.1. Variant 1: AI vs Manually Built Scenarios . . . . .                    | 132 |
| A.9.2. Variant 2: Multiple Blue and Red Team Scenarios (Status Quo) . . . . . | 133 |
| A.9.3. Decision . . . . .                                                     | 133 |
| A.10. Variant Decision: AI Integration . . . . .                              | 134 |
| A.10.1. Variant 1: AI-agent in NSAK . . . . .                                 | 134 |
| A.10.2. Variant 2: NSAK as MCP Server . . . . .                               | 135 |
| A.10.3. Decision . . . . .                                                    | 136 |
| A.11. Workshops . . . . .                                                     | 137 |
| A.11.1. Brainstorm Scenario Selection . . . . .                               | 137 |
| A.11.2. NSAK Scenario: Network Discovery - Red Team . . . . .                 | 139 |
| A.11.3. NSAK Scenario: TLS Downgrade Poodle - Red Team . . . . .              | 140 |
| A.11.4. NSAK Scenario: Honeypot - Blue Team . . . . .                         | 143 |
| A.11.5. NSAK Scenario: Traffic Capture - Red Team . . . . .                   | 143 |
| A.11.6. NSAK Scenario: Intrusion Detection - Blue Team . . . . .              | 144 |
| A.11.7. NSAK Scenario: AI - Purple Team . . . . .                             | 145 |
| A.12. Self-Hosted LLM . . . . .                                               | 147 |
| A.12.1. Hardware Setup . . . . .                                              | 147 |
| A.12.2. System Setup . . . . .                                                | 147 |
| A.12.3. Evaluation . . . . .                                                  | 150 |
| A.12.4. Self-Hosted Base Setup . . . . .                                      | 152 |
| A.12.5. Nginx & Let's Encrypt - Reverse Proxy with SSL Termination . . . . .  | 154 |
| A.12.6. Ollama - LLM Server . . . . .                                         | 155 |
| A.12.7. Open WebUI - Web Chat for LLM Servers . . . . .                       | 158 |
| A.12.8. Drawio MCP - Diagram Tool . . . . .                                   | 161 |
| A.12.9. Netbox - Network Inventory System . . . . .                           | 164 |
| A.12.10. Grafana / Loki - Logging System . . . . .                            | 167 |
| A.13. Writing Quality . . . . .                                               | 172 |
| A.14. Test-Environment-Code . . . . .                                         | 173 |
| A.14.1. container-lab . . . . .                                               | 173 |
| A.14.2. Flaws and Vulnerabilities . . . . .                                   | 179 |
| A.15. Reconnaissance Scenario: Source Listings . . . . .                      | 181 |

# 1. Introduction

According to the World Economic Forum's Global Risks Report 2026, cyber insecurity is listed among the top ten global risks and is projected to rise significantly from position 30 to position 5 in the long term, highlighting the growing adverse threat of AI technologies [1].

Complementing this global perspective, the CrowdStrike Global Threat Report highlights a significant rise in AI-enabled adversaries [2].

The absence of, or understaffing of, security experts, paired with the emerging complexity of systems and business models cite Figueredo2024RCVaR, highlights the increasing need for automated tools and frameworks capable of identifying and mitigating emerging cyber threats.

To address this, the *Network Swiss Army Knife (NSAK)*, a modular, open-source framework, is designed to orchestrate security scenarios consisting of reusable attack drills. The framework aims to support automated security testing and adversary emulation in controlled network environments.

The development of the NSAK framework builds upon previous project work, in which the core architecture, simplified configuration management, and a basic scenario execution model were developed. [3].

This thesis further extends the framework by enhancing configuration management for operating the NSAK device, establishing a reproducible test environment, and investigating how agentic AI can support automated red-team and blue-team scenarios.

The thesis is structured as follows: **Chapter 2** reviews related work and analyzes how selected security frameworks integrate AI into their products. **Chapter 3** introduces the fundamental concepts and outlines the key design decisions that form the basis of this work. **Chapter 4** describes the research methodology, including the approach used for evaluation and data collection. **Chapter 5** presents the implementation of the proposed solution, followed by **Chapter 6**, which evaluates its effectiveness in a virtual and physical test environments. Finally, **Chapter 7** discusses the results and their implications, while **Chapter 8** concludes the thesis and highlights potential directions for future work.

Match the structure overview to the actual structure

sharpen research question and contribution

## 1.1. Motivation

**Hypothesis:** AI can enhance red and blue team simulations in a network environment by generating and executing scenarios within the Network Swiss Army Knife (NSAK) framework and assisting in the evaluation of detection results. AI can enhance red and blue team simulations in a network environment by generating and executing scenarios within the NSAK framework and assisting in the evaluation of detection results.

## 1.2. Predecessor Project: Project 2

This section summarizes the predecessor project that led to the bachelor's thesis. The NSAK is a cybersecurity framework that operates with the following resources and concepts: *Devices*, which are physical or virtual machines capable of running the NSAK. The security operation is executed in a network environment and can involve one or more attacks or defense scenarios. The smallest unit in the framework is a drill, which can be orchestrated in a scenario and chained with other drills.

While the modularity and proof of concept were confirmed, several aspects like enhanced configuration management, Graphical User Interface (GUI) or a Representational State Transfer Application Programming Interface (REST API) are mentioned for future work.

The goals of this thesis directly address these open points and are defined in [Section 1.3](#).

Proof  
Reading

## 1.3. Project Goals

For a goal to be considered **S.M.A.R.T.**, it must fulfill the following criteria:

- ▶ **Specific:** What do we want to achieve?
- ▶ **Measurable:** What are the success criteria?
- ▶ **Achievable:** How do we plan to achieve this goal?
- ▶ **Relevant:** Why is it important to achieve this goal?
- ▶ **Time-bound:** Goals categorized as *must* have to be achieved within the thesis, while *should* or *could* goals are considered optional or nice to have.

### 1.3.1. Research Goals

#### Compare Frameworks (must)

To identify the USP (unique selling point) of the NSAK framework, we will analyze and compare existing security emulation frameworks. First, relevant frameworks such as MITRE Caldera and Atomic Red Team will be identified and reviewed. Based on this analysis, the identified concepts will be compared with the architecture of the NSAK framework. The comparison will focus on how scenarios and drills are defined, how configuration and parameter handling are structured, how reusable components are organized, how automation is achieved, and which AI capabilities they provide.

The analysis and comparison will be based on the following properties:

- ▶ Concepts
- ▶ Modularity
- ▶ Configurability
- ▶ Automation
- ▶ AI Capabilities

#### Success criteria:

- ▶ At least two security emulation frameworks are identified and reviewed
- ▶ A matrix comparing the specified properties of the selected frameworks is created
- ▶ The USP of the NSAK framework is identified

#### AI Integration Research (must)

The goal of this research is to investigate how AI (artificial intelligence) can support automated security testing and adversary emulation. The research evaluates potential integration points where AI could enhance the scenario execution, automation, decision-making, and analysis capabilities of the NSAK framework.

#### Success criteria:

- ▶ How is AI currently used in cybersecurity automation?
- ▶ Which tasks in adversary emulation can be supported or automated by AI?

- ▶ How could AI interact with the NSAK scenarios and drills (e.g., Agents, MCP)?
- ▶ What architectural changes would be required to integrate AI into NSAK?

### 1.3.2. Framework

#### Configuration Management (must)

The goal is to extend the NSAK framework, so that scenarios and drills can expose configurable parameters to increase the flexibility of scenarios and reusability of drills. Parameters can be provided via command-line arguments or YAML Ain't Markup Language (YAML) files and are resolved by the framework into a unified run configuration used for scenario execution.

#### Success criteria:

- ▶ Scenarios and drills can expose configurable parameters
- ▶ The available parameters can be listed by the Command Line Interface (CLI) with the `--help` option
- ▶ Parameters can be provided via CLI arguments or YAML files
- ▶ The framework resolves the inputs into a run configuration (internal data structure).

#### REST API (could)

Adding a REST Application Programming Interface (API) to the NSAK framework was already proposed in the predecessor project, as it greatly enhances interoperability with other tools. The priority of this goal depends heavily on the outcome of the goal 1.3.1 and on whether we want to fulfill the goal 1.3.2, which is prioritized as “could”. If the evaluated AI integration requires a REST API (e.g., MCP), this goal needs to be prioritized as “must”; it will be prioritized as “could”.

#### Success criteria:

- ▶ A REST API specification with feature parity to the CLI is created.
- ▶ The REST API is implemented according to the specification.
- ▶ The API is documented (OpenAPI) and usable by external clients.

- ▶ The REST API remains optional and does not replace the CLI-based interaction.
- ▶ Authentication and Authorization are explicitly out of scope.

### MCP Server (could)

The MCP Protocol allows the AI to interact with applications in a standardized way - and could therefore be necessary for implementing AI-based scenarios. If the outcome of the research task [1.3.1](#) highlights the necessity to specify/expose an MCP-Server to interact seamlessly with the NSAK framework, the implementation is a (must), otherwise a (could).

- ▶ A MCP-Server is specified and implemented
- ▶ Documentation how to connect to the MCP is provided.
- ▶ A form of authentication is provided.

### Web GUI (could)

The goal is to design and implement a web-based graphical user interface that allows users to interact with the NSAK framework in a user-friendly way. The Web GUI will enable management of scenarios and drills, triggering scenario execution, and visualization of execution results and system status.

### Success criteria:

- ▶ A simple Web GUI that interacts with the REST API has been implemented.
- ▶ The Web GUI displays a list of available scenarios and their states.
- ▶ The scenario's results can be viewed on the Web GUI
- ▶ Scenarios can be configured and executed via the Web GUI
- ▶ Authentication and authorization are explicitly out of scope.

### 1.3.3. Scenarios

#### AI Scenarios (must)

The goal is to implement artificial intelligence according to the result of 1.3.1 and to explore the use of AI to support red, blue, and purple team exercises in the NSAK framework. The AI will be evaluated for its potential to assist in executing attacks in a scenario and analyzing detection results. The findings will help assess how AI could enhance and coordinate red and blue team activities across a network via the NSAK device.

#### Success criteria:

- ▶ The AI is integrated into NSAK as a scenario or as part of the framework, depending on the result of 1.3.1.
- ▶ The AI can execute CLI or Python tools to directly or indirectly interact with the attached network.
- ▶ An operator can interact with or stop the running AI.
- ▶ The result artifacts are stored and can be presented or reused for further scenarios.

#### Reproducible Test Environment (must)

The goal is to design and implement a reproducible and safe test environment for the NSAK framework that enables consistent execution and evaluation of scenarios and drills. The environment should provide a stable basis for testing and comparison of manual and AI-assisted scenarios.

#### Success criteria:

- ▶ The test environment can be set up reproducibly using a defined configuration or setup procedure with Infrastructure as code (e.g., Ansible)
- ▶ NSAK scenarios and drills can be executed consistently within the environment.
- ▶ The environment supports both manual and AI-assisted scenarios.
- ▶ Scenario results can be collected and analyzed for evaluation purposes.



### Network Discovery – Red Team (must)

The goal is to implement a reference red team scenario for network discovery within the NSAK framework. The scenario should simulate reconnaissance activities to identify hosts, services, and network characteristics, and the results should be comparable to the AI-Agent approach.

#### Success criteria:

- ▶ A red team network discovery scenario is implemented in the NSAK framework.
- ▶ Discovered network information is collected as artifacts.
- ▶ The scenario can be executed reproducibly within the test environment.
- ▶ The results can be compared with those generated by the AI-assisted approach.

### Intrusion Detection – Blue Team (should)

The goal is to define and implement a blue team reference scenario for intrusion detection within the NSAK framework. The scenario should allow observation and evaluation of suspicious network activity to assess detection capabilities in the test environment and compare the results with those of the AI approach.

#### Success criteria:

- ▶ A blue-team intrusion-detection scenario is implemented in the NSAK framework.
- ▶ Detection events or alerts can be observed and documented during scenario execution.
- ▶ Detection results can be compared with results from AI-assisted scenario execution.

### 1.3.4. Evaluation Goals

#### AI Scenario Evaluation (must)

The objective is to determine the effectiveness of the AI-based scenario introduced by the goal 1.3.3 within the network environment defined by the goal 1.3.3. The evaluation will assess whether the AI scenario successfully executed the expected red and blue team activities for the provided prompts and the quality of the results.

#### Success criteria:

- ▶ A methodical approach for the assessment is evaluated and documented.
- ▶ The results of the AI scenario for red and blue team activity are assessed.
- ▶ There is a conclusion on the effectiveness of the AI scenario for red and blue team activities.

#### AI vs Engineered Scenarios (must/should)

The goal of this evaluation is to compare the quality of AI-based red (network reconnaissance) (must) and blue (network intrusion detection) (should) team scenarios with that of manually engineered scenarios.

The comparison focuses on criteria such as execution reliability, reproducibility, and coverage. The results will help determine the potential benefits and limitations of AI-supported scenarios and provide insights into how AI could enhance future versions of the NSAK framework.

#### Success criteria:

- ▶ A methodical approach for the comparison is evaluated and documented.
- ▶ The results of the AI scenario are compared against the respective results of the red team scenario (must)
- ▶ The results of the AI scenario are compared against the respective results of the blue team scenario (should)
- ▶ There is a conclusion on the effectiveness of the AI scenario compared to manually engineered scenarios.

### 1.3.5. Out of Scope / Optional Feature

At the project management meeting on 05.03.26 (see Table [A.17](#) in Appendix A.7), the initial project scope was discussed. During planning, the team proposed using AI-based scenarios and comparing them with the implemented versions, which required reassessing certain milestones.

Running AI-based scenarios does not require a REST API or GUI. A reproducible test setup and further research are needed to properly evaluate this approach.

## 2. Related Research

### 2.1. Framework Comparison

Previous research, such as ‘SOK: A Survey of Open-Source Threat Emulators’, compares multiple types of threat emulation frameworks using both quantitative and qualitative metrics [4]. Among the evaluated frameworks, Atomic Red Team, Mitre Caldera, and Metasploit received high scores in categories as environments and operator, environment configuration and scenario execution.

In the following section, the selected frameworks are briefly introduced, and the unique selling point of the NSAK is highlighted. As a point of reference, the cybersecurity community increasingly relies on the concepts of MITRE ATT&CK, which has established itself as a standard framework in the field since its introduction in 2014 [5, 6].

The MITRE ATT&CK is a globally accessible knowledge base of adversary tactics and techniques based on real-world observations [5]. The knowledge base is structured into TTP (Tactics, Techniques, and Procedures), which are organized into 14 attack behavior groups. The TTP is getting continuously updated and extended, which helps to defend against the capabilities of threat actors.

The evaluation of the proposed frameworks is based on information obtained from official documentation [7–9] and the findings of Zilberman et al. [4].

| Characteristic           | Metasploit Framework                                             | Atomic Red Team                                            | MITRE Caldera                                 |
|--------------------------|------------------------------------------------------------------|------------------------------------------------------------|-----------------------------------------------|
| Primary Purpose          | Exploitation and penetration testing framework                   | Security testing through small atomic ATT&CK tests         | Automated adversary emulation platform        |
| Relation to MITRE ATT&CK | Partial mapping via modules                                      | Direct mapping to ATT&CK techniques                        | Built around ATT&CK adversary emulation       |
| Automation Level         | Medium (scripts and modules)                                     | Low (manual execution of tests)                            | High (automated attack chains)                |
| Main Functionality       | Exploit development, payload delivery, post-exploitation         | Execute small ATT&CK technique tests                       | Simulate adversary campaigns across hosts     |
| Architecture             | Modular framework with exploits, payloads, and auxiliary modules | Collection of scripts/tests organized by ATT&CK techniques | Client-server platform                        |
| Typical Use Case         | Penetration testing and vulnerability validation                 | Detection validation and security control testing          | Red team automation and purple team exercises |
| Difficulty Level         | Medium-High                                                      | Low-Medium                                                 | Medium-High                                   |

Table 2.1.: Comparison of Metasploit, Atomic Red Team, and MITRE Caldera ([4, 7–9])

As shown in Table 2.1, the key characteristics of the frameworks differ greatly. **Atomic Red Team** provides libraries of attack scripts that implement individual Tactics, Techniques, and Procedures of MITRE ATT&CK. Each attack can be executed independently or chained together into a larger attack scenario. **Mitre Caldera** is a complex, agent-based adversary simulation tool, whereas **Metasploit** is a framework that provides exploits, payloads and auxiliary modules. Further details of the frameworks are presented in the following sections. 2.1.12.1.22.1.3

Table formatting

### 2.1.1. Atomic Red Team

The Atomic Red Team repository provides documentation and configuration files for each technique. For every TTP, a YAML configuration file defines the execution parameters, while an enhanced Markdown file contains a human-readable description of the attack. This documentation includes a short description of the technique, execution instructions, and additional contextual information [8].

The framework can be executed directly from the command line and provides

features such as logging, detailed error messages, and automatic prerequisite checks. If a technique cannot be executed, for example due to operating system constraints, the framework reports the issue and can optionally attempt to install the required prerequisites [8].

*Zilberman et al.* recommends deeper cybersecurity knowledge to use Atomic Red Team framework effectively [4]. The Open Source mentality, expandability and easily usable scripts can be a great asset for an independent framework such as the NSAK.

### 2.1.2. Caldera

Mitre Caldera is also built on the MITRE ATT&CK knowledge base and functions as a command-and-control center. A central server controls agents running on the target system, executes attacks called abilities, and sends the operation results back to the server. Every Ability is a single attack step built on the Tactics, Techniques, and Procedures. The framework ships with a graphical user interface and is primarily used for red team attack simulation, purple teaming, and detection engineering. The additional features, like a training plugin for educational purposes, include facts that can be configured as dynamic safe values, which can be stored and parsed in later steps [7].

With its provided interface and functionality, Caldera is most suitable as a threat emulator, offering a wide range of built-in features, including lateral movement and Command and Control (CNC) communication [4].

### 2.1.3. Metasploit

Metasploit is a widely used penetration testing framework for testing or attacking systems and exploiting known vulnerabilities. Metasploit follows a sequence of steps to identify a vulnerability in a target system, select a vulnerability for the attack, select an exploit that uses this vulnerability, and deliver a payload that runs on the target system. The configuration via the CLI requires knowledge of the target network and system and requires partial configuration [9, 10].

According to Zilberman et al., an operator requires extensive cybersecurity knowledge to effectively use the framework. To support the operator, the GUI “Armitage”, which is not further evaluated in this work, can be used [4].

### 2.1.4. NSAK

The NSAK framework, as an independent universal Network Swiss Army Knife, is not confined and can leverage, if needed, different sets of attack frameworks.

An approach could be to use an AI-Enhanced Network Discovery Scenario, which provides the necessary information to configure a Metasploit attack that leads to an attack operation in Caldera and an atomic red team detection script. Hardware-based isolation and network position provide flexibility and should help reduce the operational costs of red and blue teams by reducing simple configuration overhead [3].

## 2.2. AI-Research

### 2.2.1. AI and Language Models

Language Model (LM) describes the technique for advancing language intelligence in machines and its development history can be divided into several stages. The first stage, Statistical Language Model (SLM) emerged around the 1990s, and the basic idea is to assume, that such a model can predict the next word. But the fact that an exponential number of transition probabilities are needed, caused a lack of accuracy and led limited success.

To overcome these limitations, Neural Language Model (NLM) was introduced to learn a distributed representation of words and capture more complex patterns, leading to a major breakthrough. Building on this, the advancement of pre-trained language models Pre-trained Language model (PLM), which are first trained on large datasets and then fine-tuned for specific tasks, represented another significant development.

Building on this development, Large Language Model (LLM)s further increases the scale of PLMs by expanding parameter count and computational power [11].

In large language models, parameters are the trainable numerical values that define how the model processes input and generates output. During training, these parameters are continually adjusted to enable the model to learn patterns and language context [12]. The models most of us use today, such as ChatGPT, Claude and Gemini are pre-trained LLMs used for language inference.

These advances led to MCP, a protocol that allows tools and services to expose an interface specifically for AI. Further, the introduction of AI-agents is enabling a new level of autonomy for achieving concrete goals with LLMs. The combination of these technologies is now resulting in systems, where LLMs empower autonomous agents capable of decision-making to execute complex tasks.

### 2.2.2. MCP

Anthropic, a US software company focused on AI research, introduced the protocol in 2024. The MCP was introduced as an open source standard that enables AI assistants to connect to external systems, including databases, business tools, and development environments [13]. Dynamic tool discovery and schema negotiation is helping clients and agents to access these functions in a standardized manner and allows autonomous interaction. Additionally, MCP allows information to flow in both directions, so the model can send requests, and tools can send updates or events.

This architectural approach shifts the AI integration from static toward an interoperable ecosystem of AI-enabled services. On the other hand, the flexible design and open source approach of MCP unlocks multiple novel and systematic security risks that are in need of enhanced security and governance [14].

### 2.2.3. Agentic-AI

Agentic Artificial Intelligence (Agentic-AI) can be understood as an advanced machine intelligence that is capable of perceiving, reasoning, and acting with a certain level of autonomy. Traditional systems and agents are more limited to executing predefined commands, whereas Agentic-AI is not restricted in this way. Due to their ability to adapt in real time, they can continuously refine their knowledge, techniques, and data. The ability to access multiple inputs, from natural language processing to sensor data, allows Artificial Intelligence Agent (AI Agent)s to develop a good understanding of their environment [15].

Multi Agent System (MAS) defines a way in which multiple specialist agent coordinate and cooperate with each other to solve a task, that would not be solvable to one agent alone.

Therefore, the agents need a form of interaction that allows them to act reactively, often through a shared mental model, which helps create a common understanding of the complex tasks to be solved. To implement a MAS efficiently, modern design principles such as encapsulation and modularity are important, which are helping to balance resource usage of the system and LLM [16]

Huang et al. provide evidence that Agentic-AI and MAS are capable of solving complex tasks across various sectors, including security and banking. However, these systems also introduce new potential vulnerabilities, both within AI Agent systems themselves and through their interactions with external environments.

The identified vulnerabilities can be categorized into two types: accidental and deliberate. Accidental vulnerabilities include software bugs, logical errors introduced by agents, hardware malfunctions, and poor data quality. Deliberate



attacks, on the other hand, encompass adversarial attacks targeting LLMs, data poisoning, and model theft.

To mitigate these risks, a comprehensive framework consisting of governance mechanisms, proactive monitoring, regulatory compliance, and rigorous testing is required to ensure the secure and reliable operation of agentic systems [17].

### Fazit MCP and Agentic AI

In the cybersecurity podcast Risky Business, it is mentioned that MCP is “dead”. The open-source strategy of Anthropic (2.2.2) has led to the development of multiple MCP implementations that were not designed with a strong focus on security, but rather on rapid adoption of a novel technique.

This provocative claim is based on the observation that, in order to handle the excessive MCP tooling, AI Agents may instead rely on direct access to a root shell, as it provides a more efficient way of interaction [18].

### Evaluated Frameworks and MCP

Caldera uses a MCP plugin as a bridge between the model and the REST API, allowing the agent to plan and issue commands in a manner similar to a human operator.

Caldera introduces three main research areas. Firstly, a planner determines the order and selection of abilities and can dynamically construct new adversaries based on user requirements through prompt-based interaction.

Secondly, an adversary factory model is used to generate adversaries tailored to the operator’s specific use cases.

Thirdly, a forward-deployed agent is enhanced with Retrieval Augmented Generation (RAG) for Cyber Threat Intelligence (CTI), leveraging Structured Threat Information Expression (STIX)-based data to create reproducible and testable adversaries [19].

The MCP-Server, provided by **Metasploit**, enables a LLM to interact with the Metasploit Framework dynamically. The Interface covers module information (a Metasploit module is a known exploit), exploitation workflows, payload generation, session management and handler management to help manage an attack. The documentation includes a security warning that emphasizes the importance of reviewing every prompt, using a test environment, and being aware of post-exploitation commands. **Metasploit** is a powerful exploitation framework, and the simplicity that AI could bring can also harm. [20]

How and why is this relevant for our thesis? how could we implement it? To which findings should we pay attention? What can we ignore for now?

Elaborate on this, because we will actually not use MCP

I think we are to focused on the term MCP now. We should speak of tool-calling, MCP is just one

The **Atomic Red Team** MCP Server provides an interface that allows AI systems to access and use tests from Atomic Red Team. It enables searching, validating, and optionally executing attack techniques (e.g., based on MITRE ATT&CK). [21]

### Outcome Comparison

The framework comparison highlights the difference between the NSAK and the established frameworks. Each compared framework performs a specific task in a complex operation, and the NSAK device can leverage the strengths of multiple aspects of a different set of tools. Building upon these insights, the NSAK is designed to extend and integrate these capabilities into a unified and adaptable system and the change of focus to AI gets further elaborated in this thesis.

What about the other frameworks? What about AI Agents? What can we learn from their approach? How does it affect our strategy?

## 3. Concepts

### 3.1. Framework Resources

The NSAK framework is built around four resource types that were originally defined in “Project 2” 1.2 [3]. Each resource type has a distinct role in the framework and is stored as a YAML file in the resource library.

#### 3.1.1. Device

A **Device** resource describes a physical or virtual machine that is capable of running the NSAK framework. It serves as the execution host for scenarios and drills and is typically deployed at a strategic network position to give it the access required for the operations it performs.

A device exposes its network configuration to the framework by declaring its interfaces and their Internet Protocol addresses in its resource YAML. As an alternative, a special device with the id “Unknown” was added, which auto discovers the hosts network configuration. This allows scenarios and drills to resolve interface names and target addresses without hardcoding network details, making the same scenario portable across different devices and network environments.

#### 3.1.2. Environment

An **Environment** represents a specific network topology, including its infrastructure components, hosts, and services [3]. It describes the target network context in which a scenario is intended to operate. For example, a simple client-server setup over a switch, a Wireless LAN access point with connected clients, or a segmented business network.

Environments are referenced by scenarios to make explicit which network context they are designed for, helping operators select and validate the right configuration before executing an operation.

### 3.1.3. Drill

A **Drill** is a self-contained, reusable unit of execution with a specific, well-defined goal [3]. Its goal can be an active or passive attack step, network discovery, traffic capture, a configuration change, or any other discrete action relevant to a red or blue team activity.

Examples of drills include Address Resolution Protocol spoofing, host discovery via Address Resolution Protocol scan, transparent Transmission Control Protocol proxying, traffic capture with T-Shark, and network configuration tasks such as enabling IP forwarding or Network Address Translation (NAT).

Each drill declares the system and Python packages it requires as well as a typed interface that lists its input arguments and return type. The return value of one drill can be consumed as the input of a subsequent drill within a scenario, enabling data to flow through a chain of execution steps.

### 3.1.4. Scenario

A **Scenario** describes a concrete red or blue team use case and orchestrates a sequence of drills targeted at one or more environments [3]. Where a drill is a single technical reusable and modular building block, a scenario is an implementation of a specific attack or defense operation.

Like drills, scenarios declare a typed interface with input arguments, allowing operators to parameterise a scenario at execution time via the CLI without modifying the scenario definition itself. This separation keeps scenarios reusable across different network configurations while still allowing runtime flexibility.

## 3.2. Configuration Management

The NSAK framework distinguishes three configuration layers, each with a different scope and lifetime. Keeping them separate avoids coupling concerns that change at different rates and makes it easier to reason about where a given piece of configuration comes from.

### 3.2.1. Static Configuration

The static configuration layer involves values which are specific for a deployment and are very unlikely to change after the initial provision of the NSAK device. It mainly establishes the filesystem layout that all other parts of the framework

depend on. For example, where the resource library is located, where runtime data is written, and what the base directory of the installation is.

### 3.2.2. Runtime Configuration

The runtime configuration layer holds mutable state, which might be changed by an operator via the CLI. It is stored as a YAML file, which is currently loaded at the start of every invocation. If no configuration file exists yet, a default is created automatically on first access, so no explicit initialization step is required from the user.

### 3.2.3. Resource Configuration

The resource configuration layer is provided by the resource library rather than the framework core. Each resource declares its own configuration inside a YAML file stored alongside its implementation. The framework reads these files to discover and load resources at runtime.

### Typed Interface Arguments

Drills and scenarios declare a typed interface in their resource configuration YAML file, that lists the arguments with an explicit type annotation and an optional default value. The framework reads these annotations at runtime and automatically generates the correct CLI options for each resource. This means that adding a new argument to a drill or scenario immediately exposes it as a typed CLI option without requiring any changes to the framework or the CLI code.

### Mergeable YAML Structure

All resource YAML files follow a common structure in which resources are nested under a resource-type key (e.g. `drills:`, `scenarios:`). This structure ensures that YAML files from different parts of the library could be merged into a single in-memory data structure without key conflicts, which is a prerequisite for supporting multiple library paths and composite library configurations in the future.

## 4. Methodology

## 5. Implementation

---

summary  
over  
imple-  
mented  
parts

## 5.1. Configuration Management Implementation

The configuration management in NSAK is organised into three distinct layers, each serving a different scope and lifetime.

The **static configuration** layer resolves fixed filesystem paths at import time. It derives all relevant paths from a single base directory, which can be overridden via an environment variable to support different deployment contexts such as Docker containers or automated test environments.

The **runtime configuration** layer holds mutable state that persists across CLI invocations. It is loaded from a YAML file on every command execution and written back whenever state changes, for example, when a different device is loaded. This layer serves as the single source of truth for any information that must be shared between separate CLI calls.

The **resource configuration** layer is defined by the resource library itself. Each resource type –devices, drills, scenarios, and environments– declares its configuration inside a YAML file stored in the library. For devices, this includes network interface assignments, while drills and scenarios additionally declare their typed argument interfaces, which the CLI uses to generate the correct command options at runtime.

### 5.1.1. NSAK Configuration Management

The NSAK configuration is split into two distinct layers: a static layer that resolves fixed paths at import time, and a mutable runtime layer that is persisted to disk and loaded on every invocation.

#### Static Configuration

The static configuration is defined in `src/nsak/core/settings.py` and shown in listing 1. It derives all relevant paths from a single `BASE_PATH`, which is either supplied via the `NSAK_BASE_PATH` environment variable or resolved automatically to the repository root at import time. `RUN_PATH` points to the `run/` directory where runtime data such as the persisted config is stored, and `LIBRARY_PATHS` holds the set of directories that are searched for resource libraries. Overriding `NSAK_BASE_PATH` is the primary mechanism to run NSAK from a different base directory, e.g. inside a Docker container or during automated testing.

This section is repeating what's already stated in the concepts chapter...



```
1 import logging
2 import os
3 from pathlib import Path
4
5 from dotenv import load_dotenv
6
7 load_dotenv(os.getenv("NSAK_ENV_FILE", None))
8
9 logger = logging.getLogger()
10
11
12 def get_base_path() -> Path:
13     """
14     Returns the default base path.
15
16     :return:
17     """
18     base_path = os.getenv("NSAK_BASE_PATH", None)
19
20     if base_path is not None:
21         return Path(base_path).absolute()
22
23     return Path(__file__).resolve().parents[3].absolute()
24
25
26 def get_run_path() -> Path:
27     """
28     Returns the run path.
29
30     :return:
31     """
32     run_path = os.getenv("NSAK_RUN_PATH", None)
33
34     if run_path is not None:
35         return Path(run_path).absolute()
36
37     return (get_base_path() / "run").absolute()
38
39
40 def get_library_paths() -> set[Path]:
41     """
42     Returns a list of library paths.
43
44     :return:
45     """
46     library_paths = set()
47
48     default_path = (BASE_PATH / "lib").absolute()
49     if default_path.exists():
50         library_paths.add(default_path)
51
52     library_path = os.getenv("NSAK_LIBRARY_PATH", None)
53     if library_path is not None:
54         library_path = Path(library_path).absolute()
```

## Runtime Configuration

The runtime configuration is defined as the “Config” dataclass in `src/nsak/core/configuration`, and holds all state that can change during normal operation, currently the debug flag and the active “Device”. It is persisted as a YAML file at `run/config.yaml` and loaded on every CLI invocation.

To avoid loading the configuration on import, the module-level `config` object is wrapped in a `lazy_object_proxy.Proxy`. The proxy defers the actual `Config.load()` call until the object is first accessed. Since the root CLI group imports `config` on every command invocation (listing 3), this ensures the configuration is always initialised before any command logic runs. If `run/config.yaml` does not yet exist, `Config.init()` is called automatically to create a default configuration and write it to disk, eliminating the need for manual setup on first run.

---

```

1  from __future__ import annotations
2
3  import logging
4  import uuid
5  from copy import deepcopy
6  from dataclasses import dataclass, field
7  from datetime import datetime
8  from pathlib import Path
9  from zoneinfo import ZoneInfo
10
11 from nsak.core.configuration.ai_configuration import
12     ↳ AiConfiguration
13 from nsak.core.configuration.container_info import ContainerInfo
14 from nsak.core.configuration.drawio_mcp_configuration import
15     ↳ DrawioMCPConfiguration
16 from nsak.core.configuration.email_configuration import
17     ↳ EmailConfiguration
18 from nsak.core.configuration.loki_configuration import
19     ↳ LokiConfiguration
20 from nsak.core.device import Device
21 from nsak.core.settings import RUN_PATH
22
23 logger = logging.getLogger(__name__)
24
25 def get_default_timezone() -> ZoneInfo:
26     """
27     Return the default timezone.
28     """
29     return ZoneInfo("Europe/Zurich")
30
31 @dataclass(kw_only=True)
32 class Configuration:
33     """
34     Class for loading and saving the configuration.
35     """
36
37     from ..device.device_manager import DeviceManager
38
39     debug: bool = field(default=True)
40     device: Device = field(
41         default_factory=DeviceManager.get_default,
42     )
43     container: ContainerInfo = field(
44         default_factory=ContainerInfo.load,
45     )
46     run_uuid: uuid.UUID = field(default_factory=uuid.uuid4)
47     email: EmailConfiguration | None = field(
48         default=None,
49     )
50     ai: AiConfiguration | None = field(default=None)
51     loki: LokiConfiguration | None = field(default=None)
52     drawio_mcp: DrawioMCPConfiguration | None =

```

```
1 import click
2
3 from .ai_agent import ai_agent_group
4 from .benchmark import benchmark_group
5 from .config import config_group
6 from .device import device_group
7 from .drill import drill_group
8 from .environment import environment_group
9 from .scenario import kill_all_scenarios, scenario_group
10
11
12 @click.group()
13 def cli() -> None:
14     """
15     CLI root.
16     """
17     # Load the configuration for initialization
18     from nsak.core import config # noqa
19
20
21 cli.add_command(scenario_group)
22 cli.add_command(drill_group)
23 cli.add_command(device_group)
24 cli.add_command(environment_group)
25 cli.add_command(ai_agent_group)
26 cli.add_command(config_group)
27 cli.add_command(benchmark_group)
28 cli.add_command(kill_all_scenarios, name="killswitch")
```

---

Listing 3: NSAK CLI Initialisation

Mutations to the runtime configuration always follow the same pattern: a manager method (e.g. `DeviceManager.load()`) updates the relevant field on the config object and calls `config.save()` to persist the change. Custom YAML representers are registered for `IPInterface` and `PosixPath` so that these types are serialised as plain strings rather than Python-specific tags. Listing 4 shows an example of the persisted `run/config.yaml` after loading the `bananapi_r4` device.

```
1 debug: true
2 device:
3   id: bananapi_r4
4   name: BananaPI R4
5   path: <base-path>/lib/devices/bananapi_r4
6   author: vonal3@bfh.ch
7   repository: <repository-url>
8   configuration:
9     network:
10      ethernets:
11        lan1@eth0:
12          addresses:
13            10.10.10.30/24:
14              is_target: true
15              is_management: false
```

---

Listing 4: Persisted Runtime Configuration (run/config.yaml)

### 5.1.2. Device Configuration Management

In the context of “Project 2”[1.2](#), this resource type was already defined conceptually, and the BananaPI R4 was added as an example in the library under `lib/devices/bananapi_r`. But the resource type was not actually added to the NSAK core, and to manage device configurations, we need to implement the Device resource type first.

#### Create New Device Resource Type

Similar to the other resource types, a YAML specification for defining devices in the library and multiple classes for representing and managing devices in the core are required. Because the existing resource YAML specifications do not allow merging YAML files, this implementation proposes a new base structure that fulfills this trait. Section [5.1.2](#) describes the necessary changes to the existing YAML specifications. The listing [5](#) shows the initial specification for a device resource, which will get extended later on with a section for configuration management.

---

```

1 devices:
2   <str:resource-id>:
3     id: <str:device-id>
4     name: <str:device-name>
5     description: <str:device-description>
6     author: <str:author-email>
7     repository: <str:resource-repository>

```

---

Listing 5: Device Resource YAML Specification

In the table 5.1 below, the required Python classes are listed and described. The classes will be located under `src/nsak/core/device` within their respective files.

| Class         | File Name         | Description                                                                                     |
|---------------|-------------------|-------------------------------------------------------------------------------------------------|
| Device        | device.py         | Dataclass representing a device resource at runtime, closely reflecting the YAML specification. |
| DeviceLoader  | device_loader.py  | Loads and returns Device instances from the resource library.                                   |
| DeviceManager | device_manager.py | Implements business logic and exposes an API for use by the CLI.                                |

Table 5.1.: Device Resource Classes

In “Project 2”[1.2](#), it became already clear that the described classes share a lot of boilerplate code between each other. For this reason, the base classes described in [5.1](#) were added to eliminate duplicated logic, thereby improving compliance with the Don’t Repeat Yourself principles. The classes will be located under `src/nsak/core/resource` within their respective files and the adoption of the existing classes is described in section [5.1.2](#).

| Class           | File Name           | Description                                                                   |
|-----------------|---------------------|-------------------------------------------------------------------------------|
| Resource        | resource.py         | Base class for all resources: Scenario, Drill, Environment and Device.        |
| ResourceLoader  | resource_loader.py  | Base class which handles searching and loading resources from the filesystem. |
| ResourceManager | resource_manager.py | Base class which defines common methods for all resource manager classes.     |

Table 5.2.: Resource Python Classes

To complete the integration of the device resource, the new functionality must be exposed via CLI. The available bash commands are shown in the following listing 6.

---

```
1 # Show usage, options and subcommands for the device resource
2 nsak device --list
3
4 # List all devices found in the resource library
5 nsak device list
6
7 # Output:
8 › nsak device list
9 ID           NAME           PATH
10 bananapi_r4  BananaPI R4  <repository>/lib/devices/bananapi_r4
```

---

Listing 6: NSAK CLI: List Devices

### Device Configuration Management Implementation

To allow a device to expose its configuration options, the YAML specification needs to be extended accordingly. For now, the focus is only on network configuration, as this is the most important information for the scenarios and drills. The specification was designed so that other sections can be added later. The following excerpt 7 shows the specification for the configuration section of the device resource YAML.

---

```
1 devices:
2   <str:resource-id>:
3     # ...
4     configuration:
5       network:
6         ethernets:
7           <str:interface-name>:
8             addresses:
9               <str:cidr>:
10                is_target: <boolean>
11                is_management: <boolean>
```

---

Listing 7: Device Resource Configuration YAML Specification

The following code snippet, located under `src/nsak/core/device/device_loader.py` 8

shows how the YAML configuration gets parsed to the respective datastructures defined in `src/nsak/core/network/configuration.py` and `src/nsak/core/device/device.py`.



---

```

20 class DeviceLoader(ResourceLoader[Device]):
21     """
22     Finds and loads devices from the library paths.
23     """
24
25     ResourceClass = Device
26
27     @classmethod
28     def _parse_device_configuration(
29         cls, data: dict[str, Any]
30     ) -> DeviceConfiguration | None:
31         if "configuration" not in data:
32             return None
33         network = data["configuration"].get("network") or {}
34         if network == "auto":
35             ethernets = {}
36             for interface in get_network_interfaces():
37                 ethernets[interface.name] = {
38                     "addresses": {
39                         ip: {
40                             "is_target": True,
41                             "is_management": False,
42                         }
43                         for ips in interface.ips.values()
44                         for ip in ips
45                     }
46                 }
47         else:
48             ethernets = network.get("ethernets") or {}
49
50         return DeviceConfiguration(
51             raw=data["configuration"],
52             network=NetworkConfiguration(
53                 ethernets={
54                     interface: EthernetConfiguration(
55                         name=interface,
56                         addresses={
57                             ip: IPConfiguration(
58                                 ip=ip_interface(ip),
59
60                                 ↪ is_target=address.get("is_target",
61                                 ↪ False),
62
63                                 ↪ is_management=address.get("is_management",
64                                31 ↪ False),
65                             )
66                             for ip, address in
67                                 ↪ ethernet.get("addresses",
68                                 ↪ {}).items()
69                         },
70                     ),
71                 }
72             )
73         for interface, ethernet in ethernets.items()
74     )
75

```

```
12 @dataclasses.dataclass(frozen=True, kw_only=True)
13 class IPConfiguration:
14     """
15     Represents an ip configuration on an interface.
16     """
17
18     ip: IPInterface
19     is_target: bool
20     is_management: bool
21
22
23 @dataclasses.dataclass(frozen=True, kw_only=True)
24 class EthernetConfiguration:
25     """
26     Represents an ethernet node.
27     """
28
29     name: str
30     addresses: dict[str, IPConfiguration]
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593
594
595
596
597
598
599
600
601
602
603
604
605
606
607
608
609
610
611
612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701
702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781
782
783
784
785
786
787
788
789
790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896
897
898
899
900
901
902
903
904
905
906
907
908
909
910
911
912
913
914
915
916
917
918
919
920
921
922
923
924
925
926
927
928
929
930
931
932
933
934
935
936
937
938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962
963
964
965
966
967
968
969
970
971
972
973
974
975
976
977
978
979
980
981
982
983
984
985
986
987
988
989
990
991
992
993
994
995
996
997
998
999
```

---

Listing 9: Device Configuration Datastructures

The following listing shows the newly available CLI commands for managing device configurations [10](#):

---

```
1 # List all subcommands for the Device resource
2 nsak device --help
3
4 # List all available devices
5 nsak device list
6
7 # Show the device details and configuration
8 nsak device show <str:device-id>
9
10 # Load a device into in the NSAK configuration
11 nsak device load <str:device-id>
12
13 # Show the currently loaded device stored in the NSAK
14   ↪ configuration
15 nsak device loaded
16
17 # Reset the loaded device
18 nsak device unload
```

---

Listing 10: Device Configuration CLI Management

### Refactor Existing Resources

The implementation of the device resource [5.1.2](#) introduced an improved base structure for the resource YAML specifications and three base classes which are abstracting the common structure and functionality. This section describes the resulting changes in the existing code and resource YAML specifications.

As already stated in section [5.1.2](#) all resource type specific classes were refactored to inherit from their respective base classes as described in table [5.1](#). With the example of the `EnvironmentLoader`, the diff below [5.1](#) illustrates how the abstraction of common properties and methods allows for the removal of a lot of boilerplate code.




Figure 5.1.: EnvironmentLoader Base Class Refactoring Diff

The structure of all resources was aligned by nesting the resource under their respective resource type key (scenarios, drills, environments, and devices), so that we could merge all YAML resource files into a single library data structure without key conflicts. In the YAML specification [11](#) below, the new base structure derived from the device resource [5.1.2](#) is illustrated. All existing resource YAML files in the library were adjusted accordingly, and thanks to the introduction of the ResourceLoader base class, the required adjustments in the code were minimal.

```
1 # Resource type version was removed, as it was not used anyway
2
3 # scenarios|drills|environments|devices
4 <str:resource-type-key>:
5     # unique resource identifier
6     <str:resource-id>:
7         # The "metadata" key is removed and the children are now
8         ↪ nested directly under the resource
9         name: <str:name>
10        description: <str:description>
11        author: <str:email>
12        repository: <str:link>
13        # ... resource type specific properties, which
14        ↪ previously were at the top level
```

---

Listing 11: Mergeable Resource YAML Specification

### 5.1.3. Scenario & Drill Configuration Management

Both the scenario and drill resource types were already fully implemented in “Project 2”[1.2](#). Their YAML specifications already contained an interface section describing argument and return types, but this information was purely documentary and had no effect at runtime.

As part of this project, the `interface.arguments` section was extended [12](#) so that each argument carries an explicit `type` annotation and an optional default value.

```
1 <str:resource-type-key>:
2     <str:resource-id>:
3         # Unchanged properties...
4
5     interface:
6         arguments:
7             <str:argument-name>:
8                 type: <str:type-annotation>
9                 default: <any:default-value>
10        return_type: <str:type-annotation>
```

---

Listing 12: Updated Interface Specification

The CLI reads these annotations at runtime to generate typed command options

for each drill and scenario automatically, without requiring any additional hand-written code. The `discover_hosts` drill in listing 13 illustrates this: its single interface argument of type `str` is translated directly into a `--interface` option on the generated CLI command.

---

```
1 drills:
2   discover_hosts:
3     name: Discover Hosts
4     author: vonal3@bfh.ch
5     repository:
6       ↪ https://gitlab.ti.bfh.ch/gausf1-vonal3/swiss-army-knife-network-sniffe
7
8     dependencies:
9       system:
10        - arp-scan
11        - nmap
12       python: [ ]
13
14     interface:
15       arguments:
16         interface:
17           type: str
18         subnet:
19           type: str
20           default: null
21       return_type: NetworkDiscoveryResultMap
```

---

Listing 13: `discover_hosts` Drill YAML

Scenarios follow the same interface convention and additionally declare the drills and environments they depend on. The `mitm` scenario in listing 14 chains the `discover_hosts`, `arp_spoof`, and `transparent_tcp_proxy` drills together and exposes a single typed interface argument, which the CLI surfaces in the same way as for individual drills.

```
1 scenarios:
2   mitm:
3     id: mitm
4     name: MITM Scenario
5     author: vonal3@bfh.ch
6     repository:
7       ↪ https://gitlab.ti.bfh.ch/gausfl-vonal3/swiss-army-knife-network-sniffer
8
9     drills:
10       - discover_hosts
11       - arp_spoof
12       - transparent_tcp_proxy
13
14     environments:
15       - simple_tcp_client_server
16
17     dependencies:
18       system:
19         - arp-scan
20         - dsniff
21         - iptables
22       python:
23         - scapy
24
25     interface:
26       arguments:
27         interface:
28           type: str
29       return_type: None
```

---

Listing 14: mitm Scenario YAML

## 5.2. Reproducible Test Environment

Unless stated otherwise, all referenced implementation artifacts in this section originate from the NSAK repository [3].

The scope of this thesis is confined to the AI based reconnaissance and optionally, the intrusion detection scenarios [A.3.1](#), including the design and implementation of corresponding drills.

To ensure a controlled and safe evaluation of the implementation, a simulated test environment has to be established. This approach prevents unintended interference with production systems and ensures reproducibility during the development and testing phases.

For this exact purpose the “Environment”[3.1.2](#) resource type was implemented into the NSAK framework. In the predecessor project [1.2](#), a simplified network topology comprising a client (Alice) and a server (Bob) is used to emulate and analyze a Man In The Mittle (MITM) scenario. This environment is described in a Docker Compose file, located under: `lib/environments/simple_tcp_client_server/compose.yaml`.

The whole setup can be simulated with the following command: `nsak environment simulate simple_tcp_client_server mitm`.

...

While Docker or Podman provides a general-purpose containerization platform, they abstract away most of the lower levels of the network stack and thus lack the support for modeling network topologies. In the existing implementation this was partially addressed, by using the `macvlan` network driver, to enable level 2 networking. In contrast, Containerlab builds upon Docker Compose and optimize its functionalities for networking purposes.

The primary advantages of Containerlab are:

- ▶ It’s ability to create a reproducible containerized network topology based on a simple configuration file
- ▶ Support multiple interfaces nodes which closely resemble real-world network devices (in comparison Docker typically operates with a single network interface).
- ▶ The rapid creation, extension and teardown of complete network environments, ensure that each experiment starts from a clean and controlled state.

In summary, while OCI containers serves as the underlying container engine, Containerlab provides a higher-level network simulation.

The following listing [15](#) shows an excerpt of a Containerlab topology.yaml file.

Maybe it would make sense to describe this in the project 2 section instead



Basically the nodes are defined and with virtual Ethernet links connected. The bridge represents a multi port switch.

---

```
1 topology:
2   nodes:
3     node1:
4       kind: linux
5       image: alpine:3.19
6       cap-add:
7         - NET_ADMIN
8       sysctls:
9         net.ipv4.ipforward: 1
10
11     node2:
12       kind: linux
13
14   links:
15     - endpoints: [ "node1:eth1", "bridge:nd1" ]
```

---

Listing 15: Excerpt of the Containerlab topology definition

### 5.2.1. Network Topology

The network topology shown in Figure 5.2 illustrates the IP addressing scheme and the interconnections between the individual components.

The `external` node represents an external attacker interacting with the network from the Wide Area Networks (WAN).

The `nsak` container can be deployed in different network segments, enabling the simulation of various attacker positions and reconnaissance scenarios.

update  
without  
external  
node




Figure 5.2.: topology-test-environment

### 5.2.2. DMZ-Server

To evaluate reconnaissance scenarios, common services must run on endpoints.

Berkley Internet Name Domain (BIND) is a complete implementation of the DNS protocol that can be configured through the `named.conf` file as an authoritative name server and resolver.

In the test lab illustrated in 5.2, the functionality of the web server and DNS server was combined for simplicity.

During the De-Militarized Zone (DMZ)-server container build process, a setup shell script is executed.

In the first part, the authoritative BIND `named.conf` file, DNS options, and master zones for forward and reverse lookups are created.

In the forward zone, domain names map to IP addresses, while in the reverse zone, IP addresses map to domain names 56.

add  
filepath  
and  
combine  
with  
other  
sections

```

; <<>> 016 9.18.44 <<>> @172.16.1.10 vlab.local AXFR
; (1 server found)
;; global options: +cmd
vlab.local.      300    IN      SOA     ns1.vlab.local. hostmaster.vlab.local.vlab.local. 2024031501 3600 900 604800 300
vlab.local.      300    IN      NS      ns1.vlab.local.
alice.vlab.local. 300    IN      A       192.168.10.100
bob.vlab.local.  300    IN      A       192.168.10.101
ctrl.vlab.local. 300    IN      A       192.168.10.5
file.vlab.local. 300    IN      A       192.168.10.5
ns1.vlab.local.  300    IN      A       172.16.1.10
printer.vlab.local. 300    IN      A       192.168.10.50
web.vlab.local.  300    IN      A       172.16.1.10
vlab.local.      300    IN      SOA     ns1.vlab.local. hostmaster.vlab.local.vlab.local. 2024031501 3600 900 604800 300
;; Query time: 0 msec
;; SERVER: 172.16.1.10#53(172.16.1.10) (TCP)
;; WHEN: Sat Apr 04 13:34:18 UTC 2026
;; XFR size: 10 records (messages 1, bytes 333)

```

Figure 5.3.: Authoritative Zone Transfer Snippet

The output of a full zone transfer illustrated in figure 5.3 displays the resolved IP addresses and domain names. In our test environment, the restriction of Authoritative Zone Transfer (Full Zone Transfer) requests to only trustworthy IPs is intentionally omitted; the devices are deliberately left unhardened.

In the second part, a minimal NGINX web server with HTTP and HTTPS support is created. Like in real setups, HTTP traffic on port 80 is redirected to HTTPS on port 443. To provide an artifact for reconnaissance purposes, a file named `{secret.txt}` containing sensitive data is placed on the server.

The LAN zone of the network, consists of 4 devices.

### 5.2.3. Domain Control and File Server

The internal server, reachable at 192.168.10.5, combines the role of domain controller and file server into a single node.

It exposes Secure Shell (SSH) on port 22, Server Message Block (SMB) on ports 139 and 445, and Lightweight Directory Access Protocol (LDAP) on port 389.

The LDAP service is intentionally misconfigured to permit anonymous read access to the entire `dc=lab,dc=local` directory tree, exposing user accounts and group memberships without authentication.

The SMB service provides three shares: `public` is readable without credentials, while `finance` and `it` require valid user accounts.

Sensitive files such as payroll exports and budget summaries are placed in the restricted shares as intentional post-exploitation artifacts 57.

#### 5.2.4. Workstations

Two workstations reside on the LAN segment: Alice at 192.168.10.100, assigned to the Finance department, and Bob at 192.168.10.101, assigned to the IT department.

Both nodes expose SSH on port 22 and SMB on port 445, and both SSH banners intentionally disclose the organization name.

On Alice, the home directory of user `asmith` contains a confidential finance report and a drive mapping file that reveals the path to the central file share on the domain server.

On Bob, the home directory of user `bjones` contains a `scripts` directory with a host inventory that lists all internal IP addresses and their respective roles [58](#).

#### 5.2.5. Printer

The printer service, reachable at 192.168.10.50, provides two different interfaces. The first web interface is administrative and is available on port 80. It shows internal printer metadata, such as device name and contact information. The printer version is intentionally disclosed in this admin view. The second interface, a REST API, listens on port 631 at the `{/jobs}` endpoint, where it displays a predefined list of print jobs. All printer requests are logged in `{/var/logs/printer_sim_logs}` inside the container [59](#).

#### 5.2.6. Introduced Vulnerabilities

For reconnaissance purposes, several vulnerabilities and flaws were introduced, which are partly mentioned in the previous sections and are presented in detail in Listing [61](#). These flaws are necessary to compare the results of the larger and smaller models as well as the implemented vs the AI-enhanced scenarios.

### 5.3. AI –Scenarios

This section describes the design and implementation of the milestone AI Scenarios [A.3.1](#), which is based on the user story map workshop AI Purple Team [A.11.7](#). In the workshop we defined the features which such a scenario may provide and categorized them in “Must”, “Could” and “Should”. Because of time constraints, we set the focused only on the user stories categorized as “Must”. The main goal is to implement an AI reconnaissance scenario, which is described in the [5.3.8](#). But on the way it made sense to also create an intermediary scenario which can run custom prompts [5.3.7](#). Further, we already explored the AI based intrusion detection scenario [5.3.9](#) which should be seen as proof of concept, because the corresponding milestone and goals are categorized as “Should”.

#### 5.3.1. AI Agent

An AI agent is an autonomous system that combines a large language model (LLM) with a set of tools it can invoke to accomplish a goal. Rather than just generating text, the agent follows a reasoning loop: it plans which tool to call, executes it, observes the result, and continues until the task is complete [\[22\]](#).

The project team chose [LangChain](#) as the agentic framework without comparing it with alternative solutions, for the following reasons:

- ▶ Free and open-source
- ▶ Python package with native integration into the project’s toolchain
- ▶ Suggested by various sources and has a good reputation in the community
- ▶ Extensive documentation and active community support

For the integration into the NSAK framework we decided to implement an `AiAgent` class as the central abstraction, which wraps a `BaseChatModel` together with a set of `LangChain` tools and delegates the reasoning loop to `LangChain`’s built-in `create_agent()`.

#### AI Configuration

Before any AI based scenario can run, the `AiConfiguration` entry must be set in the runtime configuration, containing the model, provider and credentials. Listing [16](#) shows the dataclass, which holds two fields: `model` and `base_url`. For authentication, the basic auth can be provided in the `base_url` field, otherwise the `api_key` must be set.

Describe the three different abstraction levels: Deep-Agent, LangChain and Lang-Graph (maybe include cheat-sheet with proper citation)

```
1 from __future__ import annotations
2
3 import dataclasses
4
5
6 @dataclasses.dataclass(kw_only=True)
7 class AiConfiguration:
8     """
9     Configuration for the AI agent backend.
10    """
11
12    provider: str
13    model: str
14    base_url: str
15    api_key: str | None = dataclasses.field(default=None)
```

---

Listing 16: AI Configuration Dataclass

With the CLI command `nsak config set ai` the configuration can be set interactively, alternatively the values can be set individually like in the following example:

- ▶ `nsak config set ai.provider "ollama"`
- ▶ `nsak config set ai.model "qwen2.5:7b-instruct"`
- ▶ `nsak config set ai.base_url "http://user:pass@localhost:11434"`
- ▶ `nsak config set ai.api_key "secret"`

### Provider Abstraction

The `AiAgent.create_model()` method parses the provider prefix and dispatches to the appropriate LangChain chat model class via `PROVIDER_MAP`, as shown in listing 17. This means any supported provider can be swapped out by changing a single configuration value, without touching scenario code. Most providers allow to set the “temperature” of a model, where lower values produce more deterministic responses, while higher values increase creativity and diversity [23]. For our intents and purposes always set `temperature=0`, as we want to make agent behavior as deterministic as possible.

```
27 )
28 from nsak.core.configuration import Configuration, config
29
30 logger = logging.getLogger(__name__)
31
32
33 class TrackToolCallMiddleware(AgentMiddleware): # type: ignore
34     """
35     Tool tracking middleware.
36     """
37
38     def __init__(self, tools_available: Iterable[str]) -> None:
39         """
40         Init tool tracking middleware.
41         """
42         self.tools_called: dict[str, list[str]] = {
43             tool: [] for tool in set(tools_available)
44         }
45
46     async def awrap_tool_call(
47         self,
48         request: ToolCallRequest,
49         handler: Callable[[ToolCallRequest],
50             Awaitable[ToolMessage | Command[Any]]],
51     ) -> ToolMessage | Command[Any]:
52         """
```

---

Listing 17: AI Provider Map and MCP Client Initialisation

### 5.3.2. LangChain Concepts

LangChain provides a framework for developing applications powered by large language models (LLMs) and agent-based systems. It enables developers to combine modular components and external integrations, simplifying AI application design while maintaining flexibility as underlying technologies evolve [24].

The primary LangChain concepts used in the NSAK framework are:

- ▶ `BaseChatModel` - provider-agnostic interface for the underlying LLM.
- ▶ `BaseTool` - interface for callable tools that the agent can invoke.
- ▶ `create_agent()` - builds a LangGraph-based ReAct agent graph that wires

the model and tools together.

- ▶ messages - a list of conversation messages (user, AI, tool) maintained across the agent's reasoning loop.



---

```

1  import asyncio
2  import logging
3  from pprint import pformat
4  from typing import Any, AsyncGenerator, Awaitable, Callable,
    ↪ Iterable, Sequence
5
6  from langchain.agents import create_agent
7  from langchain.agents.middleware import AgentMiddleware
8  from langchain.agents.structured_output import ToolStrategy
9  from langchain_anthropic.chat_models import ChatAnthropic
10 from langchain_core.callbacks import AsyncCallbackHandler
11 from langchain_core.language_models import BaseChatModel
12 from langchain_core.messages import ToolMessage
13 from langchain_core.outputs import LLMResult
14 from langchain_core.tools import BaseTool
15 from langchain_mcp_adapters.client import MultiServerMCPClient
16 from langchain_ollama.chat_models import ChatOllama
17 from langchain_openai.chat_models import ChatOpenAI
18 from langgraph.prebuilt.tool_node import ToolCallRequest
19 from langgraph.types import Command
20
21 from nsak.core.ai.tools import (
22     cli_tool,
23     generate_drill_tools,
24     host_configuration,
25     human_interaction_hook,
26     send_email,
27 )
28 from nsak.core.configuration import Configuration, config
29
30 logger = logging.getLogger(__name__)
31
32
33 class TrackToolCallMiddleware(AgentMiddleware): # type: ignore
34     """
35     Tool tracking middleware.
36     """
37
38     def __init__(self, tools_available: Iterable[str]) -> None:
39         """
40         Init tool tracking middleware.
41         """
42         self.tools_called: dict[str, list[str]] = {
43             tool: [] for tool in set(tools_available)
44         }
45
46     async def awrap_tool_call(
47         self,
48         request: ToolCallRequest,
49         handler: Callable[[ToolCallRequest],
50             ↪ Awaitable[ToolMessage | Command[Any]]],
51     ) -> ToolMessage | Command[Any]:
52         """
53         Track async tool calls

```

The `AiAgent.__init__()` method calls `create_model()` to instantiate the provider-specific chat model and then passes it to `create_agent()` together with the tool list and a fixed system prompt that establishes the cybersecurity simulation context. The `run()` method invokes the agent once and streams all resulting messages (user, AI, tool calls, tool results) as formatted strings. The `run_interactive()` method extends this into an interactive loop, where after each agent response the operator is prompted for the next instruction, and the full conversation history is retained in the messages list until the operator types an empty line or `exit`.

### Tool Calling Concepts

Tool calling is the mechanism by which the LLM declares that it wants to invoke an external function. The framework intercepts the model's output, parses the tool call, executes the corresponding Python function, MCP, or Skill, and appends the result to the message history before asking the model to continue. This loop repeats until the model produces a final text response with no tool calls.

In NSAK, tools can be provided through several complementary mechanisms:

- ▶ **LangChain Tools** - native Python functions decorated with `@tool`, which LangChain automatically wraps with type-safe JSON schemas for the model. This is the primary mechanism used for NSAK-specific tools (host configuration, CLI access, e-mail, human interaction hook). Further, all NSAK drills are wrapped as LangChain tools via `generate_drill_tools()`, giving the agent type-safe access to all drill capabilities.
- ▶ **MCP** - the Model Context Protocol allows external tool servers to expose additional capabilities; used here for the draw.io integration (see 5.3.2).
- ▶ **CLI Access** - the `cli_tool` gives the agent direct shell access, allowing it to run any installed binary (nmap, tshark, etc.).

### Tool Calling Models

Considering its limited size of 7 billion parameters, we were quite happy with our first test model “qwen2.5:7b-coder”. But as it turns out, not all models are capable of calling tools. Often, multiple variants of the same model are available, and they don't only differ in size, but also for different strengths or capabilities such as tool calling. The tool calling variants are often have the suffix “instruct”. Luckily, “qwen2.5:7b-coder” has such variant called “qwen2.5:7b-instruct”.

The `create_ai_agent()` factory function shown in listing 19 is the central assembly point used by all three AI scenarios. It reads the ai configuration block, assembles the tool list (always including `cli_tool`, `host_configuration`, and

Explain the issue with limited context windows and high token usage of certain MCP tools, such as drawio mcp

`send_email`), optionally adds the human interaction hook or drill tools, loads MCP tools with graceful fallback, and returns a configured `AIAgent` instance.

```
173     model: str,
174     base_url: str,
175     api_key: str | None,
176     tools: Sequence[BaseTool | Callable[..., Any] |
        ↳ dict[str, Any]] | None = None,
177     middleware: Sequence[AgentMiddleware] | None = None,
178     debug: bool = False,
179     system_prompt: str | None = None,
180     response_format: Any = None, # noqa: ANN401
181 ) -> None:
182     """
183     Initializes the AiAgent and the connection to its
        ↳ backend model.
184     """
185     self.usage = UsageCallback()
186
187     if system_prompt is not None:
188         self.system_prompt = system_prompt
189
190     if tools is None:
191         tools = []
192
193     self.track_tool_call_middleware =
        ↳ TrackToolCallMiddleware(
194         tools_available=[tool.name for tool in tools] #
        ↳ type: ignore
195     )
196
197     middlewares: list[AgentMiddleware] =
        ↳ [self.track_tool_call_middleware]
198     if middleware is not None:
199         middlewares.extend(list(middleware))
200
201     if model == "gpt-oss:120b" and response_format is not
        ↳ None:
202         # Hack for gpt-oss:120b as langchain uses
        ↳ ProviderStrategy for this model, which fails.
203         response_format = ToolStrategy(response_format)
204
205     self.model = AiAgent.create_model(
206         provider,
```

---

Listing 19: create\_ai\_agent() Factory Function

The `enable_drill_tools` flag, when set to `True`, calls `generate_drill_tools()`, which dynamically wraps every NSAK drill as a type-safe LangChain tool. This gives the agent access to structured capabilities such as `discover_hosts` and `port_scan` without requiring CLI invocation.

### MCP Integration

The MCP [13] is an open protocol that standardises how AI models connect to external tool servers. Rather than implementing tool calling ad-hoc, MCP servers expose a well-defined interface that any compatible LLM framework can consume [25]. LangChain Agents must be invoked in an event loop to properly support MCP, this forced us to refactor the implementation “`ai_agent.py`” to be async. Further we need to add support for async drills and scenarios, as otherwise we could not run the AI agent in those contexts.

NSAK integrates an external draw.io MCP server to allow the AI agent to create and manipulate diagrams during a scenario run. The client is initialised at module level (listing 17, lines 41–51 using `MultiServerMCPClient` from `langchain-mcp-adapters`, pointing to `https://drawio.hiube.ch/mcp` with Basic authentication.

MCP tools are loaded asynchronously inside `create_ai_agent()` and appended to the tool list. If the server is unavailable, the exception is caught and logged as a warning so the scenarios can still run without diagram support.

### LangChain Tool: Host Configuration

The `host_configuration` tool shown in listing 20 gives the agent a structured view of the host’s current configuration. It returns a dictionary containing the loaded device configuration with interface names, Internet Protocol addresses, and the `is_target` / `is_management` flags loaded from the device YAML. This bootstraps the agent’s situational awareness without requiring it to run `ip addr` or parse raw network output - the agent knows from the start which interfaces are part of the attack surface and which are for management traffic.

---

```

9  @tool # type: ignore[misc]
10 def host_configuration() -> dict[str, Any]:
11     """
12     Returns the host network configuration as a dictionary.
13
14     Use this tool first to understand the local network setup
15     ↪ before running any scans or attacks.
16
17     The returned dictionary has the following structure:
18     {
19         "debug": bool,
20         "device": {
21             "id": str,
22             "name": str,
23             "description": str,
24             "configuration": {
25                 "network": {
26                     "ethernets": {
27                         "<interface_name>": {
28                             "name": str,
29                             "is_up": bool,          # False means
30                             ↪ the interface is DOWN, skip it
31                             "is_target": bool,      # True means
32                             ↪ this interface is used for
33                             ↪ attacks/scans
34                             "is_management": bool # True means
35                             ↪ this interface is used for
36                             ↪ device access
37                             "addresses": {
38                                 "<cidr>": {
39                                     "ip": str,          # e.g.
40                                     ↪ "10.10.10.5/24" - use
41                                     ↪ this as nmap source IP
42                                     "is_target": bool,
43                                     "is_management": bool,
44                                 }
45                             },
46                         },
47                     },
48                 },
49             },
50         }
51     }
52
53     Guidance:
54     - Use `is_target=True` interfaces/IPs when running nmap
55     ↪ scans or attacks
56     - Skip interfaces where `is_up=False`
57     - Use `is_management=True` interfaces for device access or
58     ↪ data extraction
59     - `configuration` may be None if the device has not been
60     ↪ configured yet
61     """
62     return ConfigurationSerializer.serialize(configuration())

```

### LangChain Tool: CLI Tool

The `cli_tool` shown in listing 21 allows the AI agent to execute arbitrary shell commands with a configurable timeout (default 120 s). It returns a tuple of (`exit_code`, `stdout`, `stderr`), giving the agent full feedback to diagnose failures and adapt its next action. The scope is intentionally broad: the agent can invoke `nmap`, `tshark`, `apt`, or any other binary available in the execution environment. Besides the execution of drills, this flexibility is what makes the AI scenarios powerful, but it is also why the kill switch described in 5.3.6 is a necessary safety measure.

```

10 @tool # type: ignore[misc]
11 def cli_tool(command: str, timeout: int = 120) -> tuple[int,
    ↳ str, str]:
12     """
13     Run an arbitrary CLI command and return its output.
14
15     Use nmap for network scanning, e.g.:
16     - "nmap -sV 10.10.10.1" (service version detection)
17     - "nmap -sC -sV -oN output.txt 10.10.10.1" (default scripts
    ↳ + versions)
18     - "nmap -p 1-1000 10.10.10.1" (port range)
19
20     Use nmap NSE scripts for targeted service enumeration with
    ↳ --script <name> or --script <cat>:
21     - Single script: "nmap --script http-title -p 80
    ↳ 10.10.10.1"
22     - Multiple: "nmap --script http-title,http-headers -p
    ↳ 80,443 10.10.10.1"
23     - With args: "nmap --script smb-security-mode
    ↳ --script-args smbuser=guest -p 445 10.10.10.1"
24
25     Service-to-script mapping – run these when the corresponding
    ↳ service is detected:
26     HTTP / HTTPS / http-alt (80, 443, 8080...):
27         http-title, http-headers, http-robots.txt
28         e.g. "nmap --script
    ↳ http-title,http-headers,http-robots.txt -p 80
    ↳ 10.10.10.1"
29
30     DNS / domain (53):
31         dns-zone-transfer, dns-brute
32         e.g. "nmap --script dns-zone-transfer,dns-brute -p 53
    ↳ 10.10.10.1"
33
34     SMTP / SMTPS (25, 465, 587):
35         smtp-commands, smtp-enum-users
36         e.g. "nmap --script smtp-commands,smtp-enum-users -p
    ↳ 25 10.10.10.1"
37
38     FTP (21):
39         ftp-anon, ftp-ls
40         e.g. "nmap --script ftp-anon,ftp-ls -p 21 10.10.10.1"
41
42     SMB / NetBIOS (139, 445):
43         smb-security-mode, smb2-security-mode
44         e.g. "nmap --script
    ↳ smb-security-mode,smb2-security-mode -p 139,445
    ↳ 10.10.10.1"
45
46     LDAP / LDAPS (389, 636):
47         ldap-rootdse
48         e.g. "nmap --script ldap-rootdse -p 389 10.10.10.1"
49

```



### 5.3.3. Human Interaction Hook

#### LangChain Tool: Human Interaction Hook

Even though LangChain has its own concept, called [Human-in-the-loop](#), for now a simple python based tool was implemented as shown in 22, to achieve the same goal. When enabled via the interactive flag in `create_ai_agent()`, the agent can pause its execution and ask the operator a question, then incorporate the answer into its next reasoning step.

```

4 @tool # type: ignore[misc]
5 def human_interaction_hook(question: str) -> str:
6     """
7     Ask the human operator a question and return their answer.
8
9     Use this tool when you need information, confirmation, or
10    ↪ guidance that
11    cannot be determined from the available tools or the current
12    ↪ context.
13
14    Args:
15        question: The question or request to present to the
16        ↪ human operator.
17
18    Returns
19    -----
20        The human operator's response as a string.
21    """
22    return input(f"[Human Input Required]\n{question}\n> ")

```

Listing 22: LangChain Tool: Human Interaction Hook

### 5.3.4. Logging

#### LangChain Agent Logs

Explain the general python logging setup and implementation of NSAK

Explain that the LangChain agent already provides extensive logging, where we

## Remote Logging System: Grafana / Loki

Uses the package [python-logging-loki](#). We went for this method over more conventional approaches, because we can provide more metadata in form of “tags”.

instead of collecting the logs from stdout Filters: Regex: `\[(AI)|(Prompt)|(Tool)\]`

## A.12.10

```

6 @dataclasses.dataclass(kw_only=True)
7 class LokiConfiguration:
8     """
9     Configuration for logging to loki.
10    """
11
12    url: str
13    username: str
14    password: str

```

Listing 23: Loki Logging Configuration

Com-  
plete this  
section

Add a  
figure

## 5.3.5. E-Mail Alerting

The python standard library already includes packages for creating and sending E-Mails via Simple Mail Transfer Protocol (SMTP). The “EmailBackend” implemented to the NSAK framework is an opinionated wrapper around those packages and leverages the previously introduced 5.1 to set up necessary configuration. If no SMTP configuration is provided as specified in 24, the system will write the E-Mails as logs instead. The working tool can be seen in actions in the test run of the AI Reconnaissance scenario 5.3.8, where the operator instructs the agent to send the results via E-Mail.

```
15 @dataclasses.dataclass(kw_only=True)
16 class EmailConfiguration:
17     """
18     SMTP Configuration for sending emails.
19     """
20
21     host: str
22     port: int
23     encryption: EmailEncryptionType = EmailEncryptionType.NONE
24     username: str
25     password: str
26     sender: str
27     receiver: str
```

---

Listing 24: E-Mail Alerting Configuration

```
28 class EmailBackend:
29     """
30     Email Backend used for sending mails via SMTP.
31     """
32
33     backend: smtplib.SMTP | LoggerEmailBackend
34
35     def __init__(self, email_config: EmailConfiguration | None)
36     ↪ -> None:
37         """
38         Set up the SMTP client with the correct configuration
39         ↪ for encryption.
40         """
41         self.config = email_config
42         self._init_backend()
```

---

Listing 25: E-Mail Alerting Implementation

### LangChain Tool: E-Mail Alerting

The `send_email` tool shown in listing 26 exposes the `EmailBackend` to the AI agent. During a scenario run the agent can autonomously send e-mails to operators - for example, when the AI intrusion detection scenario identifies a suspicious host or the operator instructed the AI reconnaissance scenario to send a re-

port. The recipient of the mail is defined via the configuration management and cannot be changed by the agent, so that the operator always knows which mails where sent and who received them. Optional `cc_recipients` and `attachments` parameters allow the agent to inform additional stakeholders or attach diagrams and reports. Each email comes with some predefined For better traceability a template is defined, which is filled with meta information including the current scenario and a unique ID, which can be used to associate the mail in grafana. Further, the subject line is always prefixed with [NSAK], so that operators can filter them in their inbox.

---

```
6 # Hardcoded for now!
7 grafana_link_template =
8     ↳ "https://grafana.hiube.ch/a/grafana-lokiexplore-app/explore/nsak_run_uuid/"
9     template = """
10     Scenario: %(scenario)s\n
11     Run ID: %(run_uuid)s\n
12     Datetime: %(timestamp)s\n
13     \n
14     See logs in Grafana: %(grafana_link)s\n
15     \n
16     -----\n
17     \n
18     %(message)s\n
19     \n
20     """
21
22 @tool # type: ignore[misc]
23 def send_email(
24     subject: str,
25     message: str,
26     cc_recipients: list[str] | None = None,
```

---

Listing 26: LangChain Tool: E-Mail Alerting

### 5.3.6. CLI Kill Switch

It's possible that an AI-Agent starts doing unexpected, unwanted or even destructive actions. This was also identified as a possible risk in the risk assessment [A.7.8](#). This kind of misbehavior, combined with elevated permissions and tool calling capabilities via direct CLI access, could lead to a wide range of issues:

- ▶ The AI-Agent is stuck in a loop while using most system resources.
- ▶ The AI-Agent misconfigures the system, rendering it inaccessible or even unusable.
- ▶ The AI-Agent burns too many API credits because of a bug, an inefficient request, or a huge task.
- ▶ The AI-Agent starts an attack against a production network.

For these reasons we want to be able to kill a specific or all scenarios as fast as possible. As a first measure, the project team decided to add this functionality as simple CLI commands:

- ▶ `nsak scenario kill <scenario-id>`: Kill a specific scenario (sends SIGTERM to the container)
- ▶ `nsak scenario killswitch`: Kill all running scenarios
- ▶ `nsak killswitch`: Alias for `nsak scenario killswitch`

#### 5.3.7. Scenario: AI Agent

The AI Agent scenario is the most generic of the three: it accepts a free-form prompt argument and runs the AI agent in an interactive loop, maintaining the full conversation history across turns. This makes it the primary entry point for ad-hoc AI-assisted tasks where no dedicated scenario exists yet. The operator can guide the agent step by step, observe its reasoning, and issue follow-up instructions until they type an empty line or `exit`. Listing 27 shows the scenario entrypoint and listing 28 shows its YAML interface declaration.

```
1  """
2  Scenario entrypoint for AI based network reconnaissance.
3  """
4  import logging
5
6  from nsak.core import create_ai_agent
7
8  logger = logging.getLogger(__name__)
9
10 async def run(prompt: str, interactive: bool = True) -> None:
11     """
12     Scenario, which runs an AI agent with a custom prompt in an
13     ↪ interactive loop.
14
15     After each agent response the operator is prompted for the
16     ↪ next instruction.
17     Enter an empty line or 'exit' to stop the session.
18
19     :return: None
20     """
21
22     logger.info("Starting scenario: AI Agent")
23
24     ai_agent = await create_ai_agent(interactive)
25     result = ai_agent.run_interactive(prompt)
26
27     async for line in result:
28         logger.warning(line)
```

---

Listing 27: Scenario: AI Agent - scenario.py

---

```

1 scenarios:
2   ai_agent:
3     id: ai_agent
4     name: AI Agent (custom prompt)
5     author: vonal3@bfh.ch
6     repository:
7       ↪ https://gitlab.ti.bfh.ch/gausf1-vonal3/swiss-army-knife-network-sniffe
8
9     drills:
10
11   environments:
12
13   dependencies:
14     system:
15     python:
16
17   interface:
18     arguments:
19     prompt:
20       type: str
21     interactive:
22       type: bool
23       default: True
24     return_type: None

```

---

Listing 28: Scenario: AI Agent - scenario.yaml

## Scenario: AI Agent - Test Run

```
nsak scenario execute ai_agent --prompt "Scan the target network and
report all findings" --interactive True
```

## 5.3.8. Scenario: AI Reconnaissance

The AI Reconnaissance scenario automates host and service discovery on a specified network interface. It encodes a structured five-step plan directly into the initial prompt: the agent is instructed to call `host_configuration` to determine the Internet Protocol addresses and subnets on the interface, then use the `cli` tool to invoke `nmap` for host scanning, then again for service enumeration, and finally to report all findings including vulnerability indicators. If the interactive

Describe and add a reference to the result of the test runs and the limited tool calling capabilities of smaller models

flag is set, the last step invites the operator to issue follow-up instructions via the `human_interaction_hook`.

This prompt-driven approach means the scenario does not contain any network scanning logic itself - the agent decides which `nmap` flags to use, how to interpret results, and what to scan next. Listing 29 shows the scenario entrypoint and listing 30 shows its YAML interface declaration.



---

```

1  """
2  Scenario entrypoint for AI based network reconnaissance with
   ↳ structured output.
3
4  This currently works only with frontier models.
5  """
6  import json
7  import logging
8  from dataclasses import dataclass
9
10 from nsak.core import create_ai_agent, config, AiAgent
11 from nsak.core.network import EnumerateServicesResult
12 from nsak.core.network.enumerate_services_result import
   ↳ EnumerateServicesResultEntry
13 from
   ↳ nsak.core.scenario_results.ai_reconnaissance_scenario_result
   ↳ import AIReconnaissanceScenarioResult
14 from nsak.core.scenario_results.network_discovery_result import
   ↳ NetworkDiscoveryTable, NetworkDiscoveryRow
15 from nsak.core.scenario_results.reconnaissance_scenario_result
   ↳ import ReconnaissanceScenarioResult
16
17 logger = logging.getLogger(__name__)
18
19 reconnaissance_prompt_template = """
20 Steps:
21 1. Retrieve the network configuration with the
   ↳ `host_configuration` tool.
22 2. Discover all subnets, hosts and services with nmap on the
   ↳ following interface: %(interface)s
23 3. Enumerate all services based on the result of the network
   ↳ discovery result with service-specific nmap NSE scripts
24 4. Write a markdown formatted assessment of your findings
25 5. Return the result as structured output: Result
26
27 """
28
29 @dataclass(frozen=True, kw_only=True)
30 class Result:
31     """
32     Adhoc result type.
33     """
34     result: ReconnaissanceScenarioResult
35     assessment: str
36
37     63
38     async def run_reconnaissance_agent(
39         interface: str,
40         interactive: bool = False,
41     ) -> tuple[Result, AiAgent]:
42         """
43         Run an agent which returns a NetworkDiscoveryResultMap.
44         """
45         prompt = reconnaissance_prompt_template % {"interface":
   interface}

```

```

1 scenarios:
2   ai_reconnaissance:
3     id: ai_reconnaissance
4     name: AI Reconnaissance
5     author: vonal3@bfh.ch
6     repository:
7       ↪ https://gitlab.ti.bfh.ch/gausf1-vonal3/swiss-army-knife-network-sniffer
8
9   drills:
10
11  environments:
12    - simple_tcp_client_server
13
14  dependencies:
15    system:
16    python:
17
18  interface:
19    arguments:
20      interface:
21        type: str
22      interactive:
23        type: bool
24        default: False
25    return_type: AIReconnaissanceScenarioResult

```

Listing 30: Scenario: AI Reconnaissance - scenario.yaml

## Scenario: AI Reconnaissance - Test Run

```
nsak scenario execute ai reconnaissance --interface eth0
```

## 5.3.9. Scenario: AI Intrusion Detection

The AI Intrusion Detection scenario is a proof of concept, as it was specified as “Should” goal. It demonstrates a two-phase approach: first it performs host discovery using `DrillManager.execute("discover_hosts")`, then it captures live network traffic using `tshark` (a system dependency declared in the scenario YAML). Up to 200 packets are captured over a 60-second window and extracted as Comma Separated Values (CSV) fields (frame number, source and destination

Describe and add a reference to the result of the first testrun

Media Access Control, source and destination Internet Protocol, Transmission Control Protocol/User Datagram Protocol (UDP) destination port, protocol stack, and frame length).

Both the host discovery table and the packet CSV are injected into a structured PROMPT\_TEMPLATE that instructs the AI to identify at most five suspicious events, each with source/destination Internet Protocol and Media Access Control, protocol, and the reason for suspicion, and to identify any malicious hosts. By constraining the number of events the prompt keeps the agent's output concise and actionable. Listing 31 shows the full scenario entrypoint and listing 32 shows its YAML interface declaration.

---

```

1  """
2  Scenario entrypoint for AI based intrusion detection.
3  """
4  import logging
5  import subprocess
6
7  from nsak.core import create_ai_agent, DrillManager
8  from nsak.core.network import NetworkDiscoveryResultMap
9
10 logger = logging.getLogger()
11
12 PACKET_COUNT = 200
13 CAPTURE_DURATION = 60 # seconds
14 PROMPT_TEMPLATE = """
15 You are analyzing network traffic for intrusion detection on
16   ↪ interface %(interface)s.
17
18 Network Discovery Result Map:
19 %(host_discovery_result_map)s
20
21 Captured packets (CSV):
22 %(csv_summary)s
23
24 Identify at most 5 suspicious events if any. For each, output:
25 - src/dst IP and MAC
26 - protocol
27 - reason it is suspicious
28
29 Identify malicious hosts (MAC and/or IP) if any.
30
31 Be concise. Do not repeat packet data.
32 """
33
34 async def run(interface: str, interactive: bool = False) ->
35   ↪ None:
36     """
37     Scenario, which listens to network traffic and tries to
38     ↪ detect intruders.
39
40     :return: None
41     """
42     logger.info("Starting scenario: AI Intrusion Detection")
43
44     host_discovery_result_map: NetworkDiscoveryResultMap =
45     ↪ DrillManager.execute("discover_hosts",
46     ↪ interface=interface)
47
48     # Capture traffic and extract fields including MAC addresses
49     result = subprocess.run(
50         [
51             "tshark",
52             "-i", interface,
53             "-c", str(PACKET_COUNT),
54             "-a", f"duration:[CAPTURE_DURATION]"
55         ]
56     )

```

```

1 scenarios:
2   ai_intrusion_detection:
3     id: ai_intrusion_detection
4     name: AI Intrusion Detection
5     author: vonal3@bfh.ch
6     repository:
7       ↪ https://gitlab.ti.bfh.ch/gausf1-vonal3/swiss-army-knife-network-sniffer
8
9   drills:
10
11  environments:
12    - simple_tcp_client_server
13
14  dependencies:
15    system: [ tshark ]
16    python:
17
18  interface:
19    arguments:
20      interface:
21        type: str
22      interactive:
23        type: bool
24        default: False
25      return_type: None

```

Listing 32: Scenario: AI Intrusion Detection - scenario.yaml

## Scenario: AI Intrusion Detection - Test Run

```
nsak scenario execute ai_intrusion_detection --interface eth0
```

## 5.3.10. LangChain Agent Chat UI

We did not add the [LangChain Agent Chat UI](#) because we made some implicit architecture decisions which prevent easy integration.

If we decide to implement it regardless, we would need to redesign and refactor the AI Agent implementation, for example by switching to a callback-based output model. A useful reference for such a migration is available at [this LangChain UI tutorial](#).

Describe and add a reference to the result of the first testrun

Move this to the LangChain part

## 5.4. Red Team: Reconnaissance Scenario

The reconnaissance scenario is an engineered red team scenario that systematically maps an unknown network. It chains four drills in sequence: host discovery, port scanning, and service enumeration. If no interface is specified by the operator, the scenario first discovers all active physical interfaces and, if any lacks an IP address, requests one via Dynamic Host Configuration Protocol (DHCP) before scanning. The scenario results serve as the reproducible baseline for comparison with the AI-driven reconnaissance approach. The full scenario implementation is provided in Listing 62.

### 5.4.1. Goal

The goal of the scenario is to identify all reachable hosts, their open ports and running services in a target network without prior knowledge of the topology. The drill chain is designed to work generically across different network environments: the operator can either provide a specific interface and subnet or let the scenario discover the available interfaces automatically.

### 5.4.2. Discover Hosts Drill

The *Discover Hosts* drill is the first active step of the scenario. For local network segments, it runs `arp -scan --localnet` on the given interface, which sends Address Resolution Protocol requests and collects the Internet Protocol address, Media Access Control address, and vendor of each responding host. If an explicit subnet is provided, the drill falls back to `nmap -sn` instead, which uses Internet Control Message Protocol (ICMP) and Transmission Control Protocol probes and therefore works across Layer 3 boundaries. Media Access Control addresses are only available for hosts on the same Layer 2 segment; remote hosts are recorded without a Media Access Control. The full implementation is provided in Listing 64.

#### Drill Interface

The drill exposes two arguments: a required `interface` and an optional `subnet`. If `subnet` is omitted, the drill defaults to a local ARP scan; if provided, it switches to the `nmap`-based subnet scan. The resource configuration is shown in Listing 63.

### 5.4.3. Enumeration

After host discovery, the scenario runs two further drills that progressively deepen the collected information.

#### Port Scan

The *Port Scan* drill receives the `NetworkDiscoveryResultMap` produced by host discovery and runs `nmap -sV --open` against each discovered Internet Protocol. For every open port, the drill records the port number, transport protocol (Transmission Control Protocol/UDP), service name, and detected version string and appends a `NetworkService` entry to the existing result map. The full implementation is provided in Listing 65.

#### Service Enumeration

The *Enumerate Services* drill takes the enriched result map and runs targeted Nmap Scripting Engine (NSE) scripts for each discovered service. Well-known services such as HTTP, DNS, FTP, SMTP, SMB, and LDAP are mapped to a specific set of scripts (e.g. `http-title`, `dns-zone-transfer`, `smb-security-mode`). Unknown services fall back to the generic `banner` script, which captures the initial response of any Transmission Control Protocol service. The drill returns a mapping of `ip:port` keys to lists of finding strings extracted from the NSE output. The full implementation is provided in Listing 66.

### 5.4.4. Output Artifacts

The primary output of the scenario is the `NetworkDiscoveryResultMap`, which aggregates all results per interface. Each entry maps an interface name to a `NetworkDiscoveryResult` containing the list of discovered `NetworkService` objects with their endpoints (IP, MAC, port, protocol, version). After port scanning, the scenario prints the result as a human-readable table. The service enumeration findings are printed separately as a structured list grouped by `ip:port`.

## 5.5. Benchmark Suite

As described in the evaluation criteria 6.1.1 there will be at least 120 evaluation runs, which will be a lot of data to analyze and assess. For this reason, we decided to implement a benchmarking suite, consisting of a benchmark harness that executes benchmark runs and summarizes their results into a report. The idea is to fully automate the aggregation of the **quantitative criteria** (duration, token usage) and to prepare the reports needed for the **qualitative criteria** (correctness, creativity, consistency) as good as possible. The key concepts for this implementation are described the table below 5.3:

| Term             | Concept                                      |
|------------------|----------------------------------------------|
| BenchmarkManager | The automation harness                       |
| BenchmarkRun     | Represents a single execution                |
| BenchmarkResult  | The result data type of a single execution   |
| BenchmarkReport  | The collected output aggregating all results |

Table 5.3.: Benchmark Suite Concepts



## 6. Evaluation

| Aspect         | Traditional Recon                                          | AI Recon                                                |
|----------------|------------------------------------------------------------|---------------------------------------------------------|
| File           | lib/scenarios/reconnaissance/scenario.py                   | lib/scenarios/ai_reconnaissance/scenario.py             |
| Return value   | NetworkDiscoveryResultMap (structured)                     | None (LangChain stream only)                            |
| Tools          | arp-scan, nmap (Drills)                                    | LangChain tool_calls (cli, host_configuration)          |
| Async          | sync                                                       | async                                                   |
| Determinism    | deterministic (fixed drill order in code)                  | non-deterministic (LLM selects tools autonomously)      |
| Error handling | explicit in code (e.g. DHCP fallback, interface discovery) | implicit at model level (may skip steps or hallucinate) |

Table 6.1.: Comparison between Traditional Recon and AI Recon

**Limitation:** Traditional Recon returns a structured `NetworkDiscoveryResultMap`, whereas AI Recon returns `None` (log stream only). A direct comparison of results therefore requires either parsing the AI output or documenting this asymmetry as a limitation.

### 6.1. Evaluation Setup

#### 6.1.1. Criteria

**Qualitative criteria:**

- ▶ **Correctness and Task completion/completeness:** Are the scenarios completed fully and are the results correct?
- ▶ **Creativity: AI-specific** Does AI Recon identify relationships or suggest follow-up steps not covered by Traditional Recon?
- ▶ **Consistency network discovery results:** How much does the output vary across  $n = 10$  runs with an identical prompt, per model and environment?

**Quantitative criteria:**

- ▶ **Speed:** Time from scenario start to complete output of all discovered hosts and services
- ▶ **Token usage: AI-specific** Number of tokens consumed per run (AI Recon only)

With  $3 \text{ models} \times 3 \text{ environments} \times n = 10 \text{ runs}$ , this yields 90 measurements for AI Recon.

### 6.1.2. Environments

1. Virtual environment (containerlab)
2. Physical environment
3. BFH lab

### 6.1.3. Models

1. Self-hosted local LLM (Ollama)
2. BFH LLM
3. Frontier model (Anthropic Claude)

We try to fix most of the parameters which influence the success criteria therefore we use a fixed engineered prompt for all models and environments

## 6.2. Todos

1. section for test runs evaluation and implementation
2. human interaction loop and dynamic mcp
3. test harness for automated runs
4. unified result data structure with evaluation criteria, avg speed, tokenusage, summary
5. graphical presentation of result tabel
6. Evaluation Mail Wenger und Urs

### 6.3. Virtual Environment (Containerlab)

Table 6.2.: Nonlinear Model Results

| Kind               | Avg Duration | Avg Tokens | Avg Tool Calls | Hosts | Findings | Services |
|--------------------|--------------|------------|----------------|-------|----------|----------|
| Frontier (claude)  | 262 s        | 83,277     | 20             | 4     | 3        | 9        |
| BFH (gpt-oss:120b) | 383 s        | 20,496     | 5              | 3     | 4        | 9        |
| qwen:30b           | 422 s        | 28,122     | 7              | 3     | 4        | 8        |
| recon              | -            | -          | -              | 10    | 10       | 10       |

**Overview** Table 6.2 summarises the quantitative results of each model across ten runs in the virtual Containerlab environment, compared against the deterministic **recon** baseline. The frontier model (Claude Opus) completed reconnaissance fastest at an average of 262 s per run while also issuing the most tool calls (avg. 20), reflecting a deeper, iterative enumeration strategy. Both the BFH model (gpt - oss : 120b) and the local model (qwen3 : 30b) were slower and used far fewer tool calls (5 and 7 respectively), which is consistent with their shallower service enumeration depth visible in the lower host and findings counts. The reference **recon** scenario, which executes a fixed, deterministic script, serves as the theoretical upper bound with all 10 hosts, findings, and services identified.


Figure 6.1.: Mean run duration per model (Containerlab,  $n = 10$  runs each).

Figure 6.1 visualizes the mean run duration per model. The frontier model is the

fastest despite performing the most tool calls and consuming the most tokens. The local `qwen3:30b` model is the slowest on average (422 s), closely followed of the `gpt:oss:120b`




Figure 6.2.: Mean total token usage per model (Containerlab,  $n = 10$  runs each).

Figure 6.2 shows mean token consumption. The frontier model consumes roughly four times as many tokens as the BFH model (83,277 vs. 20,496), which directly reflects the richer context it accumulates: longer tool outputs, repeated re-analysis of scan results, and verbose final assessments. Notably, `qwen3:30b` and the `gpt:oss:120b` have an almost equal token consumption.




Figure 6.3.: Total tokens vs. run duration, annotated with tool-call count per run (Containerlab).

Figure 6.3 plots total token consumption against run duration for every individual run, with tool-call counts annotated per data point. The frontier model occupies the upper-right region (high tokens, moderate-to-high duration) with a wide spread, reflecting run-to-run variation in how deeply the model chose to enumerate individual services. The BFH and local models cluster in the lower-left (fewer tokens, shorter-to-mid duration), confirming their more uniform but shallower execution profiles.

figure of  
services  
and  
ports

## 6.4. Physical Lab

Table 6.3.: Nonlinear Model Results

| Kind               | Avg Duration | Avg Tokens | Avg Tool Calls | Hosts | Findings | Services |
|--------------------|--------------|------------|----------------|-------|----------|----------|
| Frontier (claude)  | 262 s        | 83,277     | 20             | 4     | 3        | 9        |
| BFH (gpt-oss:120b) | 383 s        | 20,496     | 5              | 3     | 4        | 9        |
| qwen:30b           | 422 s        | 28,122     | 7              | 3     | 4        | 8        |
| recon              | -            | -          | -              | 10    | 10       | 10       |

**Overview** Table 6.2 summarises the quantitative results of each model across ten runs in the virtual Containerlab environment, compared against the deterministic recon baseline. The frontier model (Claude Opus) completed reconnaissance fastest at an average of 262 s per run while also issuing the most tool calls (avg. 20), reflecting a deeper, iterative enumeration strategy. Both the BFH

model (`gpt-oss:120b`) and the local model (`qwen3:30b`) were slower and used far fewer tool calls (5 and 7 respectively), which is consistent with their shallower service enumeration depth visible in the lower host and findings counts. The reference `recon` scenario, which executes a fixed, deterministic script, serves as the theoretical upper bound with all 10 hosts, findings, and services identified.




Figure 6.4.: Mean run duration per model (Containerlab,  $n = 10$  runs each).

Figure 6.1 visualizes the mean run duration per model. The frontier model is the fastest despite performing the most tool calls and consuming the most tokens. The local `qwen3:30b` model is the slowest on average (422 s), closely followed of the `gpt.oss:120b`




Figure 6.5.: Mean total token usage per model (Containerlab,  $n = 10$  runs each).

Figure 6.2 shows mean token consumption. The frontier model consumes roughly four times as many tokens as the BFH model (83,277 vs. 20,496), which directly reflects the richer context it accumulates: longer tool outputs, repeated re-analysis of scan results, and verbose final assessments. Notably, qwen3:30b and the gpt:oss:120b have an almost equal token consumption.




Figure 6.6.: Total tokens vs. run duration, annotated with tool-call count per run (Containerlab).

Figure 6.3 plots total token consumption against run duration for every individual run, with tool-call counts annotated per data point. The frontier model occupies the upper-right region (high tokens, moderate-to-high duration) with a wide spread, reflecting run-to-run variation in how deeply the model chose to enumerate individual services. The BFH and local models cluster in the lower-left (fewer tokens, shorter-to-mid duration), confirming their more uniform but shallower execution profiles.

## 6.5. Cyberlab bfh

Table 6.4.: Nonlinear Model Results

| Kind               | Avg Duration | Avg Tokens | Avg Tool Calls | Hosts | Findings | Services |
|--------------------|--------------|------------|----------------|-------|----------|----------|
| Frontier (claude)  | 262 s        | 83,277     | 20             | 4     | 3        | 9        |
| BFH (gpt-oss:120b) | 383 s        | 20,496     | 5              | 3     | 4        | 9        |
| qwen:30b           | 422 s        | 28,122     | 7              | 3     | 4        | 8        |
| recon              | -            | -          | -              | 10    | 10       | 10       |

**Overview** Table 6.2 summarises the quantitative results of each model across ten runs in the virtual Containerlab environment, compared against the deterministic recon baseline. The frontier model (Claude Opus) completed reconnaissance fastest at an average of 262 s per run while also issuing the most tool calls (avg. 20), reflecting a deeper, iterative enumeration strategy. Both the BFH



model (`gpt-oss:120b`) and the local model (`qwen3:30b`) were slower and used far fewer tool calls (5 and 7 respectively), which is consistent with their shallower service enumeration depth visible in the lower host and findings counts. The reference `recon` scenario, which executes a fixed, deterministic script, serves as the theoretical upper bound with all 10 hosts, findings, and services identified.




Figure 6.7.: Mean run duration per model (Containerlab,  $n = 10$  runs each).

Figure 6.1 visualizes the mean run duration per model. The frontier model is the fastest despite performing the most tool calls and consuming the most tokens. The local `qwen3:30b` model is the slowest on average (422 s), closely followed of the `gpt.oss:120b`




Figure 6.8.: Mean total token usage per model (Containerlab,  $n = 10$  runs each).

Figure 6.2 shows mean token consumption. The frontier model consumes roughly four times as many tokens as the BFH model (83,277 vs. 20,496), which directly reflects the richer context it accumulates: longer tool outputs, repeated re-analysis of scan results, and verbose final assessments. Notably, qwen3:30b and the gpt:oss:120b have an almost equal token consumption.




Figure 6.9.: Total tokens vs. run duration, annotated with tool-call count per run (Containerlab).

Figure 6.3 plots total token consumption against run duration for every individual run, with tool-call counts annotated per data point. The frontier model occupies the upper-right region (high tokens, moderate-to-high duration) with a wide spread, reflecting run-to-run variation in how deeply the model chose to enumerate individual services. The BFH and local models cluster in the lower-left (fewer tokens, shorter-to-mid duration), confirming their more uniform but shallower execution profiles.

## 7. Discussion

## 8. Conclusion

# Bibliography

- [1] World Economic Forum. The global risks report 2026, 2026.
- [2] CrowdStrike. Global threat report 2025, 2025.
- [3] Frank Gauss and Lukas von Allmen. Nsak (network swiss army knife). <https://github.com/nsak-team/nsak>, 2026. Accessed: 2026-01-10.
- [4] Polina Zilberman, Rami Puzis, Sunders Bruskin, Shai Shwarz, and Yuval Elovici. Sok: A survey of open-source threat emulators. arXiv preprint, 2020.
- [5] Blake E. Strom, Andy Applebaum, Doug Miller, Adam Nickels, Cody Pennington, and Chris Thomas. Mitre att&ck: Design and philosophy. Technical report, MITRE Corporation, 2020.
- [6] Bader Al-Sada, Alireza Sadighian, and Gabriele Oligeri. Mitre att&ck: State of the art and way forward. *ACM Computing Surveys*, 2023.
- [7] MITRE Corporation. Mitre caldera documentation, 2024.
- [8] Red Canary. Atomic red team, 2024.
- [9] Rapid7. Metasploit framework documentation, 2024.
- [10] David Kennedy, Jim O’Gorman, Devon Kearns, and Mati Aharoni. *Metasploit: The Penetration Tester’s Guide*. No Starch Press, 2011.
- [11] Wayne Xin Zhao, Kun Zhou, Junyi Li, Tianyi Tang, Xiaolei Wang, Yupeng Hou, Yingqian Min, Beichen Zhang, Junjie Zhang, Zican Dong, et al. A survey of large language models. *arXiv preprint arXiv:2303.18223*, 1(2):1–124, 2023.
- [12] Mohammad Shoeybi, Mostofa Patwary, Raul Puri, Patrick LeGresley, Jared Casper, and Bryan Catanzaro. Megatron-lm: Training multi-billion parameter language models using model parallelism, 2020.
- [13] Anthropic. Model context protocol. <https://www.anthropic.com/news/model-context-protocol>, 2024. Accessed: 2026-03-18.
- [14] Mohammed Mehedi Hasan, Hao Li, Emad Fallahzadeh, Gopi Krishnan Rajbahadur, Bram Adams, and Ahmed E. Hassan. Model context protocol (mcp) at first glance: Studying the security and maintainability of mcp servers, 2025.

- [15] K Huang. Define ai agent. In Ken Huang, editor, *Agentic AI: Theories and Practices*. 2025.
- [16] K Huang and J Huang. Define ai agent. In Ken Huang, editor, *Agentic AI: Theories and Practices*. 2025.
- [17] J Huang, K Huang, K Jackson, and C Hughes. Ai agent safety and security considerations. In Ken Huang, editor, *Agentic AI: Theories and Practices*. 2025.
- [18] Risky Business Feature Podcast. ...mcp is dead. <https://risky.biz/>, 2026. Podcast episode , Accessed: 2026-03-18.
- [19] Ethan Michalak and Mark Perry. Connecting llms to caldera with the mcp plugin. <https://medium.com/@mitrecaldera/connecting-llms-to-caldera-with-the-mcp-plugin-dal450254827>, 2025. MITRE Caldera Blog, Accessed: 2026-03-16.
- [20] GH05TCREW. Metasploit mcp server. <https://github.com/GH05TCREW/MetasploitMCP>, 2025. Accessed: 2026-03-16.
- [21] Hare Sudhan. Atomic red team mcp server. <https://github.com/cyberbuff/atomic-red-team-mcp>, 2025. Accessed: 2026-03-16.
- [22] Ajay Bandi, Bhavani Kongari, Roshini Naguru, Sahitya Pasnoor, and Sri Vidya Vilipala. The rise of agentic ai: A review of definitions, frameworks, architectures, applications, evaluation metrics, and challenges. *Future Internet*, 17(9), 2025.
- [23] Anthropic. Model temperature, 2026. Accessed: 2026-05-16.
- [24] LangChain AI. Langchain: The agent engineering platform. <https://github.com/langchain-ai/langchain>, 2026. Accessed: 2026-05-02.
- [25] Xinyi Hou, Yanjie Zhao, Shenao Wang, and Haoyu Wang. Model context protocol (mcp): Landscape, security threats, and future research directions. *ACM Trans. Softw. Eng. Methodol.*, February 2026. Just Accepted.
- [26] Advisera. Iso 27001 risk assessment, treatment, and management: The complete guide, 2025. Accessed: 2026-03-04.
- [27] Simplr. General recommended vram guidelines for llms, 2025. Online; accessed 10 March 2026.

## List of Figures

|                                                                                                       |     |
|-------------------------------------------------------------------------------------------------------|-----|
| 5.1. EnvironmentLoader Base Class Refactoring Diff . . . . .                                          | 34  |
| 5.2. topology-test-environment . . . . .                                                              | 40  |
| 5.3. Authoritative Zone Transfer Snippet . . . . .                                                    | 41  |
| 6.1. Mean run duration per model (Containerlab, $n = 10$ runs each). . .                              | 73  |
| 6.2. Mean total token usage per model (Containerlab, $n = 10$ runs each). .                           | 74  |
| 6.3. Total tokens vs. run duration, annotated with tool-call count per<br>run (Containerlab). . . . . | 75  |
| 6.4. Mean run duration per model (Containerlab, $n = 10$ runs each). . .                              | 76  |
| 6.5. Mean total token usage per model (Containerlab, $n = 10$ runs each). .                           | 77  |
| 6.6. Total tokens vs. run duration, annotated with tool-call count per<br>run (Containerlab). . . . . | 78  |
| 6.7. Mean run duration per model (Containerlab, $n = 10$ runs each). . .                              | 79  |
| 6.8. Mean total token usage per model (Containerlab, $n = 10$ runs each). .                           | 80  |
| 6.9. Total tokens vs. run duration, annotated with tool-call count per<br>run (Containerlab). . . . . | 81  |
| A.1. Checklist and Burndown Sprint1 . . . . .                                                         | 101 |
| A.2. Sprint Retro 1 . . . . .                                                                         | 102 |
| A.3. Roadmap Sprint Planning 2 . . . . .                                                              | 104 |
| A.4. Milestone Description and Charts: Compare Frameworks . . . . .                                   | 105 |
| A.5. Milestone Description and Charts: Research AI-Integration . . . . .                              | 106 |
| A.6. Milestone Description and Charts: Configuration Management . . .                                 | 107 |
| A.7. Sprint Retro 2 . . . . .                                                                         | 108 |
| A.9. Milestone Description: Reproducible Test Environment . . . . .                                   | 109 |
| A.8. Roadmap Sprint Planning 3 . . . . .                                                              | 110 |
| A.10. Milestone Description and Charts: Research AI-Integration . . . . .                             | 111 |
| A.11. Retro Sprint 3 . . . . .                                                                        | 112 |
| A.12. Milestone Description: Reproducible Test Environment . . . . .                                  | 114 |
| A.13. Retro Sprint 4 . . . . .                                                                        | 115 |
| A.14. Roadmap Sprint Planning 5 . . . . .                                                             | 116 |
| A.15. Risks Brainstorming . . . . .                                                                   | 123 |
| A.16. Risks Brainstorming . . . . .                                                                   | 124 |
| A.17. Risk Assessment . . . . .                                                                       | 125 |
| A.18. Variant 1: AI vs Manually Built Scenarios . . . . .                                             | 132 |
| A.19. Variant 2: Multiple Blue and Red Team Scenarios . . . . .                                       | 133 |



|                                                                |     |
|----------------------------------------------------------------|-----|
| A.20. AI Integration Variant 1: AI-agent in NSAK . . . . .     | 134 |
| A.21. Variant 1, AI Agent in NSAK . . . . .                    | 134 |
| A.22. AI Integration Variant 2: NSAK as MCP server . . . . .   | 135 |
| A.23. Variant 2 NSAK MCP Server . . . . .                      | 135 |
| A.24. Brainstorm Scenarios . . . . .                           | 137 |
| A.25. Clustering possible Scenarios . . . . .                  | 138 |
| A.26. NSAK Scenario: Network Discovery - Red Team . . . . .    | 140 |
| A.27. NSAK Scenario: TLS Downgrade Poodle - Red Team . . . . . | 141 |
| A.28. NSAK Scenario: TLS Downgrade Poodle -Topology . . . . .  | 142 |
| A.29. NSAK Scenario: Honeypot - Blue Team . . . . .            | 143 |
| A.30. NSAK Scenario: Traffic Capture - Red Team . . . . .      | 144 |
| A.31. NSAK Scenario: Intrusion Detection - Blue Team . . . . . | 145 |
| A.32. NSAK Scenario: AI - Purple Team . . . . .                | 146 |
| A.33. FRITZ!Box DynDNS Setup . . . . .                         | 149 |
| A.34. Infomaniak DynDNS Setup . . . . .                        | 149 |
| A.35. FRITZ!Box DynDNS Setup . . . . .                         | 150 |
| A.36. FRITZ!Box DynDNS Setup . . . . .                         | 151 |

## List of Tables

|                                                                                         |     |
|-----------------------------------------------------------------------------------------|-----|
| 2.1. Comparison of Metasploit, Atomic Red Team, and MITRE Caldera ( [4, 7–9]) . . . . . | 11  |
| 5.1. Device Resource Classes . . . . .                                                  | 28  |
| 5.2. Resource Python Classes . . . . .                                                  | 28  |
| 5.3. Benchmark Suite Concepts . . . . .                                                 | 70  |
| 6.1. Comparison between Traditional Recon and AI Recon . . . . .                        | 71  |
| 6.2. Nonlinear Model Results . . . . .                                                  | 73  |
| 6.3. Nonlinear Model Results . . . . .                                                  | 75  |
| 6.4. Nonlinear Model Results . . . . .                                                  | 78  |
| A.1. Key Scrum Meetings . . . . .                                                       | 92  |
| A.2. Issue Board . . . . .                                                              | 93  |
| A.3. Definition of Ready Criteria for Tickets . . . . .                                 | 93  |
| A.4. Definition of Done Criteria for Tickets . . . . .                                  | 94  |
| A.5. Project Milestone Overview in GitLab . . . . .                                     | 95  |
| A.6. Sprint Planning 2 . . . . .                                                        | 103 |
| A.7. Sprint Planning 3 . . . . .                                                        | 109 |
| A.8. Sprint Planning 4 . . . . .                                                        | 113 |
| A.9. Sprint Planning 5 . . . . .                                                        | 116 |
| A.10. Risk categories used for the project risk assessment . . . . .                    | 121 |
| A.11. Risk evaluation criteria . . . . .                                                | 121 |
| A.12. Risk probability scale . . . . .                                                  | 122 |
| A.13. Risk impact scale . . . . .                                                       | 122 |
| A.14. Risk level classification (probability and impact are interchangeable) . . . . .  | 122 |
| A.15. Risk Treatment Strategies . . . . .                                               | 124 |
| A.16. Meeting Notes: Checkpoint 1 . . . . .                                             | 126 |
| A.17. Meeting Notes: Checkpoint 2 . . . . .                                             | 127 |
| A.18. Meeting Notes: Checkpoint 3 . . . . .                                             | 128 |
| A.19. Meeting Notes: Checkpoint 4 . . . . .                                             | 129 |
| A.20. Meeting Notes: Checkpoint 5 . . . . .                                             | 130 |
| A.21. Meeting Notes: Checkpoint 6 . . . . .                                             | 130 |
| A.22. Hardware configuration of the self-hosted LLM inference server . . . . .          | 147 |

## List of Listings

|     |                                                                   |     |
|-----|-------------------------------------------------------------------|-----|
| 1.  | <a href="#">NSAK Static Configuration</a>                         | 23  |
| 2.  | <a href="#">NSAK Runtime Configuration</a>                        | 25  |
| 3.  | <a href="#">NSAK CLI Initialisation</a>                           | 26  |
| 4.  | <a href="#">Persisted Runtime Configuration (run/config.yaml)</a> | 27  |
| 5.  | <a href="#">Device Resource YAML Specification</a>                | 28  |
| 6.  | <a href="#">NSAK CLI: List Devices</a>                            | 29  |
| 7.  | <a href="#">Device Resource Configuration YAML Specification</a>  | 29  |
| 8.  | <a href="#">DeviceLoader Configuration Parsing</a>                | 31  |
| 9.  | <a href="#">Device Configuration Datastructures</a>               | 32  |
| 10. | <a href="#">Device Configuration CLI Management</a>               | 33  |
| 11. | <a href="#">Mergeable Resource YAML Specification</a>             | 35  |
| 12. | <a href="#">Updated Interface Specification</a>                   | 35  |
| 13. | <a href="#">discover_hosts Drill YAML</a>                         | 36  |
| 14. | <a href="#">mitm Scenario YAML</a>                                | 37  |
| 15. | <a href="#">Excerpt of the Containerlab topology definition</a>   | 39  |
| 16. | <a href="#">AI Configuration Dataclass</a>                        | 44  |
| 17. | <a href="#">AI Provider Map and MCP Client Initialisation</a>     | 45  |
| 18. | <a href="#">LangChain: AI Agent</a>                               | 47  |
| 19. | <a href="#">create_ai_agent() Factory Function</a>                | 50  |
| 20. | <a href="#">LangChain Tool: Host Configuration</a>                | 52  |
| 21. | <a href="#">LangChain Tool: CLI Tool</a>                          | 54  |
| 22. | <a href="#">LangChain Tool: Human Interaction Hook</a>            | 55  |
| 23. | <a href="#">Loki Logging Configuration</a>                        | 56  |
| 24. | <a href="#">E-Mail Alerting Configuration</a>                     | 57  |
| 25. | <a href="#">E-Mail Alerting Implementation</a>                    | 57  |
| 26. | <a href="#">LangChain Tool: E-Mail Alerting</a>                   | 58  |
| 27. | <a href="#">Scenario: AI Agent - scenario.py</a>                  | 60  |
| 28. | <a href="#">Scenario: AI Agent - scenario.yaml</a>                | 61  |
| 29. | <a href="#">Scenario: AI Reconnaissance - scenario.py</a>         | 63  |
| 30. | <a href="#">Scenario: AI Reconnaissance - scenario.yaml</a>       | 64  |
| 31. | <a href="#">Scenario: AI Intrusion Detection - scenario.py</a>    | 66  |
| 32. | <a href="#">Scenario: AI Intrusion Detection - scenario.yaml</a>  | 67  |
| 33. | <a href="#">UFW configuration</a>                                 | 148 |
| 34. | <a href="#">Restart SWAG Container</a>                            | 150 |
| 35. | <a href="#">Self-Hosted Base Setup</a>                            | 152 |

|     |                                                                              |     |
|-----|------------------------------------------------------------------------------|-----|
| 36. | <a href="#">Nginx &amp; Let's Encrypt Instructions</a>                       | 153 |
| 37. | <a href="#">Nginx &amp; Let's Encrypt Compose File</a>                       | 154 |
| 38. | <a href="#">Ollama Instructions</a>                                          | 155 |
| 39. | <a href="#">Ollama Nginx Configuration</a>                                   | 156 |
| 40. | <a href="#">Ollama Enable Debug Log</a>                                      | 157 |
| 41. | <a href="#">Open WebUI Instructions</a>                                      | 158 |
| 42. | <a href="#">Open WebUI Compose File</a>                                      | 159 |
| 43. | <a href="#">Open WebUI Nginx Configuration</a>                               | 160 |
| 44. | <a href="#">Drawio Instructions</a>                                          | 161 |
| 45. | <a href="#">Drawio Compose File</a>                                          | 162 |
| 46. | <a href="#">Drawio Nginx Configuration</a>                                   | 163 |
| 47. | <a href="#">Netbox Instructions</a>                                          | 164 |
| 48. | <a href="#">Netbox Compose File</a>                                          | 165 |
| 49. | <a href="#">Netbox Nginx Configuration</a>                                   | 166 |
| 50. | <a href="#">Grafana / Loki Instructions</a>                                  | 167 |
| 51. | <a href="#">Grafana / Loki Compose File</a>                                  | 168 |
| 52. | <a href="#">Grafana / Loki Nginx Configuration</a>                           | 169 |
| 53. | <a href="#">Grafana / Loki Configuration</a>                                 | 170 |
| 54. | <a href="#">Grafana / Loki Alloy Configuration</a>                           | 171 |
| 55. | <a href="#">Firewall</a>                                                     | 173 |
| 56. | <a href="#">Web-Server and DNS</a>                                           | 174 |
| 57. | <a href="#">File Server and Domain Control</a>                               | 175 |
| 58. | <a href="#">Work-Stations Bob and Alice</a>                                  | 176 |
| 59. | <a href="#">Printer</a>                                                      | 177 |
| 60. | <a href="#">Links</a>                                                        | 177 |
| 61. | <a href="#">Introduced Vulnerabilities in the Virtual Enterprise Network</a> | 179 |
| 62. | <a href="#">Reconnaissance Scenario Implementation</a>                       | 181 |
| 63. | <a href="#">Discover Hosts Drill: Resource Configuration</a>                 | 182 |
| 64. | <a href="#">Discover Hosts Drill: Implementation</a>                         | 183 |
| 65. | <a href="#">Port Scan Drill: Implementation</a>                              | 184 |
| 66. | <a href="#">Enumerate Services Drill: Implementation</a>                     | 185 |

## Glossary

This document is incomplete. The external file associated with the glossary ‘main’ (which should be called `thesis.gls`) hasn’t been created.

Check the contents of the file `thesis.gls`. If it’s empty, that means you haven’t indexed any of your entries in this glossary (using commands like `\gls` or `\glsadd`) so this list can’t be generated. If the file isn’t empty, the document build process hasn’t been completed.

If you don’t want this glossary, add `nomain` to your package option list when you load `glossaries-extra.sty`. For example:

```
\usepackage[nomain]{glossaries-extra}
```

Try one of the following:

- ▶ Add `automake` to your package option list when you load `glossaries-extra.sty`. For example:

```
\usepackage[automake]{glossaries-extra}
```

- ▶ Run the external (Lua) application:

```
makeglossaries-lite.lua "thesis"
```

- ▶ Run the external (Perl) application:

```
makeglossaries "thesis"
```

Then rerun  $\text{\LaTeX}$  on this document.

This message will be removed once the problem has been fixed.

## A. Project Management

### A.1. Stakeholders

Hj. Wenger: Instructor and Network Specialist

Urs Keller: Expert and Entrepreneur

### A.2. Agile Management Process

The project followed an interactive development approach using Gitlab for issue tracking and planning. We defined our project management in a Scrum-like manner, but we individualized the process according to the projects scope and our needs. Work items where organized into epics, milestones and issues. Each iteration lasts two weeks. We decided to hold the following Scrum ceremonies after each sprint:

| <b>Sprint Planning</b>             | <b>Sprint Review</b>        | <b>Sprint Retro</b>     | <b>Checkpoint</b>        |
|------------------------------------|-----------------------------|-------------------------|--------------------------|
| Select milestone with sprint goal  | Close open tickets          | Reflect on sprint       | Present sprint results   |
| Check issues for DoR and obstacles | Gather team feedback        | Identify improvements   | Get stakeholder feedback |
| Ensure issues are clear            | Validate delivered features | Discuss what went well  | Define next tasks        |
| Create sprint backlog              | Review progress             | Discuss issues/problems | Discuss issues/problems  |

Table A.1.: Key Scrum Meetings

The workflow used for ticket tracking consisted of the following stages:

## A.2.1. Issue Lifecycle:

| Buckets | Description                                                           |
|---------|-----------------------------------------------------------------------|
| Backlog | List of all tasks and features planned for the project.               |
| Todo    | Tasks selected for the current sprint that are ready to be worked on. |
| Q&A     | Tasks waiting for clarification, review, or testing.                  |
| Done    | Completed tasks that meet the defined acceptance criteria.            |

Table A.2.: Issue Board

## A.2.2. Definition of Ready (DoR)

| Category   | Criteria                                                                                                                                                                     |
|------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Clarity    | <ul style="list-style-type: none"><li>▶ Clear and precise title</li><li>▶ Objective or problem described in two to three sentences</li><li>▶ Scope clearly defined</li></ul> |
| Scope      | <ul style="list-style-type: none"><li>▶ Affected module specified (drill, scenario, core, container, frontend)</li><li>▶ External dependencies identified</li></ul>          |
| Validation | <ul style="list-style-type: none"><li>▶ Testable acceptance criteria defined</li><li>▶ Expected behavior clearly specified</li></ul>                                         |

Table A.3.: Definition of Ready Criteria for Tickets

### A.2.3. Definition of Done (DoD)

| Category       | Criteria                                                                                                                                                                                                                                                        |
|----------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Implementation | <ul style="list-style-type: none"><li>▶ All acceptance criteria fulfilled</li><li>▶ Container execution verified</li><li>▶ Pre-commit hooks pass</li><li>▶ Logging and error handling consistent</li></ul>                                                      |
| Testing        | <ul style="list-style-type: none"><li>▶ Feature or bug fix tested</li><li>▶ Acceptance criteria verifiable (automated or manual)</li><li>▶ Relevant edge cases considered</li><li>▶ Execution reproducible</li></ul>                                            |
| Documentation  | <ul style="list-style-type: none"><li>▶ README updated if required</li><li>▶ Architecture changes documented</li><li>▶ Thesis relevance briefly noted</li><li>▶ Limitations documented</li><li>▶ Required documentation stored for thesis or appendix</li></ul> |

Table A.4.: Definition of Done Criteria for Tickets

## A.3. Project Structure

This section describes the projects structure and list the goals as milestones linked to their respective page in GitLab.

### A.3.1. Milestones

A milestone represent crucial steps toward the overall project goal. Per our definition a milestone has at least one epic and the mandatory milestones are listed in [Table A.5](#)



### GitLab Milestones

Project Management  
 Research AI Integration  
 Research: Compare Frameworks  
 NSAK Framework: Configuration Management  
 NSAK Scenario: Network Discovery - Red Team  
 NSAK Scenario: AI Scenarios  
 Reproducible Test Environment

Table A.5.: Project Milestone Overview in GitLab

### Project Planning

19.02.2026 – 10.03.2026

This is the initial milestone which is used to plan the whole project and bachelor thesis.

- ✓ The bachelor thesis is planned as an agile (scrum-like) project
- ✓ A risk analysis for the project was conducted and the identified risks are assessed and addressed
- ✓ All meetings, deadlines, scrum ceremonies and sprints are added to the Outlook calendar
- ✓ The goals of the thesis are defined and categorized as “must”, “should” and “could”
- ✓ The goals are created as milestones and divided into epics, categorized by the labels: Project Management, Research, Framework, Implementations, Documentation
- ✓ We know how we coordinate the project with the stakeholders (instructor and expert)

### Research: AI Integration

10.03.2026 – 22.03.2026

The goal of this research milestone was to investigate how AI can support automated security testing and adversary emulation. The research evaluated potential integration points where AI could enhance scenario execution, automation,

decision-making, and analysis capabilities of the NSAK framework. The milestone was extended by three days following input from the expert and instructor to clarify the intended interaction between LLMs and NSAK.

- ☐ How is AI currently used in cybersecurity automation?
- ☐ Which tasks in adversary emulation can be supported or automated by AI?
- ☐ How could AI interact with the NSAK scenarios and drills (e.g. Agents, MCP)?
- ☐ What architectural changes would be required to integrate AI into NSAK?

### Research: Compare Frameworks

*10.03.2026 – 19.03.2026*

This milestone focused on the research and comparison of existing security automation frameworks that could potentially support or influence the NSAK scenario and drill execution model. The evaluation included frameworks such as MITRE CALDERA and Atomic Red Team, examining their execution models, configuration and parameter handling, extensibility mechanisms, and the structure of reusable test definitions.

The expected outcomes were:

- ☐ A short comparison of selected frameworks
- ☐ A summary of relevant architectural concepts
- ☐ An initial assessment of how these frameworks relate to NSAK

### NSAK Framework: Configuration Management

*10.03.2026 – 19.03.2026*

This milestone defined the configuration model and execution framework for NSAK scenarios. It introduced a structured approach for configuring scenarios and drills, including typed argument declarations, and the generation of a resolved run configuration. The goal was to provide a flexible and reproducible configuration system that allows scenarios and drills to expose parameters and execute in a defined order.

- ✓ Drills and scenarios can expose runtime parameters as configuration objects
- ✓ The configuration objects can be set via CLI arguments
- ✓ A configuration management system is implemented in the NSAK framework

## NSAK Scenario: Network Discovery – Red Team

02.04.2026 – 16.04.2026

Network discovery is one of the most essential Red Team activities in the reconnaissance phase. This milestone covered the implementation of a scenario that identifies network participants and services, producing a structured network map that can be consumed by subsequent drills such as Address Resolution Protocol spoofing.

### Goals (Must):

- ☐ The scenario identifies the available ports and interfaces on the device
- ☐ The scenario discovers the available subnets and obtains valid addresses
- ☐ The scenario discovers addresses, addresses and services (port scan)
- ☐ The scenario supports at least a fast full-search mode
- ☐ The scenario returns a network map as a structured data type

### Goals (Could):

- ☐ A human interaction hook allows the operator to select the discovery mode
- ☐ Additional modes: slow full search, distinct search

## NSAK Scenario: AI Scenarios

19.03.2026 – 16.04.2026

This milestone explored what happens when an AI agent is given full access to the NSAK framework for both Red and Blue Team activities. The scenario is intentionally open-ended, reflecting the research nature of the milestone.

### Goals (Must):

- ☐ The operator can provide a custom prompt or select a predefined Red or Blue Team prompt
- ☐ The scenario runs an AI agent with a remote connection
- ☐ An operator can interact with the agent via CLI at all times
- ☐ The scenario logs all relevant actions and network traffic locally
- ☐ The AI can trigger an e-mail alert to notify an operator
- ☐ The scenario has a kill switch to stop execution at any time

### Goals (Should):

- ☐ Human interaction hook via Web GUI
- ☐ Remote logging and sniffing
- ☐ Kill switch with full device shutoff

### Reproducible Test Environment

19.03.2026 – 02.04.2026

The goal of this milestone was to design and implement a reproducible and safe test environment for the NSAK framework, providing a stable basis for consistent execution and evaluation of both manual and AI-assisted scenarios.

- ☐ The test environment can be set up reproducibly using Infrastructure as Code (e.g. Ansible).
- ☐ NSAK scenarios and drills can be executed consistently within the environment.
- ☐ The environment supports both manual and AI-assisted scenarios.
- ☐ Scenario results can be collected and analyzed for evaluation purposes.

### A.3.2. Epics

consist of multiple stories and contribute to fulfilling milestones.

### A.3.3. Issues

Each Story is assigned a weight, sprint, milestone and epic and are usually represented by an “Issue” in gitlab.

## A.4. Sprints / Iterations

### A.4.1. Planning

“Task” and “Issue”, the terms are often used interchangeably with “Stories”. Stories must meet the [A.2.2](#) before implementation and [A.2.3](#) before closing.

## A.5. Planning and Estimation

The following estimation strategy was used:

- ▶ 1 hour corresponds roughly to 1 weight point
- ▶ Iteration length: 2 weeks
- ▶ Weights per iteration: 40h per person = 80 weight per iteration
- ▶ Sprint planning duration: 1.5 hours per iteration
- ▶ Sprint review duration: 1 h per iteration
- ▶ Sprint retro duration: 1 h per iteration

## A.6. Project Roadmap

To keep track of the overall progress in the project, the project team uses the “Roadmap” feature in GitLab. After each sprint planning meeting we create a new screenshot of the roadmap and add it to the according sprint section in [A.7](#). The first sprint planning is an exception, because at this point the diagram would have been empty.

## A.7. Sprint Ceremonies

### A.7.1. Sprint 1: 19.02 –10.03.2026

#### Sprint Planning 1: 19.02.2026

For the planning of sprint 1 we derived all issues from the milestone project management [A.5](#) and had an overall weight from 89.

#### Sprint Review 1: 10.03.2026

The milestone was successfully reached, and the board was cleared in preparation for the next sprint. All issues were reviewed and moved to the “Done” column.

**Milestone Review: Project Management** As shown in burndown and burnup charts [A.1](#) of the Project Planning [A.3.1](#) milestone, all deliverables were fulfilled. However, the deadlines of the milestone had to be moved to the 10.03.2026, because the project planning required more time than initially expected. It should also be noted that the charts abruptly go down or up respectively, as we closed a lot of the tickets after we reviewed them together just before we held this meeting.

## Project Planning

- ☒ The bachelor thesis is planned as an agile (scrum like) project
- ☒ A risk analysis for the project was conducted
- ☒ The identified risks are assessed and addressed
- ☒ All meetings, deadlines, scrum ceremonies and sprints are added to the outlook calendar
- ☒ The goals of the thesis are defined and categorized as "must", "should" and "can"
- ☒ The goals are created as milestones and divided into epics, which are categorized by the labels: "Project Management", "Research", "Framework", "Implementations", "Documentation"
- ☒ We know how we coordinate the project with the stake holders (instructor and expert)

Display by Count Weight

### Burndown chart



### Burnup chart




Figure A.1.: Checklist and Burndown Sprint1

## Sprint Retro 1: 10.03.2026

This sprint retro was held in Microsoft Whiteboard as it provides simple templates, the result is illustrated in [A.2](#).

### Key Takeaways:

- ▶ Maintain clear and consistent communication with the instructor.
- ▶ Be aware of maintaining a healthy work-life balance during the project.
- ▶ Apply lean project management practices to keep the process efficient and focused.




Figure A.2.: Sprint Retro 1



## A.7.2. Sprint 2: 05.03 –18.03.2026

## Sprint Planning 2: 05.03.2026

The current state of the project roadmap is shown in the following screenshot [A.3](#).

| Week   | Responsible | Task                                                      |
|--------|-------------|-----------------------------------------------------------|
| Week 1 | Lukas       | Configuration management implementation                   |
| Week 1 | Frank       | Research and comparison of frameworks                     |
| Week 1 | Team        | AI workshop on Thursday and writing of the related issues |
| Week 2 | Team        | AI research and evaluation                                |
| Week 2 | Team        | Preparation and pitch of the project deliverables         |

Table A.6.: Sprint Planning 2

## A. Project Management




Figure A.3.: Roadmap Sprint Planning 2

### Sprint Review 2: 19.03.2026

We got some helpful inputs related to both milestones from the tutor and expert in the checkpoint meeting [A.8.3](#) earlier this day. We decided to still consider the milestone to be successfully reached as the planned work package and deliverables were finalized and presented. The resulting follow-up tasks are planned for

the next week and as a consequence we will hold the next sprint planning [A.7.3](#) a week later than planned. All issues were reviewed and moved to the “Done” column in preparation for the next sprint. It must be noted, that because we usually review and close the tickets at the end of the sprint, the charts are not representative in most cases.

Milestone Review: Compare Frameworks    The Compare Frameworks [A.3.1](#) milestone is completed as illustrated in [A.4](#).

Research: Compare Frameworks (must)

Description

This milestone focuses on the research and comparison of existing security automation frameworks that could potentially support or influence the NSAK scenario and drill execution model.

The objective of this research is to understand how established frameworks structure and execute security tests, manage configuration, and organize reusable components.

The evaluation will include frameworks such as MITRE CALDERA and Atomic Red Team, which represent different approaches to security automation.

List of pre-work/papers:

The comparison will focus on key aspects relevant for NSAK, including:

- execution models (scenarios, abilities, tests)
- configuration and parameter handling
- extensibility and plugin mechanisms
- structure and reuse of test definitions

The results of this research will provide the foundation for a framework gap analysis, where the capabilities of the evaluated frameworks are compared against the NSAK requirements.

The outcome of this milestone will be:

- ☒ a short comparison of selected frameworks
- ☒ a summary of relevant architectural concepts
- ☒ an initial assessment of how these frameworks relate to NSAK

The findings will serve as input for the subsequent gap analysis milestone.

Display by

Count

Weight

Burndown chart

| Date        | Remaining | Guideline |
|-------------|-----------|-----------|
| 10 Mar 2026 | 30        | 30        |
| 12 Mar 2026 | 25        | 25        |
| 14 Mar 2026 | 25        | 20        |
| 16 Mar 2026 | 25        | 15        |
| 18 Mar 2026 | 0         | 0         |

Burnup chart

| Date        | Total | Completed |
|-------------|-------|-----------|
| 10 Mar 2026 | 30    | 0         |
| 12 Mar 2026 | 30    | 0         |
| 14 Mar 2026 | 30    | 0         |
| 16 Mar 2026 | 30    | 0         |
| 18 Mar 2026 | 30    | 30        |

Figure A.4.: Milestone Description and Charts: Compare Frameworks

Milestone Review: Research AI-Integration    As shown in [A.5](#), the milestone Research AI-Integration [A.3.1](#) is completed.

## Research AI-Integration

AI Integration (must) The goal of this research is to investigate how AI (artificial intelligence) can support automated security testing and adversary emulation. The research evaluates potential integration points where AI could enhance the scenario execution, automation, decision-making, and analysis capabilities of the NSAK framework.


Success criteria:

- ☒ How is AI currently used in cybersecurity automation?
- ☒ Which tasks in adversary emulation can be supported or automated by AI?
- ☒ How could AI interact with the NSAK scenarios and drills? (e.g. Agents, MCP)
- ☒ What architectural changes would be required to integrate AI into NSAK?

Milestone stretched to 22.03 due to input from the expert and instructor to clarify the interaction with LLM and NSAK.

Display by Count Weight

### Burndown chart



### Burnup chart




Figure A.5.: Milestone Description and Charts: Research AI-Integration

**Milestone Review: Configuration Management** The burnup and burndown charts below [A.6](#) are showing that the milestone Configuration Management [A.3.1](#) is implemented as planned.

## NSAK Framework: Configuration Management (must)

- ✓ A configuration management is implemented into the NSAK Framework

### Setup

Current configuration approach: ENV vars, CLI flags or hardcoded New configuration approach:


- ✓ Drills and scenarios can expose build and runtime parameters as a configuration objects
- ✓ The configuration objects can be set interactively during `nsak scenario build` or `nsak scenario run`
- ✓ Alternatively the configuration objects can be set as CLI parameters (e.g. JSON)
- ✓ Alternatively the configuration objects can be provided via file (e.g. JSON)

This milestone defines the configuration model and execution framework for NSAK scenarios. It introduces a structured approach for configuring scenarios and drills, including drill definitions, reusable presets, and the generation of a resolved run configuration.

The goal is to provide a flexible and reproducible configuration system that allows scenarios to expose parameters, merge configuration layers, and execute drills in a defined order.

Display by ☐ Count ☒ Weight

### Burndown chart



### Burnup chart




Figure A.6.: Milestone Description and Charts: Configuration Management

## Sprint Retro 2: 19.03.2026

This sprint retro was held in Microsoft Whiteboard as it provides simple templates, the result is illustrated in A.7.

### Key Takeaways:

- ▶ The long project planning phase greatly helped us with the projects overview and minimizes the planning needed for each sprint.
- ▶ The “focus weeks” really boosted the projects progress, and we are looking forward to the next “focus weeks”.
- ▶ Because we did not review tasks (Quality Assurance (QA)) on a regular basis, all this work needed to be done at the end of an iteration.
- ▶ Reviewing on a more regular basis would also lead to more accurate burnup and burndown charts.
- ▶ We decided to not include the charts in the future, because it makes no sense to adjust our process only for the sake of good-looking charts.

## 4L's retrospective

Use this template to conduct a retrospective on your team's last sprint. Fill out each of the quadrants while considering the following questions: What did your team like? What did your team lack? What did the team learn? What does the team wish for?



Strategy

| LIKED                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      | LEARNED                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <p><b>Focus Project Time</b></p> <p><b>Workshop and communication</b></p> <p><b>Focus on key deliverables pragmatic</b></p> <p><b>Constructively meeting with expert and instructor with shared thesis</b></p> <p><b>Agenda und Result</b></p> <p><b>Two weeks of focus helped to boost progress and really get into the topics</b></p> <p><b>Output of the Research Milestones (AI and Framework Comparison)</b></p> <p><b>Planning helped splitting tasks and responsibilities efficiently Review takes time too</b></p> | <p><b>Review fluently to improve charts and processes and improve backward tracability</b></p> <p><b>Show only graphics with clear intension</b></p> <p>maybe prepare links in the chapter to jump to the right section for sharring inside the thesis</p> <p><b>Plan less tasks in a sprint, its always more work than we think</b></p> <p><b>Review takes time too</b></p> <p><b>Make diagrams more simple (stakeholders may not have the same context)</b></p> |
| LACKED                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     | LONGED FOR                                                                                                                                                                                                                                                                                                                                                                                                                                                        |
| <p><b>Life around the project needed more time as expected</b></p> <p><b>constantly of reviewing the assigned tasks</b></p> <p><b>actual view of roadmap needs to be displayed in pm part of thesis</b></p> <p><b>Time for documentation</b></p> <p><b>Syncing understanding of the AI tech stack with stakeholders</b></p>                                                                                                                                                                                                | <p><b>More proficience with scientific writing</b></p> <p><b>citation help</b></p> <p><b>Faster implementation of concepts, code and documentation</b></p>                                                                                                                                                                                                                                                                                                        |

Figure A.7.: Sprint Retro 2

## A.7.3. Sprint 3: 26.03 –09.04.2026

## Sprint Planning 3: 26.03.2026

Because we got some important input regarding variant decision for AI-integration, framework comparison, and configuration management [A.18](#) which led to additional work. As consequence of this, we decided to shorten the third sprint to one week, which is reflected in the following table [A.7](#) and postponed the sprint planning from the 19.03 to the 26.03. To account for the delayed start in the overall planning we adjusted the roadmap as shown in the screenshot below [A.8](#).

| Week       | Responsible | Task                                                              |
|------------|-------------|-------------------------------------------------------------------|
| Week 1 & 2 | Frank       | Reproducible Test Environment implementation                      |
| Week 1 & 2 | Lukas       | NSAK Scenario: AI Scenarios (first part): AI-Agent implementation |

Table A.7.: Sprint Planning 3

## Sprint Review 3: 09.04.2026

The completion of the following two milestones in this sprint are marking a very important point in our bachelor thesis, as we are now able to run the AI based reconnaissance scenario against the reproducible test environment. This enables us to reproducibly test and analyze the actual capabilities of the AI-agent in a realistic scenario. Further we can now demonstrate the full interplay of the technologies to the projects stakeholders and possibly other interested persons.

**Milestone Review: Reproducible Test Environment** The success criteria of the Reproducible Test Environment [A.3.1](#) milestone are completed as illustrated in [A.9](#).

**Reproducible Test Environment (must)**

Reproducible Test Environment (must) The goal is to design and implement a reproducible and safe test environment for the NSAK framework that enables consistent execution and evaluation of scenarios and drills. The environment should provide a stable basis for testing and comparison of manual and AI-assisted scenarios.

Success criteria:

- ☒ The test environment can be set up reproducibly using a defined configuration or setup procedure with Infrastructure as code (e.g., ContainerLab)
- ☒ NSAK scenarios and drills can be executed consistently within the environment.
- ☒ The environment supports both manual and AI-assisted scenarios.
- ☒ Scenario results can be collected and analyzed for evaluation purposes (Logging via Loki).

short description about the results of the milestone

Figure A.9.: Milestone Description: Reproducible Test Environment

**Milestone Review: Implement AI Scenarios (first part): AI-Agent implementation**  
As shown in [A.10](#), all “must” goals of the milestone Implement AI Scenarios [A.3.1](#)




Figure A.8.: Roadmap Sprint Planning 3



where completed. But there is still some work todo, to document the implementation and the learnings we made during the development. Further we could improve the implementation in several aspects such as tool-calling and external logging, where we can better analyze the AI-agents behavior. Because this was expected, we planned two sprints for the realization of the AI Scenarios.

### NSAK Scenario: AI Scenarios (must)

#### Description

The usage of AI in network security is still subject of recent research and with the advent of agentic AI we asked ourself: What happens and what are the possibilities when we give an AI full access to the NSAK framework. This scenario is not very well defined yet.

#### Goals

- ☒ The operator can provide an initial prompt or use a predefined Red or Blue Team prompt
- ☒ The scenario runs an AI agent with a remote connection
- ☒ An operator can interact with the agent, ideally at all times
- ☒ The scenario logs all relevant actions and network traffic
- ☒ The AI is able to trigger an alert to notify an operator depending on the prompt
- ☒ The scenario has a kill switch, so that the operator can stop it at all times

#### Tasks / Drills

Backbone:

- Initial Prompt
- Run
- Human Interaction Hook
- Logging / Sniffing
- Alerting
- Kill Switch

Must:

- ☒ Initial Prompt: Custom Prompt
- ☒ Initial Prompt: Red Team Prompt
- ☒ Initial Prompt: Blue Team Prompt
- ☒ Run: AI Agent with remote connection
- ☒ Human Interaction Hook: CLI
- ☒ Logging / Sniffing: Local
- ☒ Alerting: E-Mail
- ☒ Kill Switch: Stop Scenario

Should:

- ☐ Human Interaction Hook: Web GUI
- ☐ Logging / Sniffing: Remote
- ☐ Kill Switch: Shutoff

Could:

- ☐ Human Interaction Hook: Messenger (WhatsApp, Threema, etc.)
- ☐ Human Interaction Hook: Chats (Slack, RocketChat, etc.)
- ☐ Alerting: SMS
- ☐ Alerting: Phone
- ☐ Alerting: Pager
- ☐ Alerting: Messenger (WhatsApp, Threema, etc.)
- ☐ Alerting: Chats (Slack, RocketChat, etc.)
- ☐ Alerting: Scheduling
- ☐ Logging / Sniffing: Dashboard

Workshop Ticket:

#### Environment

This scenario would be really special, as it should be able to operate in all environments. For running an AI Agent we would need additional setup steps.

Figure A.10.: Milestone Description and Charts: Research AI-Integration

### Sprint Retro 3: 09.04.2026

The Key Take Away of [A.11](#) are to estimate the milestones more realistically. We spend a lot of time with documentation which hinders the development. The




Figure A.11.: Retro Sprint 3

upside of our continuous approach is a natural growing thesis which helps in the long run. Overall, we are quite happy with the outcome of sprint 3.

## A.7.4. Sprint 4: 10.04 –02.05.2026

## Sprint Planning 4: 09.04.2026

| Week       | Responsible   | Task                                                                              |
|------------|---------------|-----------------------------------------------------------------------------------|
| Week 1     | Frank & Lukas | Bring together and test reproducible test environment and AI Reconnaissance       |
| Week 2 & 3 | Frank         | Network Discovery - Red Team                                                      |
| Week 2     | Lukas         | NSAK Scenario: AI Scenarios (second part): Improve tool calling and documentation |
| Week 3     | Lukas         | External Logging with Grafana / Loki                                              |

Table A.8.: Sprint Planning 4

No photograph of the sprint planning roadmap was taken during this sprint, so no visual reference is available for this planning session.

## Sprint Review 4: 02.05.2026

In this sprint could finally bring together the reproducible network environment and the AI scenario, we realized the last scenario categorized as “Must” goal, extended the logging capabilities and could refine the AI agent. Now we can focus on refining the documentation and finalizing the last preparations before moving to the last phase of our thesis, the evaluation and discussion.

**Test run of the AI scenario in the reproducible network environment** After finishing the previous sprint we were excited to test the AI scenario in the reproducible network environment, as it’s marking a big milestone of our bachelor thesis. We spent more or less a whole day using different models and make fine adjustments to the setup, to see what’s possible. In the end we could instruct the AI agent to conduct network discovery, write a report in markdown, create a network diagram of the findings and send the artifacts via E-Mail. The results of this test run were added to the implementation section of the thesis [5.3.8](#).

**Milestone Review: Network Discovery - Red Team** In [A.12](#), not all “must” goals of the milestone were completed. During planning, a human interaction mode was discussed in which an operator could choose how fast or noisy the reconnaissance should be executed. This feature was not implemented, as it does not contribute to answering the thesis question. The scenario is able to discover MAC and IP addresses using arp-scan on Layer 2, and protocols and services using

nmap within the local subnet. All results are stored in a `NetworkDiscoveryResultMap` and displayed to the operator. Service enumeration using targeted NSE scripts was implemented as a dedicated drill, which runs service-specific scripts such as `http-title`, `dns-zone-transfer`, and `smb-security-mode` on all discovered ports.

The scenario could be extended with an optional static IP address, which would be useful in networks where DHCP is not available. Host discovery in remote networks is currently limited to IP addresses only, as Media Access Control addresses are not available across Layer 3 boundaries.

### NSAK Scenario: Network Discovery - Red Team (must)

#### Description

Network discovery is one of the most essential Red Team activity in the reconnaissance phase, because participants and services are identified and interesting targets can be evaluated. The resulting network map can then be used in later phases, for example an ARP spoofing attack can be executed to intercept traffic more efficiently.

This scenario covers the following bullet point from the thesis proposal: "Unterstützung typischer Vorgehensweisen aus dem Red-Team- und Penetration-Testing-Bereich, etwa durch Scan- oder Angriffssimulationen"

#### Goals

- ☒ The scenario correctly identifies the setup it can use for network discovery (ports and interfaces)
- ☒ The scenario is able to discover the available subnets
- ☒ The scenario can get valid IP addresses for the discovered subnets
- ☐ An operator can intervene to select a mode for how the network discovery should be done (optional)
- ☐ The scenario provides different modes on how the network discovery is done (selective, slow, fast)(optional)
- ☒ The scenario is able to identify clients, servers and services participating in the networks
- ☒ The scenario returns a map of the scanned networks

#### Tasks / Drills

Backbone:

- Physical Ports / Interfaces
- Subnets
- IP Address Assignment
- Human Interaction Hook
- Mode
- Devices
- Services
- Mapping

Must:

- ☒ Physical Ports / Interfaces: List ports
- ☒ Physical Ports / Interfaces: List interfaces
- ☒ Physical Ports / Interfaces: Filter interfaces
- ☒ Subnets: Scan subnets
- ☒ IP Address Assignment: DHCP - Request dynamic IP
- ☒ Mode: Fast full search
- ☒ Devices: Discover MAC Addresses
- ☒ Devices: Discover IP Addresses
- ☒ Services: Port Scan
- ☒ Mapping: Network map as datastructure
- ☐ Subnets: Filter subnets

Should: None

Could:

- ☐ IP Address Assignment: Static - Assign static IP
- ☐ Human Interface Hook: Select mode
- ☐ Mode: Slow full search
- ☐ Mode: Distinct search

Ticket: `swiss-army-knife-network-sniffer#122` (closed)

#### Environment

This scenario should work in all environments by design, as it should be able to discover as much as possible about an unknown network.

Maybe move the content to the implementation section and reference it

Figure A.12.: Milestone Description: Reproducible Test Environment



Figure A.13.: Retro Sprint 4

**Milestone Review: Implement AI Scenarios (second part): Improve tool calling and documentation** As already shown in A.10 from the previous sprint review A.7.3, all “must” goals of the milestone Implement AI Scenarios A.3.1 where completed. In this second part the AI scenarios where tested, tool calling was improved and the documentation was added to the implementation section of the thesis 5.3. Further we could already improve the AI scenarios with the findings of the extensive test run described above A.7.4.

**External Logging with Grafana / Loki** In the fifth checkpoint meeting A.8.5 which as in this sprint, we received the feedback that we should improve the logging for the AI-agent and especially which tools it is calling. We decided to set up a self-hosted logging system with Grafana and Loki, where we can collect all logs of the runs and create custom queries to make it easier to analyze the behavior of the AI-agent. The usage is described in the implementation section of the AI Scenarios 5.3.4. The self-hosted setup is documented in the addendum A.12.10.

Sprint Retro 4: 02.05.2026

### Key Takeaways:

- ▶ Overall good mood and good progress
- ▶ Use 80/20 rule for documentation during the implementation phase, so that the base structure is already there.
- ▶ Work more continuously on the Bachelor Thesis during the Week (1-2h every day)

## A. Project Management




Figure A.14.: Roadmap Sprint Planning 5

### A.7.5. Sprint 5: 03.05 –15.05.2026

| Week   | Responsible   | Task                                                                                            |
|--------|---------------|-------------------------------------------------------------------------------------------------|
| Week 1 | Frank & Lukas | Focus on thesis document and prepare evaluation                                                 |
| Week 2 | Lukas         | Build physical test environment                                                                 |
| Week 2 | Lukas & Frank | Evaluate the scenarios in containerlab and test if the drills are sufficient to test in bfh lab |

Table A.9.: Sprint Planning 5

Sprint Planning 5: 30.04.2026

Sprint Review 5: 15.05.2026

Sprint Retro 5: 15.05.2026

#### A.7.6. Sprint 6: 15.05 –28.05.2026

Sprint Planning 6: 14.05.2026

Sprint Review 6: 28.05.2026

Sprint Retro 6: 28.05.2026



#### A.7.7. Sprint 7: 29.05 –11.06.2026

Sprint Planning 7: 28.05.2026

Sprint Review 7: 11.06.2026

Sprint Retro 7: 11.06.2026

Rewrite  
in past  
tense

### A.7.8. Risk Assessment

To methodically conduct the risk assessment the section “Risk assessment” from an online guide published by Advisera [26], which describes the risk assessment process according to ISO 27001, was used as a reference.

#### Risk Assessment Process:

1. Risk categories and criteria: Define which dimensions and taxonomies are important for the assessment.
2. Risk identification: Collect possible risks in a brainstorming session.
3. Risk analysis: Assess the risks criteria on a risk matrix.
4. Risk evaluation: Create a list of risks ordered by the resulting risk level and evaluate if they are acceptable or need treatment.
5. Risk treatment: Describe the strategy and how to treat the risks.

#### Risk Categories & Criteria:

The table A.10 shows four commonly used categories, which were used to group the risks.

| Category         | Description                                                       |
|------------------|-------------------------------------------------------------------|
| Strategic Risk   | Risks related to business decisions (e.g., wrong prioritization). |
| Operational Risk | Risks related to internal processes (e.g., Scrum, QA process).    |
| Financial Risk   | Risks related to costs and expenses (e.g., hardware costs).       |
| External Risk    | Risks related to uncontrollable sources (e.g., force majeure).    |

Table A.10.: Risk categories used for the project risk assessment

The criteria for analysing the risks are summarized in the table A.11

| Criterion   | Description                                        |
|-------------|----------------------------------------------------|
| Probability | Likelihood that the risk will occur.               |
| Impact      | Severity of consequences if the risk materializes. |
| Level       | Risk level calculated as Probability + Impact.     |

Table A.11.: Risk evaluation criteria

The levels of the criteria and their respective numerical weights are listed in the tables A.12 and A.13.

| Level  | Weight | Description |
|--------|--------|-------------|
| Low    | 1      | Unlikely    |
| Medium | 2      | Possible    |
| High   | 3      | Likely      |

Table A.12.: Risk probability scale

| Level  | Weight | Description                        |
|--------|--------|------------------------------------|
| Low    | 1      | Minor inconveniences               |
| Medium | 2      | Delays milestones or reduces scope |
| High   | 3      | Threat to the thesis               |

Table A.13.: Risk impact scale

The following table [A.14](#) defines the calculation method and the resulting risk levels, used for the assessment.

| Risk Level | Weight | Calculation                  |
|------------|--------|------------------------------|
| Very Low   | 2      | Low + Low                    |
| Low        | 3      | Low + Medium                 |
| Medium     | 4      | Medium + Medium / Low + High |
| High       | 5      | Medium + High                |
| Critical   | 6      | High + High                  |

Table A.14.: Risk level classification (probability and impact are interchangeable)

### Risk Identification

The identification process was conducted with the help of a digital brainstorming, where the risks are already grouped by category. The identified risks are represented by the cards shown in [A.15](#). The respective risk categories are visualized by the color coding of the cards and described in the legend.

### Risk Brainstorming

- **Strategic Risk:** Risks related to business decisions (e.g. wrong prioritization).
- **Operational Risk:** Risks related to internal processes/procedures (e.g. Scrum, QA process).
- **Financial Risk:** Risks related to costs and expenses (e.g. hardware costs).
- **External Risk:** Risks related to uncontrollable sources (e.g. Force majeure).




Figure A.15.: Risks Brainstorming

### Risk Analysis

For analysing the risks a matrix was prepared, where the criteria “Impact” and “Probability” are representing the x- and y-axis respectively. The project team then assessed the criteria for each individual risk and placed them on the matrix accordingly. In a second step, the risks where moved if they did not fit relatively to each other. The resulting “Risk Analysis Matrix” is illustrated in the following figure [A.16](#). The background color of the cells is representing the resulting risk level of the tasks.

## Risk Matrix

|             |        |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     |
|-------------|--------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Probability | High   | <div>The NSAK framework/device itself is vulnerable</div> <div>User do not understand how to operate NSAK</div> <div>Regressions: Tested scenarios stop working, because a drill was changed</div> <div>Not enough time for testing (drills, scenarios and environments)</div> <div>Simulate or build network environments need more time as planned</div> <div>AI Credits are to expensive to build and test the AI scenario</div>                                                                                                                                                                                                                                                                 |
|             | Medium | <div>An existing framework already does everything we have planned</div> <div>The AI might go rough in the AI scenario</div> <div>Complex Network environments need more hardware than available</div> <div>Scenarios and drills need more time as planned</div> <div>Not enough time for documentation</div> <div>Scope is too large for bachelor thesis</div> <div>Drills not compatible between scenarios</div> <div>Misaligned communication lead to wrong expectations</div> <div>The AI bubble bursts and we need to switch AI provider</div> <div>Not enough time for media artifacts (Poster, Presentations, Film, Book, etc.)</div> <div>Scenario or drill to complex for automation</div> |
|             | Low    | <div>Lack of resources for attack scenarios</div> <div>Lack of resources/consensus for defense scenarios</div> <div>The NSAK hardware devices stop working</div> <div>NSAK gets abused by hackers</div> <div>One of the students gets injured or sick</div>                                                                                                                                                                                                                                                                                                                                                                                                                                         |
|             |        | <div>Low</div> <div>Medium</div> <div>High</div>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                    |
|             |        | Impact                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |

Figure A.16.: Risks Brainstorming


## Risk Evaluation &amp; Treatment

In a first step a table with the required columns was created. Then the risks and weights for the criteria were added with an ID in the format “R-000”, to easier reference them in the thesis later on. After sorting the table by the calculated risk level, the risk evaluation was conducted by deciding which treatment strategy described in A.15 we want to apply.

| Strategy        | Description                                      |
|-----------------|--------------------------------------------------|
| <b>Mitigate</b> | Reduce the likelihood or impact of a risk.       |
| <b>Transfer</b> | Move the risk to a third party.                  |
| <b>Accept</b>   | Acknowledge the risk and take no further action. |
| <b>Avoid</b>    | Eliminate or circumvent the cause of the risk.   |

Table A.15.: Risk Treatment Strategies

Finally, the project team went through the risks and discussed how to treat them effectively. The resulting “Risk Assessment” is illustrated in the following table A.17.



## A.8. Meeting Notes

### A.8.1. Checkpoint Meeting 1: 26.02.26

|                     |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |
|---------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Participants</b> | Hj. Wenger (Instructor), Frank Gauss, Lukas von Allmen                                                                                                                                                                                                                                                                                                                                                                                                                                          |
| <b>Agenda</b>       | <ul style="list-style-type: none"> <li>- Goals presentation via slides</li> <li>- Clarification whether there is a deadline or submission required for the presentation (besides adding it to the addendum).</li> <li>- Clarification whether the defense presentation should be an enhanced version of the public presentation or a separate, more technical presentation.</li> <li>- Proposal to establish a recurring meeting series every Thursday at 15:00 during the semester.</li> </ul> |
| <b>Tasks</b>        | <ul style="list-style-type: none"> <li>- Hand in a set of success criteria for the project goals: We thought we can delivery the success criteria and the project goals as milestones in gitlab, we corrected this later on.</li> <li>- Conduct a risk assessment: <a href="#">A.7.8</a></li> <li>- Contact expert via e-mail: Contacted at 21.02.26, response 03.03.26, invited to next checkpoint meeting 05.03.26, the expert takes part at the meeting every two weeks</li> </ul>           |

Add  
slides  
if we  
can re-  
produce  
them

Table A.16.: Meeting Notes: Checkpoint 1

### A.8.2. Checkpoint Meeting 2: 05.03.26

Add  
slides  
if we  
can re-  
produce  
them



|                     |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        |
|---------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Participants</b> | Hj. Wenger (Instructor), Urs Keller (Expert), Frank Gauss, Lukas von Allmen                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |
| <b>Agenda</b>       | <ul style="list-style-type: none"> <li>- Review of the project goals as milestones and project planning.</li> <li>- Review of risk assessment.</li> <li>- Discussion on the investigation and analysis of network communication between end devices and external networks (RED scenario: sniffing, filtering, and potential data exfiltration).</li> </ul>                                                                                                                                                                                                             |
| <b>Decisions</b>    | <ul style="list-style-type: none"> <li>- Variant 1 was selected for further development.</li> <li>- Open questions remain regarding the use and management of AI tokens.</li> </ul>                                                                                                                                                                                                                                                                                                                                                                                    |
| <b>Notes</b>        | <ul style="list-style-type: none"> <li>- We misunderstood how to deliver the project goals and success criteria: Instead of defining the goals as milestone we added a goal section in the thesis.</li> <li>- Discussion on the rationale behind selecting Variant 1.</li> <li>- Consideration of whether a user interface is required when MCP functionality is available - Open Research Question.</li> <li>- NSAK integrates AI capabilities.</li> <li>- The development of a Web GUI is currently considered lower priority - suggestion of Urs Keller.</li> </ul> |
| <b>Tasks</b>        | <ul style="list-style-type: none"> <li>- Send the slides of the final presentation and documentation of "Project 2" to Urs Keller (Expert): "Done" (via E-Mail).</li> <li>- Delivery the projects goals with a detailed description as part of the thesis by Saturday, 07.03.26: "Done" <a href="#">1.3</a></li> <li>- Finalize the project planning: "Done" <a href="#">A.6</a></li> <li>- Write the first few chapters of the thesis: "Done"</li> </ul>                                                                                                              |

Table A.17.: Meeting Notes: Checkpoint 2

## A.8.3. Checkpoint Meeting 3: 19.03.26

|                     |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
|---------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Participants</b> | Hj. Wenger (Instructor), Urs Keller (Expert), Frank Gauss, Lukas von Allmen                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
| <b>Agenda</b>       | <ul style="list-style-type: none"> <li>- Review Project Management and Roadmap - Review of the active milestones (Framework Comparison, AI-Research, Configuration Management)</li> <li>- Discuss Variant Decisions: AI-Integration</li> <li>- Ask about AI token usage</li> <li>- Ask about repository access</li> </ul>                                                                                                                                                                                                                                        |
| <b>Decisions</b>    | <ul style="list-style-type: none"> <li>- We will proceed using NSAK as the basis for our thesis.</li> <li>- Open questions remain regarding the use and management of AI tokens.</li> </ul>                                                                                                                                                                                                                                                                                                                                                                      |
| <b>Notes</b>        | <ul style="list-style-type: none"> <li>-</li> <li>-</li> </ul>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |
| <b>Tasks</b>        | <ul style="list-style-type: none"> <li>- Document decision that we will continue to work with the NSAK framework: 2.2.3 - Expand on variant decision for AI-Integration: - Split diagram into 2 variants: - Expand in detail how the variants are working (e.g., sequence diagram): A.21 and A.23 - Configuration management: Move metadata and others below the device node: "Done" 5.1.2 - BFH TI: Has a cluster with studio macs for AI usage -&gt; Ask if we can use them (<a href="https://mlmp.ti.bfh.ch/">https://mlmp.ti.bfh.ch/</a> via VPN)</li> </ul> |

Table A.18.: Meeting Notes: Checkpoint 3

## A.8.4. Checkpoint Meeting 4: 02.04.26

|                     |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                |
|---------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Participants</b> | Hj. Wenger (Instructor), Frank Gauss, Lukas von Allmen                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         |
| <b>Agenda</b>       | <ul style="list-style-type: none"> <li>- Overview Roadmap and active milestones</li> <li>- Reproducible Test Environment: Introduce ContainerLab, Show diagram, Physical Implementation (could)</li> <li>- AI Scenario: Demo AI-agent connection to self-hosted, LangChain Research LLM</li> <li>- Review MCP Sequence Diagrams</li> <li>- Discuss what needs to be added to the main document and what to the appendix</li> <li>- Ask if it's possible to lend more HW for physical test setup</li> <li>- Ask for access to BFH AI Cluster</li> </ul>                         |
| <b>Decisions</b>    | <ul style="list-style-type: none"> <li>-</li> <li>-</li> </ul>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |
| <b>Notes</b>        | <ul style="list-style-type: none"> <li>-</li> <li>- Alternative to container LAB: Cisco - DNS3 (not open source)</li> <li>- It would be interesting to add more services to the DMZ in the reproducible test environment</li> <li>- Tutor: The current structure is preferred, we should also ask the expert</li> <li>-</li> </ul>                                                                                                                                                                                                                                             |
| <b>Tasks</b>        | <ul style="list-style-type: none"> <li>- Ask for access to BFH AI: <a href="https://mlmp.ti.bfh.ch">https://mlmp.ti.bfh.ch</a> (Requires VPN) - If we are not allowed, ask tutor for guidance</li> <li>- Variant decisions: Add pros/cons and elaborate why we choose a variant</li> <li>- AI-agent: Design and implement logging</li> <li>- Ask the expert about the preferred document structure: Thesis vs Appendix</li> <li>- Create a HW list for tutor, date and configuration (must be planned, device configuration could be done by assistants)</li> <li>-</li> </ul> |

Table A.19.: Meeting Notes: Checkpoint 4

## A.8.5. Checkpoint Meeting 5: 16.04.26

|                     |                                                                                                                                                                                                                                                               |
|---------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Participants</b> | Hj. Wenger (Instructor), Urs Keller (Expert), Frank Gauss, Lukas von Allmen                                                                                                                                                                                   |
| <b>Agenda</b>       | <ul style="list-style-type: none"> <li>- Demo: Reproducible test environment</li> <li>- Demo: AI Agent</li> <li>- Show current state of the project goals and roadmap</li> <li>- Fix the date for bachelor thesis defense: Ideally 17or 18.04.2026</li> </ul> |
| <b>Decisions</b>    | - Defense Date: 18.06.2026, 13:30 - 15:30                                                                                                                                                                                                                     |
| <b>Tasks</b>        | <ul style="list-style-type: none"> <li>- Add a system prompt to tell the LLM to create a human-readable report of what it did, parallel to the logs:</li> <li>- Export logs to a remote server like Zabbix or Grafana:</li> </ul>                             |

Table A.20.: Meeting Notes: Checkpoint 5

## A.8.6. Checkpoint Meeting 6: 30.04.26

|                     |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     |
|---------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Participants</b> | Hj. Wenger (Instructor), Frank Gauss, Lukas von Allmen                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |
| <b>Agenda</b>       | <ul style="list-style-type: none"> <li>- Demo: Reconnaissance scenario</li> <li>- Demo: Grafana / Loki (Authentication Issue)</li> <li>- Show Roadmap and next two weeks</li> <li>- How can we compare and evaluate the AI with the Reconnaissance scenario?</li> </ul>                                                                                                                                                                                                                                                                             |
| <b>Notes</b>        | <ul style="list-style-type: none"> <li>- Take a look at nmap plugins</li> <li>- To compare and evaluate the AI with the Reconnaissance scenario, we could run both scenarios in the BFH LAB (physical or virtual). For a physical test run we would need to fix a date with HJ. Wenger.</li> <li>- Interesting metrics are: Speed, Correctness, Initial knowledge (pre-conditions), Creativity (3 - 5 interesting criteria)</li> <li>- Only configure Grafana / Loki to the point where we can easily analyze the steps and tools calls.</li> </ul> |
| <b>Decisions</b>    | -                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |
| <b>Tasks</b>        | -                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |

Table A.21.: Meeting Notes: Checkpoint 6

A.8.7. Checkpoint Meeting 7: 14.05.26

A.8.8. Checkpoint Meeting 8: 28.05.26

A.8.9. Checkpoint Meeting 9: 11.06.26

## Variant 1: AI Purple Team

|        | Research             | Framework                               | Scenarios<br>(Red- and Blue-Team)                 |                                                                   |
|--------|----------------------|-----------------------------------------|---------------------------------------------------|-------------------------------------------------------------------|
| Must   | Framework Comparison | Configuration Management<br><br>WEB GUI | AI Scenario Red Team<br><br>AI Scenario Blue Team | Intrusion Detection  Blue Team<br><br>Network Discovery  Red Team |
| Should |                      | Framework Improvements                  |                                                   |                                                                   |
| Could  | GAP Analysis         |                                         |                                                   |                                                                   |

Figure A.18.: Variant 1: AI vs Manually Built Scenarios

## A.9. Variant Decision: Thesis Goals

During the first sprint of the project, which was focused on project management and planning, we identified the very promising hypothesis that red and blue team activities could be automated and/or enhanced by AI. Pursuing this hypothesis would change the projects focus and goals and thus requires the agreement and approval of the stakeholder. To communicate this circumstance efficiently and leave the option to go with the status quo, we prepared two options to conduct a “Variant Decision”. In the following two subsections the two variants presented in the second checkpoint meeting [A.17](#) are described and illustrated. It must be noted, that the variants and their visualizations [A.18](#) [A.19](#) are momentary snapshots and will not be updated for traceability reasons. For example, the Web GUI was already deprioritized during the discussion of the variants, which is reflected in the project goals [1.3](#) but not here.

### A.9.1. Variant 1: AI vs Manually Built Scenarios

Figure [A.18](#) shows the project goals adjusted for the AI focused hypothesis, which is favored variant of the project team. In this option most of the evaluated scenarios will not be realized in favor of AI based scenarios, which also requires additional research [1.3.1](#). The basic idea is to implement scenarios for network discovery [1.3.3](#), intrusion detection [1.3.3](#) and an AI based scenario with the same capabilities [1.3.3](#). In the evaluation of the thesis we can then assess the quality of the AI based scenarios itself [1.3.4](#) and in comparison to the manually built scenarios [1.3.4](#).

## Variant 2

|        | Research             | Framework                               | Scenarios<br>(Red- and Blue-Team)         |                                                                  |
|--------|----------------------|-----------------------------------------|-------------------------------------------|------------------------------------------------------------------|
| Must   | Framework Comparison | Configuration Management<br><br>WEB GUI | Traffic Capture and Exfiltration Red Team | Intrusion Detection Blue Team<br><br>Network Discovery  Red Team |
| Should |                      | Framework Improvements                  | TLS Downgrade<br>POODLE Red Team          |                                                                  |
| Could  | GAP Analysis         |                                         | Honeypot Blue Team                        |                                                                  |

Figure A.19.: Variant 2: Multiple Blue and Red Team Scenarios

### A.9.2. Variant 2: Multiple Blue and Red Team Scenarios (Status Quo)

This is the fallback if the other option [A.9.1](#) is not approved by the stakeholders. Even if the project team does not favor this more conservative variant, it still represents a valid way to conduct an interesting project. As illustrated below [A.19](#), in this variant we propose to proceed with the previously planned scenarios for which we already conducted “User Story Mappings” documented under the workshops section [A.11](#). The focus would lie in more conventional scenarios, overall framework improvements and an in depth comparison with another framework in form of a GAP Analysis.

### A.9.3. Decision

The stakeholders and project team decided in favor of **Variant 1: AI vs Manually Built Scenarios** [A.9.1](#) in the checkpoint meeting [A.17](#).

## A.10. Variant Decision: AI Integration

During the AI-research 2.2 we identified two interesting ways to allow NSAK to interact with AI based technologies. As illustrated in A.20, the first variant purposes to implement an AI-agent directly into the framework and use it in scenarios and drills. In the second variant shown in A.22, the capabilities of NSAK and other relevant tools on the NSAK device would be exposed via MCP.

### A.10.1. Variant 1: AI-agent in NSAK

#### Variant 1: AI-agent in NSAK




Figure A.20.: AI Integration Variant 1: AI-agent in NSAK




Figure A.21.: Variant 1, AI Agent in NSAK

In this architecture, the AI agent is embedded within the NSAK device. The NSAK core establishes the operational guardrails and schema, and provides the initial



prompt to the AI Agent. The agent uses a local or remote LLM, which is accessible over the management network. It operates within an iterative loop, in which it may either invoke external tools via a command-line interface (e.g., nmap, Metasploit) or interact with them through an MCP server interface. The results of these operations are subsequently processed and returned to the NSAK core in a structured schema-compliant format.

### A.10.2. Variant 2: NSAK as MCP Server

#### Variant 2: NSAK as MCP Server




Figure A.22.: AI Integration Variant 2: NSAK as MCP server




Figure A.23.: Variant 2 NSAK MCP Server

Variant 1, in which the Agentic-AI is embedded within the NSAK device, enables the definition and enforcement of schemas and security policies within a controlled environment. Furthermore, it allows direct interaction with tools via the

command-line interface and integration with MCP servers. This results in a tightly coupled architecture with high control and consistency.

In contrast, Variant 2 relies on an external agent that accesses the NSAK device through an exposed MCP server interface. While this approach provides increased modularity, it also introduces limitations in flexibility, as the agent is restricted to the set of tools exposed by the NSAK device.

From a holistic and multi-layered architectural perspective, Variant 1 is preferred. It enables deeper system integration, higher security, and a wider range of executable tools.

### A.10.3. Decision

The project team decided in favor of **Variant 1: AI-agent in NSAK A.10.1**, as it promises to greatly increase the capabilities of NSAK regarding autonomy and automation.

enhance the decision with technical background

This is not correct anymore: we use LangChain Tools instead of MCP for tool calling, which is their own abstraction over vendor specific APIs

## A.11. Workshops

In the Project Management phase, several workshops were conducted, including a brainstorming session using Post-it notes and a User Story Mapping exercise on a whiteboard. The goal of these workshops was to clarify the scenarios a network edge device should support.

Following these workshops, the insights gathered were used to derive issues and goals, which were then mapped to detailed milestones in GitLab. All deliverables were categorized using the MoSCoW prioritization method into must, should, and could requirements. An example of this prioritization and scenario mapping is illustrated in Figure A.26. All Milestones links are provided in the URL at the top of the section.

### A.11.1. Brainstorm Scenario Selection

In this workshop, participants brainstorm and cluster potential red and blue team scenarios for practical cybersecurity exercises.




Figure A.24.: Brainstorm Scenarios




Figure A.25.: Clustering possible Scenarios

The Brainstorm session is used to derive further red/blue team scenarios and led to the pivot of AI-enabled Nsac scenario.

### A.11.2. NSAK Scenario: Network Discovery - Red Team

#### Network Discovery Scenario

##### Objectives

This workshop aims to evaluate the different steps of the network discovery process and to help participants understand how attackers identify and map potential targets within a network.

Building on this objective, Red Teams consider network discovery one of the most essential activities during the reconnaissance phase, as it helps them identify hosts and services and evaluate potential targets. Once a network map is created, the team can use it in later phases. For example, they might execute an ARP spoofing attack, which allows them to intercept network traffic more efficiently.

In relation to these activities, this scenario addresses the following point from the thesis proposal:

‘Unterstützung typischer Vorgehensweisen aus dem Red-Team- und Penetration-Testing-Bereich, etwa durch Scan- oder Angriffssimulationen.’

The workshop led to the milestone *NSAK Scenario: Network Discovery – Red Team (must)*, in which each step for discovering and mapping a network is documented and structured for later implementation.




Figure A.26.: NSAK Scenario: Network Discovery - Red Team

### A.11.3. NSAK Scenario: TLS Downgrade Poodle - Red Team

#### TLS Downgrade Poodle Milestone

##### Objectives

This workshop demonstrates how attackers can exploit TLS downgrade vulnerabilities, such as POODLE, to intercept encrypted communication.

In this scenario, the infamous POODLE attack is implemented, in which a malicious actor tries to convince a client and a server to use the vulnerable TLS version 1.1 instead of 1.2 or 1.3 to encrypt HTTP traffic. This attack requires the attacker to sit between the client and the server (a man-in-the-middle, or MITM), allowing them to intercept and control the traffic. When the client initiates a TLS handshake, the attacker drops packets to the server and sends TCP reset packets to the client, eventually trying older TLS versions. Because TLS 1.1 is vulnerable to a CBC Padding attack, it is possible to create a padding oracle that gradually decrypts the cookie and hijacks the user's session.

Building on this, the scenario directly addresses the following aspect from the thesis proposal:

‘Unterstützung typischer Vorgehensweisen aus dem Red-Team- und Penetration-Testing-Bereich, etwa durch Scan- oder Angriffssimulationen’




Figure A.27.: NSAK Scenario: TLS Downgrade Poodle - Red Team






Figure A.28.: NSAK Scenario: TLS Downgrade Poodle -Topology



## A.11.4. NSAK Scenario: Honeypot - Blue Team

## Honeypot Milestone

Honeyspots can serve as an additional form of intrusion detection and help Blue Teams detect attackers in a network. The system tries to attract malicious actors by providing an attack surface as large as possible; a simple way to achieve this is to open ports. After an attacker tries to connect to any of the open ports, the system notifies an operator and may close all open ports to prevent further attempts to compromise the Honeypot. The operator can then manually intervene to identify and isolate the attacker and analyze the captured payload.

This scenario covers the following bullet point from the thesis proposal:  
 ‘Bereitstellung von Funktionen zur Erkennung, Protokollierung und Analyse verdächtiger Aktivitäten als Teil defensiver Massnahmen’




Figure A.29.: NSAK Scenario: Honeypot - Blue Team

## A.11.5. NSAK Scenario: Traffic Capture - Red Team

## Traffic Capture Milestone

Capturing, exfiltrating, and analyzing traffic from a targeted network is a key Red Team activity and crucial for initial reconnaissance and planning the next phases of the attack. Based on the gathered intelligence, network participants and services can be identified, and possible attack vectors may be evaluated. If unencrypted traffic is present on the network, session data and credentials could already have been extracted.

This scenario addresses the following bullet point from the thesis proposal: ‘Untersuchung und Auswertung der Netzwerkkommunikation zwischen Endgeräten und externen Netzen.’




Figure A.30.: NSAK Scenario: Traffic Capture - Red Team

#### A.11.6. NSAK Scenario: Intrusion Detection - Blue Team

##### Intrusion Detection Milestone

Intrusion detection is one of the core Blue Team activities and aims to identify malicious actors and misconfigured clients in production network environments. This can be achieved by introducing a system at a traffic mirroring port or as a central point in a network. The intrusion detection system can leverage a wide array of security policies, reference traffic, and other markers to effectively iden-

tify anomalies. As soon as suspicious activity is detected or a policy is violated, automated measures may be initiated, and security engineers will be automatically alerted for manual intervention.

This scenario covers the following bullet point from the thesis proposal: “Bereitstellung von Funktionen zur Erkennung, Protokollierung und Analyse verdächtiger Aktivitäten als Teil defensiver Massnahmen”



Figure A.31.: NSAK Scenario: Intrusion Detection - Blue Team

#### A.11.7. NSAK Scenario: AI - Purple Team

##### AI Purple Team Milestone

Even though we were not quite hooked by the current hype around AI during the predecessor project 1.2, we now see a huge potential in integrating it into the NSAK-framework, especially in because of the tool-calling capabilities of agentic AI. In our point of view, the possibility that common blue and red team activities, which require a lot of know how and time, can be simplified or speed up is really promising. This scenarios would explore if it's possible that intrusion detection or network reconnaissance tasks can be assisted or maybe even fully automated by an AI-agent.




Figure A.32.: NSAK Scenario: AI - Purple Team

## A.12. Self-Hosted LLM

During the risk assessment [A.7.8](#) we identified that the AI tokens required to use cloud based LLMs might be very expensive, and it's really hard to estimate the costs. One of the possibilities to avoid this risk was the usage of self-hosted LLMs, and luckily we already had existing hardware at home which is suitable for running smaller LLMs. Besides eliminating the costs for AI tokens, at least during the development phase, it also allows the quick evaluation of different AI Models. Depending on the quality and performance of the smaller models running in the self-hosted setup we can consider not using a cloud based LLM at all.

### A.12.1. Hardware Setup

| Component     | Specification                                          |
|---------------|--------------------------------------------------------|
| CPU           | AMD Ryzen 5 3600 (6 cores / 12 threads, up to 4.2 GHz) |
| GPU           | NVIDIA GeForce RTX 3070 (8 GB VRAM)                    |
| System Memory | 32 GB RAM                                              |

Table A.22.: Hardware configuration of the self-hosted LLM inference server.

The 8 GB VRAM of the RTX 3070 allows us to run models with up to 7 billion parameters, depending on the specific model [\[27\]](#).

### A.12.2. System Setup

The system setup described in [35](#) will configure and expose the Ollama API and OpenWeb UI via NGINX, which provided a chat GUI to interact with Ollama. In contrast to Ollama, which runs directly on the host, the other services are run in containers managed by rootless podman and defined the docker-compose.yml [37](#). The NGINX default.conf [43](#) configures the web server to terminate the SSL connection and act as reverse proxy for the “Open WebUI” and the Ollama API under the directory /llm/.

Services:

- ▶ Ollama: LLM manager, exposing an API for prompts and other interactions (host).
- ▶ Swag: NGINX web server with Certbot for automated SSL certificate generation (containerized).
- ▶ Open WebUI: Web based chat interface for LLM managers (containerized).

Refactor the section: Move all information in their respective subsections, etc.

## System Hardening

It is recommended to put a proper firewall before the services we want to expose to the internet, but for our purposes we can just disable the port forwarding when we don't use the AI. It still makes sense to perform a minimal system hardening as shown in the following listing [33](#), before exposing the services to the internet.

---

```
1  # Install UFW
2  sudo pacman -Syu ufw
3  sudo ufw status
4
5  # Allow SSH for remote access
6  sudo ufw allow ssh
7
8  # Allow HTTP for certbot HTTP challenge
9  sudo ufw allow http
10
11 # Allow HTTPS for Open WebUI
12 sudo ufw allow https
13
14 # Allow default Ollama port
15 sudo ufw allow 11434/tcp
16
17 # Make sure to run this command after adding the "allow" rules!
18 sudo ufw enable
19
20 sudo ufw status
```

---

Listing 33: UFW configuration

## Expose Open WebUI and Ollama

For exposing the services from the home network via HTTPS to the internet, we need a domain name which points to the home router. To achieve this we can use DynDNS, which is provided for free by Infomaniak where the domain “hiube.ch” was registered, and is supported by the FRITZ!Box used in the home network. The screenshots [A.34](#) and [A.33](#) are showing the respective configurations.




Figure A.33.: FRITZ!Box DynDNS Setup




Figure A.34.: Infomaniak DynDNS Setup

Now that the domain name “ai.hiube.ch” points to the router in the home network, we need to redirect the HTTP(S) traffic to the LLM server. This can be achieved with port forwarding rules on the FRITZ!Box, as shown in the screenshot [A.35](#).




Figure A.35.: FRITZ!Box DynDNS Setup

The Open WebUI and Ollama API should already be reachable via “ai.hiube.ch”, but we will receive a security warning because there is no valid SSL certificate. The easiest way to generate the certificate, is to trigger certbot by restarting the “swag” service as shown in the listing below 34.

```
1 # Login as user ai and restart the swag service
2 podman-compose restart swag
```

Listing 34: Restart SWAG Container

A.12.3. Evaluation

For quickly evaluate models and assess their capabilities we can now use the Open WebUI served under “ai.hiube.ch”. The screenshot A.36 shows a quick test with “qwen2.5-coder:7B”.



## A. Project Management

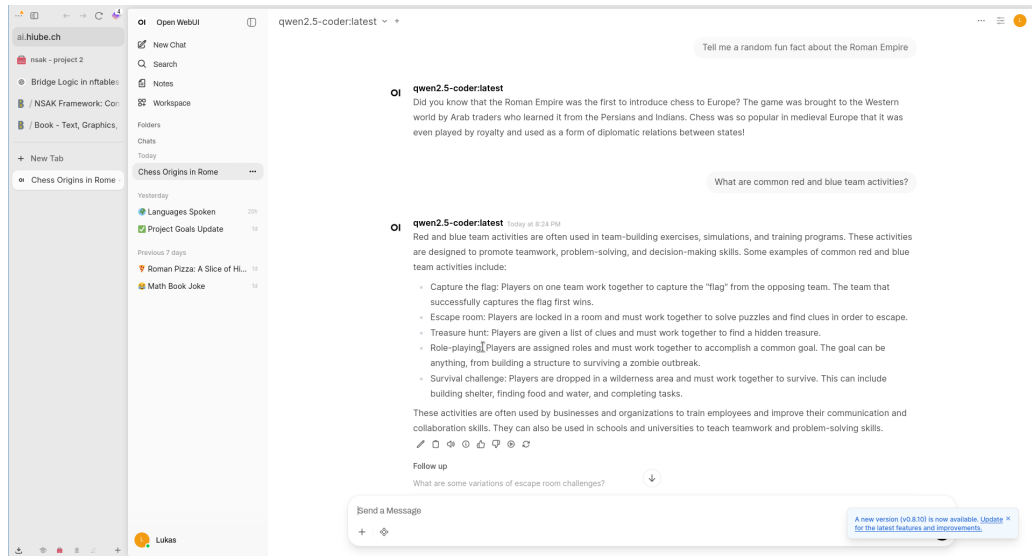


Figure A.36.: FRITZ!Box DynDNS Setup

Later on we can use the Ollama API served under `/llm/` as a backend for AI agents.

#### A.12.4. Self-Hosted Base Setup

---

```
1  # Login as ai user
2  ssh ai@10.10.10.20
3
4  # Create a docker compose file and add contents from
   ↳ lst:nginx-docker-compose.yml
5  vim docker-compose.yml
6
7  # Start the services, so that they create the volume mappings
   ↳ etc.
8  # The containers will most likely fail, but that's fine for now.
9  podman-compose up -d
10 podman-compose down
```

---

Listing 35: Self-Hosted Base Setup

---

```
1 # Login as ai user
2 ssh ai@10.10.10.20
3
4 # Create a docker compose file and add contents from
   ↪ lst:nginx-docker-compose.yml
5 vim docker-compose.yml
6
7 # Start the services, so that they create the volume mappings
   ↪ etc.
8 # The containers will most likely fail, but that's fine for now.
9 podman-compose up -d
10 podman-compose down
```

---

Listing 36: Nginx & Let's Encrypt Instructions

A.12.5. Nginx & Let's Encrypt - Reverse Proxy with SSL Termination

---

```
1 services:
2   swag:
3     image: lscr.io/linuxserver/swag:latest
4     container_name: swag
5     cap_add:
6       - NET_ADMIN
7     environment:
8       - PUID=1000
9       - PGID=1000
10      - TZ=Europe/Zurich
11      - URL=hiube.ch
12      - VALIDATION=http
13      - SUBDOMAINS=ai,
14      - EMAIL=<my-email>
15      - ONLY_SUBDOMAINS=true
16      - EXTRA_DOMAINS= #optional
17      - STAGING=false #optional
18      - DISABLE_F2B= #optional
19      - SWAG_AUTORELOAD= #optional
20      - SWAG_AUTORELOAD_WATCHLIST= #optional
21     volumes:
22       - /home/ai/swag:/config
23     ports:
24       - 443:443
25       - 80:80
26       - 443:443/udp
27     restart: unless-stopped
```

---

Listing 37: Nginx &amp; Let's Encrypt Compose File

A.12.6. Ollama - LLM Server

---

```
1 ssh ai@10.10.10.20
2
3 # Check nvidia driver is loaded and the GPU was correctly
  ↳ detected, otherwise install the according drivers
4 nvidia-smi
5
6 # Install Ollama and enable systemd service
7 sudo pacman -Syu ollama
8 sudo systemctl enable --now ollama
9 sudo systemctl status ollama
10
11 # Download and run a qwen coder 7b model for testing
12 # Later used models:
13 # - `qwen2.5:7b-instruct` (tool calling)
14 # - `gemma4:e2b` (was hyped, but is not that good in tool
  ↳ calling)
15 ollama pull qwen2.5-coder:7b
16 ollama run qwen2.5-coder
17
18 # Open the nginx default.conf and add the contents from
  ↳ lst:ollama-nginx-conf
19 vim swag/nginx/site-confs/default.conf
20
21 # Start the services, this time they should run successfully.
22 # For a local setup, certbot will fail to perform the http
  ↳ challenge for issuing the TLS certificate,
23 # as long as DynDNS and the port forwarding are not set up
  ↳ properly.
24 podman-compose up -d
```

---

Listing 38: Ollama Instructions

```
1 # Redirect all traffic to HTTPS
2 server {
3     listen 80 default_server;
4     listen [::]:80 default_server;
5
6     server_name ai.hiube.ch;
7
8     location / {
9         return 301 https://$host$request_uri;
10    }
11 }
12
13 # Using upstream enables us to restart ollama without restarting
14 ↳ the nginx
15 upstream ollama {
16     server host.docker.internal:11434;
17 }
18
19 server {
20     listen 43434 ssl;
21     listen [::]:43434 ssl;
22
23     server_name ai.hiube.ch;
24
25     include /config/nginx/ssl.conf;
26
27     location / {
28         auth_basic "Restricted";
29         auth_basic_user_file /config/nginx/.htpasswd;
30
31         proxy_pass http://ollama;
32         proxy_http_version 1.1;
33         proxy_set_header Connection "";
34         chunked_transfer_encoding off;
35         proxy_buffering off;
36         proxy_cache off;
37         proxy_set_header Host $host;
38         proxy_set_header X-Real-IP $remote_addr;
39         proxy_set_header X-Forwarded-For
40             ↳ $proxy_add_x_forwarded_for;
41         proxy_read_timeout 600s;
42         proxy_send_timeout 600s;
43     }
44 }
```

```
1  # Enable Ollama debug logs
2
3  During the development of the AI intrusion detection scenario,
4  ↪ Ollama suddenly sent an empty response terminating the
5  ↪ agent.
6  For investigation, I enabled the debug logs on the AI server.
7
8  ```bash
9  sudo systemctl edit --full ollama.service
10
11 # Add the following line
12 Environment="OLLAMA_DEBUG=2"
13
14 # Restart the service
15 sudo systemctl restart ollama.service
16
17 # The logs can now be listed with:
18 journalctl -f -b -u ollama
19 ```
20
21 ## Container logs of the empty Ollama response:
22
23 -> The last log indicates a container restart (scenario was
24 ↪ terminated by empty response)
25
26 ```bash
27 > sudo podman logs -f
28 ↪ simple_tcp_client_server_ai_intrusion_detection_1
29 WARNING:nsak.core.scenario.scenario_manager:EXEC SCENARIO0: AI
30 ↪ Intrusion Detection
31 INFO:root:Starting scenario: AI Intrusion Detection
32 WARNING:nsak.core.drill.drill_manager:EXEC DRILL: Discover Hosts
33 DEBUG:httpcore.connection:connect_tcp.started host='ai.hiube.ch'
34 ↪ port=43434 local_address=None timeout=None
35 ↪ socket_options=None
36 DEBUG:httpcore.connection:connect_tcp.complete
37 ↪ return_value=<httpcore._backends.sync.SyncStream object at
38 ↪ 0x7f6434e86f90>
39 DEBUG:httpcore.connection:start_tls.started
40 ↪ ssl_context=<ssl.SSLContext object at 0x7f6434eff4d0>
41 ↪ server_hostname='ai.hiube.ch' timeout=None
42 DEBUG:httpcore.connection:start_tls.complete
43 ↪ return_value=<httpcore._backends.sync.SyncStream object at
44 ↪ 0x7f6434eab110>
45 DEBUG:httpcore.http11:send_request_headers.started
46 ↪ request=<Request [b'POST']>
47 DEBUG:httpcore.http11:send_request_headers.complete
48 DEBUG:httpcore.http11:send_request_body.started request=<Request
49 ↪ [b'POST']>
50 DEBUG:httpcore.http11:send_request_body.complete
51 DEBUG:httpcore.http11:receive_response_headers.started
52 ↪ request=<Request [b'POST']>
53 DEBUG:httpcore.http11:receive_response_headers.complete
```

### A.12.7. Open WebUI - Web Chat for LLM Servers

---

```
1 ssh ai@10.10.10.20
2
3 # Open the nginx default.conf and add the contents from
4   ↪ lst:open-webui-nginx-conf
5 vim swag/nginx/site-confs/default.conf
6
7 # Start the services
8 podman-compose up -d
```

---

Listing 41: Open WebUI Instructions



```
1 services:
2   # ...
3
4   open-webui:
5     image: ghcr.io/open-webui/open-webui:main
6     container_name: open-webui
7     volumes:
8       - /home/ai/open_webui:/app/backend/data # Use a named
9         ↪ volume for persistence
10    environment:
11      OLLAMA_BASE_URL: http://host.docker.internal:11434
12    restart: always
```

---

Listing 42: Open WebUI Compose File

```
1  # Open WebUI
2  upstream open_webui {
3      server open-webui:8080;
4  }
5
6  server {
7      listen 443 ssl default_server;
8      listen [::]:443 ssl default_server;
9
10     server_name ai.hiube.ch;
11
12     include /config/nginx/ssl.conf;
13
14     location / {
15         proxy_pass http://open_webui;
16         proxy_set_header Host $host;
17         proxy_set_header X-Real-IP $remote_addr;
18         proxy_set_header X-Forwarded-For
19             ↪ $proxy_add_x_forwarded_for;
20         proxy_set_header X-Forwarded-Proto $scheme;
21         proxy_http_version 1.1;
22         proxy_set_header Upgrade $http_upgrade;
23         proxy_set_header Connection "upgrade";
24         proxy_read_timeout 900s;
25         proxy_send_timeout 900s;
26     }
27 }
```

---

Listing 43: Open WebUI Nginx Configuration

A.12.8. Drawio MCP - Diagram Tool

---

```
1 ssh ai@10.10.10.20
2
3 # https://www.drawio.com/blog/diagrams-docker-app
4 echo "unqualified-search-registries = [\"docker.io\"]" >>
   ↪ /etc/containers/registries.conf
5
6 # Add drawio.hiube.ch CNAME entry to ai.hiube.ch
7 # Add drawio.hiube.ch to linuxserver.io swag config for
   ↪ letsencrypt and add the drawio service from below
8 vim docker-compose.yaml
9
10 # Update nginx config
11 sudo vim swag/nginx/site-confs/default.conf
12
13 # Download drawio-mcp
14 git clone https://github.com/lgazo/drawio-mcp-server.git
15
16 # Restart services
17 podman-compose restart swag
```

---

Listing 44: Drawio Instructions

```
1 services:
2   # ...
3
4   drawio-mcp-server:
5     build: ./drawio-mcp-server
6     container_name: drawio-mcp-server
7     restart: unless-stopped
8     command:
9       - "node"
10      - "packages/drawio-mcp-server/build/index.js"
11      - "--transport"
12      - "http"
13      - "--editor"
14      - "--http-port"
15      - "3000"
```

---

Listing 45: Drawio Compose File

```
1 # Draw.io
2 upstream drawio {
3     server drawio:8080;
4 }
5
6 upstream drawio_mcp {
7     server drawio-mcp-server:3000;
8 }
9
10 upstream drawio_ws {
11     server drawio-mcp-server:3333;
12 }
13
14 server {
15     listen 443 ssl;
16     listen [::]:443 ssl;
17
18     server_name drawio.hiube.ch;
19
20     include /config/nginx/ssl.conf;
21
22     location / {
23         auth_basic "Restricted";
24         auth_basic_user_file /config/nginx/.htpasswd;
25
26         proxy_pass http://drawio;
27         proxy_set_header Host $host;
28         proxy_set_header X-Real-IP $remote_addr;
29         proxy_set_header X-Forwarded-For
30             ↪ $proxy_add_x_forwarded_for;
31         proxy_set_header X-Forwarded-Proto $scheme;
32         proxy_http_version 1.1;
33         proxy_set_header Upgrade $http_upgrade;
34         proxy_set_header Connection "upgrade";
35     }
36
37     location /mcp {
38         auth_basic "Restricted";
39         auth_basic_user_file /config/nginx/.htpasswd;
40
41         proxy_pass http://drawio_mcp/mcp;
42         proxy_set_header Host $host;
43         proxy_set_header X-Real-IP $remote_addr;
44         proxy_set_header X-Forwarded-For
45             ↪ $proxy_add_x_forwarded_for;
46         proxy_set_header X-Forwarded-Proto $scheme;
47         proxy_http_version 1.1;
48         proxy_set_header Upgrade $http_upgrade;
49         proxy_set_header Connection "upgrade";
50         proxy_buffering off;
51         proxy_cache off;
52         proxy_read_timeout 3600s;
53         proxy_send_timeout 3600s;
54         chunked_transfer_encoding on;
```

A.12.9. Netbox - Network Inventory System

---

```
1 ssh ai@10.10.10.20
2
3 #
4   ↳ https://learn.srlinux.dev/blog/2024/creating-a-network-digital-twin-with-c
5   ↳ https://github.com/netbox-community/netbox-docker
6 git clone -b release
7   ↳ https://github.com/netbox-community/netbox-docker.git
8 cd netbox-docker/
9
10 cp docker-compose.override.yml.example
11   ↳ docker-compose.override.yml
12
13 podman-compose up -d
14
15 podman-compose exec netbox /opt/netbox/netbox/manage.py
16   ↳ createsuperuser
17
18 cd ..
19
20 # Add netbox.hiube.ch CNAME entry to ai.hiube.ch
21 # Add netbox.hiube.ch to linuxserver.io swag config for
22   ↳ letsencrypt
23 vim docker-compose.yaml
24
25 # Update nginx conifg
26 sudo vim swag/nginx/site-confs/default.conf
27
28 # Restart services
29 podman-compose restart swag
```

---

Listing 47: Netbox Instructions

---

1

---

Listing 48: Netbox Compose File

```
1  # Netbox
2  upstream netbox {
3      server host.docker.internal:8000;
4  }
5
6  server {
7      listen 443 ssl;
8      listen [::]:443 ssl;
9
10     server_name netbox.hiube.ch;
11
12     include /config/nginx/ssl.conf;
13
14     location / {
15         proxy_pass http://netbox;
16         proxy_set_header Host $host;
17         proxy_set_header X-Real-IP $remote_addr;
18         proxy_set_header X-Forwarded-For
19             ↪ $proxy_add_x_forwarded_for;
20         proxy_set_header X-Forwarded-Proto $scheme;
21     }
22 }
```

---

Listing 49: Netbox Nginx Configuration



A.12.10. Grafana / Loki - Logging System

---

```
1 ssh ai@10.10.10.20
2
3 #
4   ↪ https://grafana.com/docs/loki/latest/setup/install/docker/#install-with-docker
5 mkdir loki
6 wget
7   ↪ https://raw.githubusercontent.com/grafana/loki/v3.7.0/examples/getting-started/docker-compose.yaml
8   ↪ -O loki/docker-compose.yaml
9 wget
10  ↪ https://raw.githubusercontent.com/grafana/loki/v3.7.0/examples/getting-started/alloy-local-config.yaml
11  ↪ -O loki/alloy-local-config.yaml
12 wget
13  ↪ https://raw.githubusercontent.com/grafana/loki/v3.7.0/examples/getting-started/loki-config.yaml
14  ↪ -O loki/loki-config.yaml
15
16 # Adjustments needed for rootless podman
17 systemctl --user enable --now podman.socket
18
19 # Adjust the alloy service described in the loki `compose.yaml`
20 vim loki/docker-compose.yaml
21
22 # Add grafana.hiube.ch CNAME entry to ai.hiube.ch
23 # Add grafana.hiube.ch to linuxserver.io swag config for
24   ↪ letsencrypt
25 vim docker-compose.yaml
26
27 # Update nginx config
28 sudo vim swag/nginx/site-confs/default.conf
29
30 # Restart services
31 podman-compose up -d
```

---

Listing 50: Grafana / Loki Instructions

```
1 # ~/loki/docker-compose.yaml
2 networks:
3   loki:
4
5 services:
6   read:
7     image: grafana/loki:latest
8     command: "-config.file=/etc/loki/config.yaml -target=read"
9     ports:
10      - 127.0.0.1:3101:3100
11      - 7946
12      - 9095
13     volumes:
14      - ./loki/loki-config.yaml:/etc/loki/config.yaml
15     depends_on:
16      - minio
17     healthcheck:
18      test: [ "CMD", "/usr/bin/loki", "-health" ]
19      start_period: 30s
20      interval: 10s
21      timeout: 5s
22      retries: 5
23     networks: &loki-dns
24     loki:
25       aliases:
26        - loki
27
28   write:
29     image: grafana/loki:latest
30     command: "-config.file=/etc/loki/config.yaml -target=write"
31     ports:
32      - 127.0.0.1:3102:3100
33      - 7946
34      - 9095
35     volumes:
36      - ./loki/loki-config.yaml:/etc/loki/config.yaml
37     healthcheck:
38      test: [ "CMD", "/usr/bin/loki", "-health" ]
39      start_period: 30s
40      interval: 10s
41      timeout: 5s
42      retries: 5
43     depends_on:
44      - minio
45     networks:                                168
46      <<: *loki-dns
47
48   alloy:
49     image: grafana/alloy:latest
50     volumes:
51      -
52      ↪ ./loki/alloy-local-config.yaml:/etc/alloy/config.alloy:ro
53      - /run/user/1001/podman/podman.sock:/var/run/docker.sock
54     command: run --server http listen-addr=0.0.0.0:12345
```

```
1  # Grafana / Loki
2  upstream grafana {
3      server grafana:3000;
4  }
5  upstream loki-read {
6      server read:3100;
7  }
8  upstream loki-write {
9      server write:3100;
10 }
11
12 server {
13     listen 443 ssl;
14     listen [::]:443 ssl;
15
16     server_name grafana.hiube.ch;
17
18     include /config/nginx/ssl.conf;
19
20     auth_basic "Restricted";
21     auth_basic_user_file /config/nginx/.htpasswd;
22
23     location / {
24         auth_basic off;
25
26         proxy_pass http://grafana;
27         proxy_set_header Host $host;
28         proxy_set_header X-Real-IP $remote_addr;
29         proxy_set_header X-Forwarded-For
30             ↪ $proxy_add_x_forwarded_for;
31         proxy_set_header X-Forwarded-Proto $scheme;
32         proxy_http_version 1.1;
33         proxy_set_header Upgrade $http_upgrade;
34         proxy_set_header Connection $connection_upgrade;
35         proxy_read_timeout 900s;
36         proxy_send_timeout 900s;
37     }
38
39     location = /api/prom/push {
40         proxy_pass http://loki-write$request_uri;
41     }
42
43     location = /api/prom/tail {
44         proxy_pass http://loki-read$request_uri;
45         proxy_set_header Upgrade $http_upgrade;
46         proxy_set_header Connection $connection_upgrade;
47     }
48
49     location ~ /api/prom/. * {
50         proxy_pass http://loki-read$request_uri;
51     }
52
53     location = /loki/api/v1/push {
54         proxy_pass http://loki-write$request_uri;
```

```
1 # ~/loki/loki-config.yaml
2 auth_enabled: false
3
4 server:
5     http_listen_address: 0.0.0.0
6     http_listen_port: 3100
7
8 memberlist:
9     join_members: ["read", "write", "backend"]
10    dead_node_reclaim_time: 30s
11    gossip_to_dead_nodes_time: 15s
12    left_ingesters_timeout: 30s
13    bind_addr: ['0.0.0.0']
14    bind_port: 7946
15    gossip_interval: 2s
16
17 schema_config:
18     configs:
19         - from: 2023-01-01
20           store: tsdb
21           object_store: s3
22           schema: v13
23           index:
24               prefix: index_
25               period: 24h
26
27 common:
28     path_prefix: /loki
29     replication_factor: 1
30     compactor_address: http://backend:3100
31     storage:
32         s3:
33             endpoint: minio:9000
34             insecure: true
35             bucketnames: loki-data
36             access_key_id: loki
37             secret_access_key: supersecret
38             s3forcepathstyle: true
39     ring:
40         kvstore:
41             store: memberlist
42
43 ruler:
44     storage:
45         s3:
46             bucketnames: loki-ruler
47
48 compactor:
49     working_directory: /tmp/compactor
50     retention_enabled: true
51     delete_request_store: s3
52
53 limits_config:
54     retention_period: 8760h # 1 year
55     retention_stream:
56         - selector: '{job="container_logs"}'
```

```
1 # ~/loki/loki-config.yaml
2 auth_enabled: false
3
4 server:
5   http_listen_address: 0.0.0.0
6   http_listen_port: 3100
7
8 memberlist:
9   join_members: ["read", "write", "backend"]
10  dead_node_reclaim_time: 30s
11  gossip_to_dead_nodes_time: 15s
12  left_ingesters_timeout: 30s
13  bind_addr: ['0.0.0.0']
14  bind_port: 7946
15  gossip_interval: 2s
16
17 schema_config:
18   configs:
19     - from: 2023-01-01
20       store: tsdb
21       object_store: s3
22       schema: v13
23       index:
24         prefix: index_
25         period: 24h
26 common:
27   path_prefix: /loki
28   replication_factor: 1
29   compactor_address: http://backend:3100
30   storage:
31     s3:
32       endpoint: minio:9000
33       insecure: true
34       bucketnames: loki-data
35       access_key_id: loki
36       secret_access_key: supersecret
37       s3forcepathstyle: true
38   ring:
39     kvstore:
40       store: memberlist
41 ruler:
42   storage:
43     s3:
44       bucketnames: loki-ruler
45
46 compactor:
47   working_directory: /tmp/compactor
48   retention_enabled: true
49   delete_request_store: s3
50
51 limits_config:
52   retention_period: 8760h # 1 year
53   retention_stream:
54     - selector: '{job="container_logs"}'
```

### A.13. Writing Quality

Introduced from Professor David Stuckler the PEER-System: Simple and clear academic writing.

The structure is:

- ▶ P 1 POINT just one
- ▶ E Evidence and Example 1-n
- ▶ E Explanation
- ▶ R Repeat linking Links and leads to next topic

## A.14. Test-Environment-Code

### A.14.1. container-lab

---

```
sysctl:
  net.ipv4.ip_forward: 1 # required to forward packets between interfaces
exec:
  - apk add --no-cache nftables iproute2
  - ip addr add 10.0.1.1/24 dev eth1
  - ip addr add 172.16.1.1/24 dev eth2
  - ip addr add 192.168.10.1/24 dev eth3
  - ip route replace default via 10.0.1.254
  # Default FORWARD policy: drop everything not explicitly allowed
  - nft add table inet filter
  - sh -c "nft add chain inet filter forward '{ type filter hook forward priority 0; policy drop; }'"
  # Allow return traffic for already-established connections (stateful firewall)
  - nft add rule inet filter forward ct state established,related accept
  # WAN → DMZ: only HTTP/HTTPS and DNS to dmz-server
  - sh -c "nft add rule inet filter forward iifname eth1 oifname eth2 ip daddr 172.16.1.10 tcp dport
↪ '{ 80, 443 }' accept"
  - nft add rule inet filter forward iifname eth1 oifname eth2 ip daddr 172.16.1.10 udp dport 53
↪ accept
  - nft add rule inet filter forward iifname eth1 oifname eth2 ip daddr 172.16.1.10 tcp dport 53
↪ accept
  # WAN → LAN: completely blocked (no rules = default DROP applies)
  # DMZ → LAN: completely blocked
  - nft add rule inet filter forward iifname eth2 oifname eth3 drop
  # LAN → DMZ: allowed (internal users browse DMZ services)
  - nft add rule inet filter forward iifname eth3 oifname eth2 accept
  # LAN → WAN: HTTP/HTTPS only, with source NAT (masquerade)
  - sh -c "nft add rule inet filter forward iifname eth3 oifname eth1 tcp dport '{ 80, 443 }' accept"
  # NAT: source masquerade for LAN and DMZ egress to WAN
  - nft add table ip nat
  - sh -c "nft add chain ip nat postrouting '{ type nat hook postrouting priority 100; }'"
  - nft add rule ip nat postrouting ip saddr 192.168.10.0/24 oifname eth1 masquerade
  - nft add rule ip nat postrouting ip saddr 172.16.1.0/24 oifname eth1 masquerade
  # DMZ → WAN: DNS forwarding (dmz-server forwards unknown queries to 8.8.8.8)
  - nft add rule inet filter forward iifname eth2 oifname eth1 udp dport 53 accept
  - nft add rule inet filter forward iifname eth2 oifname eth1 tcp dport 53 accept

#
# DMZ Nodes (172.16.1.0/24)
#

# — DMZ: Combined Web + DNS Server —————
# Single DMZ node running both nginx (web) and BIND9 (DNS).
# Merging saves a device slot while covering two distinct attack surfaces.
#
# Web (nginx):
```

---

Listing 55: Firewall

Listing 55 shows the nf-tables rules set for segmentation of wan, dmz and lan.

```
# 443/tcp HTTPS – self-signed cert: CN=dmz-server.lab.local
# Recon: Server header leaks version; /robots.txt disallows /admin /backup
#
# DNS (BIND9):
# 53/udp+tcp queries and zone transfer
# AXFR intentionally open (misconfiguration) – reveals all internal IPs
# Recon: dig @172.16.1.10 lab.local AXFR → full zone dump
dmz-server:
  kind: linux
  image: alpine:3.19
  cmd: tail -f /dev/null
  cap-add:
    - NET_ADMIN
  binds:
    - scripts/dmz-web-server:/dmz-web-server:ro
  exec:
    - apk add --no-cache nginx bind bind-tools iproute2 openssl python3
    - ip addr add 172.16.1.10/24 dev eth1
    - ip route replace default via 172.16.1.1
    - sh /dmz-web-server/setup.sh

#
# LAN Nodes (192.168.10.0/24)
#
# — LAN: Internal Server (Samba + LDAP + SSH) —
# Combines Domain Controller and File Server into one node.
# Samba provides SMB shares; OpenLDAP provides a lightweight directory.
# No full AD provisioning – covers the same recon/attack surface at less cost.
#
```

---

Listing 56: Web-Server and DNS

Listing 56 shows a web server and DNS server node in the Containerlab YAML configuration. This node is located in the demilitarized zone (DMZ) of the network.



```
# 139/tcp NetBIOS (Samba)
# 389/tcp LDAP - anonymous base read: DC=lab,DC=local (user list exposed)
# 445/tcp SMB - shares: public (anon read), finance (auth), it (auth)
#
# Recon:
# smbclient -L //192.168.10.5 -N          → lists all shares
# smbmap -H 192.168.10.5                 → share names + permissions
# ldapsearch -x -H ldap://192.168.10.5   → user list, group memberships
ctrl-server:
  kind: linux
  image: alpine:3.19
  cmd: tail -f /dev/null
  cap-add:
    - NET_ADMIN
    - SYS_ADMIN # needed for Samba AD provisioning
  binds:
    - scripts/ctrl-file-server:/ctrl-file-server:ro
  exec:
    - apk add --no-cache samba openldap openldap-back-mdb openldap-clients openssh iproute2
    - ip addr add 192.168.10.5/24 dev eth1
    - ip route replace default via 192.168.10.1
    - sh -c 'echo "nameserver 172.16.1.10" > /etc/resolv.conf'
    - sh /ctrl-file-server/server_setup.sh

# — LAN: Workstation alice (Finance dept, user: asmith) —————
# Simulates a Finance department workstation.
# Sensitive documents on disk; drive mappings reveal share paths.
#
# 22/tcp SSH (asmith / Password123!)
# 445/tcp SMB (Samba standalone)
#
# Recon:
# SSH banner leaks org name and policy
```

---

Listing 57: File Server and Domain Control

The LAN segment consists of four nodes, as shown in the following code snippets.

---

```

# /home/asmith/Desktop/drive_mappings.txt - reveals \\ctrl-server\finance path
alice:
  kind: linux
  image: alpine:3.19
  cmd: tail -f /dev/null
  cap-add:
    - NET_ADMIN
  exec:
    - apk add --no-cache openssh samba iproute2
    - ip addr add 192.168.10.100/24 dev eth1
    - ip route replace default via 192.168.10.1
    - sh -c 'echo "nameserver 172.16.1.10" > /etc/resolv.conf'
    - ssh-keygen -A
    - sh -c 'echo "PasswordAuthentication yes" >> /etc/ssh/sshd_config'
    - sh -c 'printf "NSAK-Enterprise - Authorized Access Only\nThis system is monitored.\n" >
      ↪ /etc/ssh/banner.txt'
    - sh -c 'echo "Banner /etc/ssh/banner.txt" >> /etc/ssh/sshd_config'
    - adduser -D -s /bin/sh asmith
    - sh -c 'echo "asmith:Password123!" | chpasswd'
    - mkdir -p /home/asmith/Documents /home/asmith/Desktop
    - sh -c 'echo "Q3 2024 Finance Report - CONFIDENTIAL" > /home/asmith/Documents/Q3_Finance.txt'
    - sh -c 'printf "DOMAIN=LAB\nUSER=asmith\nSHARE=\\\\\\\\\\\\\\\\ctrl-server\\\\\\\\finance\n" >
      ↪ /home/asmith/Desktop/drive_mappings.txt'
    - /usr/sbin/sshd

# — LAN: Workstation bob (IT dept, user: bjones) —————
# Simulates an IT department workstation.
# bjones has elevated rights - scripts directory reveals full internal topology.
#
# 22/tcp  SSH  (bjones / Password123!)
# 445/tcp  SMB  (Samba standalone)
#
# Recon:
# /home/bjones/Scripts/hosts.txt - lists all internal IPs and roles
# /home/bjones/Scripts/README.txt - hints at admin access
bob:
  kind: linux
  image: alpine:3.19
  cmd: tail -f /dev/null
  cap-add:
    - NET_ADMIN
  exec:
    - apk add --no-cache openssh samba iproute2
    - ip addr add 192.168.10.101/24 dev eth1
    - ip route replace default via 192.168.10.1
    - sh -c 'echo "nameserver 172.16.1.10" > /etc/resolv.conf'
    - ssh-keygen -A
    - sh -c 'echo "PasswordAuthentication yes" >> /etc/ssh/sshd_config'
    - sh -c 'printf "Acme Corp AG - Authorized Access Only\nThis system is monitored.\n" >
      ↪ /etc/ssh/banner.txt'
    - sh -c 'echo "Banner /etc/ssh/banner.txt" >> /etc/ssh/sshd_config'
    - adduser -D -s /bin/sh bjones
    - sh -c 'echo "bjones:Password123!" | chpasswd'
    - mkdir -p /home/bjones/Documents /home/bjones/Scripts
    - sh -c 'printf "ctrl-server=192.168.10.5\nalice=192.168.10.100\nprinter=192.168.10.50\n" >
      ↪ /home/bjones/Scripts/hosts.txt'
    - sh -c 'echo "# admin tooling - see ctrl-server at 192.168.10.5" >
      ↪ /home/bjones/Scripts/README.txt'
    - /usr/sbin/sshd

# — LAN: Network Printer (HP-like, intentionally misconfigured) —————
# Simulates a typical enterprise network printer - the "forgotten device".
# Rarely patched, often misconfigured, rich source of OSINT and lateral movement.
#
# 80/tcp  HTTP  - unauthenticated web admin interface
# 161/udp  SNMP  - community string "public" (v1/v2c, full device info)
# 631/tcp  IPP   - Internet Printing Protocol (job history, no auth)
#
# Recon:

```

---

Listing 58: Work-Stations Bob and Alice

```
# curl http://192.168.10.50          → unauthenticated admin page, firmware version
# curl http://192.168.10.50:631/jobs → print job history reveals internal usernames
printer:
  kind: linux
  image: alpine:3.19
  cmd: tail -f /dev/null
  cap-add:
    - NET_ADMIN
    - NET_BIND_SERVICE # needed to bind to privileged ports 80 and 631
  binds:
    # printer_sim.py serves the web admin UI (port 80) and IPP (port 631)
    - scripts/printer_sim.py:/printer_sim.py:ro
  exec:
    - apk add --no-cache python3 net-snmp iproute2
    - ip addr add 192.168.10.50/24 dev eth1
    - ip route replace default via 192.168.10.1
    - sh -c 'echo "nameserver 172.16.1.10" > /etc/resolv.conf'
    # SNMP with open public community string – intentional misconfiguration
    - sh -c 'printf "rocommunity public\nsysName HP-LaserJet-M428fdw\nsysLocation Server Room
    ↪ B2\nsysContact it@lab.local\n" > /etc/snmp/snmpd.conf'
    - sh -c 'snmpd -f -Lo &'
    # Web admin UI (80) and IPP job listener (631)
    - sh -c 'python3 /printer_sim.py > /var/log/printer_sim.log 2>&1 &'

# — NSAK – Swiss Army Knife (Reconnaissance Node) —————
# The nsak tool itself runs here. Dual-homed to observe both LAN and DMZ.
# eth1 → LAN 192.168.10.200 (full internal network visibility)
# eth2 → DMZ 172.16.1.200 (DMZ visibility – same segment as dmz-server)
#
# Quickstart:
# docker exec -it clab-virtual-enterprise-network-lab-nsak bash
# nsak drill execute discover_hosts --interface eth1
# nsak drill execute discover_hosts --interface eth2
```

---

Listing 59: Printer

Listing 58 illustrates two workstations, Alice and Bob, the printer shown in Listing 59, and the domain controller with Samba file server shown in Listing 57.

```
# LAN segment (192.168.10.0/24)
- endpoints: [ "firewall:eth3", "lan-br:fw3" ]
- endpoints: [ "ctrl-server:eth1", "lan-br:ctrl" ]
- endpoints: [ "alice:eth1", "lan-br:A" ]
- endpoints: [ "bob:eth1", "lan-br:B" ]
- endpoints: [ "printer:eth1", "lan-br:print" ]
- endpoints: [ "nsak:eth1", "lan-br:nsak-lan" ]
```

---

Listing 60: Links

Listing 60 shows the connection between endpoints over the defined interfaces.



### A.14.2. Flaws and Vulnerabilities

---

```
1  DMZ --- dmz-server (172.16.1.10)
2
3  DNS / BIND9
4  - allow-transfer { any; } --- AXFR (zone transfer) open to everyone
5  - dnssec-validation no --- no DNSSEC validation
6  - Zone contains internal LAN addresses (ctrl, alice, bob, printer)
7
8  Web / nginx
9  - /secret.txt accessible without auth -> returns DB Password:
   ↪ {flag:exposed-SOS}
10 - Self-signed certificate (CN=dmz-server.vlab.local)
11 ---
12 LAN --- ctrl-server (192.168.10.5)
13
14 LDAP / OpenLDAP
15 - by anonymous read --- anonymous read access to all entries -> user list,
   ↪ emails, departments
16 - rootpw Password123! stored in plaintext in config
17 - User passwords stored in LDIF: userPassword: Password123!
18 ldapsearch -x -H ldap://192.168.10.5 -b "dc=lab,dc=local"
19
20 SMB / Samba
21 - ntlm auth = yes + lanman auth = yes --- outdated, weak auth protocols
22 - Share public anonymously readable (no auth) -> README.txt hints at bjones
23 - Share finance contains salary and budget data (asmith / Password123!)
24 - Share it contains dc01_build_sheet.txt with full infrastructure details
25 - All passwords: Password123!
26 ---
27 LAN --- alice (192.168.10.100) \& bob (192.168.10.101)
28
29 - SSH password auth enabled, weak password Password123!
30 - SSH banner leaks organisation name
31 - alice: /home/asmith/Documents/Q3_Finance.txt (confidential),
   ↪ /home/asmith/Desktop/drive_mappings.txt (share paths)
32 - bob: /home/bjones/Scripts/hosts.txt (all internal IPs + roles), README.txt
   ↪ hints at admin access
33 ---
34 LAN --- printer (192.168.10.50)
35
36 - Port 80: unauthenticated web admin interface, leaks firmware version
   ↪ (HP-WebServer/2.6.5), serial number, location, contact
37 - Port 631 /jobs: print job history without auth -> reveals usernames
   ↪ (asmith, bjones, admin) and document names
   (server_credentials.txt, payroll)
38 - SNMP community string public (v1/v2c) -> snmpwalk exposes full device info
39 ---
40 Firewall policy (intentional gaps)
41
42 - LAN -> DMZ fully open (no port restriction) --- LAN users can reach all
   ↪ DMZ ports, not just 80/443
43 - DMZ -> WAN: DNS forwarding to 8.8.8.8 allowed -> DNS exfiltration possible
44
```

---

Listing 61: Introduced Vulnerabilities in the Virtual Enterprise Network



## A.15. Reconnaissance Scenario: Source Listings

```
1  """
2  Scenario entrypoint for Reconnaissance .
3  """
4  import logging
5
6  from scapy.arch import get_if_addr
7
8  from nsak.core import DrillManager
9  from nsak.core.network import NetworkDiscoveryResultMap
10 from nsak.core.network.enumerate_services_result import
    ↳ EnumerateServicesResult
11 from nsak.core.scenario_results.reconnaissance_scenario_result
    ↳ import ReconnaissanceScenarioResult
12
13 logger = logging.getLogger(__name__)
14
15
16 def run(interface: str | None = None, subnet: str | None = None)
    ↳ -> ReconnaissanceScenarioResult:
17     """
18     Scenario, which runs Reconnaissance attack.
19     1. If no interface is specified, scan all active physical
        ↳ interfaces
20     2. If no interface is specified, scan all active physical
        ↳ interfaces
21     3. Discover network on ifc and subnet
22     :return: None
23     """
24     if interface is None:
25         discovered =
            ↳ DrillManager.execute("discover_network_interfaces")
26         interfaces_to_scan = [iface.name for iface in
            ↳ discovered]
27     else:
28         interfaces_to_scan = [interface]
29
30     all_results = {}
31     for iface_name in interfaces_to_scan:
32         if get_if_addr(iface_name) in ("0.0.0.0", ""):
33             DrillManager.execute("dhcp_request",
                ↳ interface=iface_name)
34
35         network_discovery_result_map: NetworkDiscoveryResultMap
            ↳ = DrillManager.execute(
36             "discover_hosts",
37             interface=iface_name,
38             subnet=subnet,
39         )
40         all_results.update(network_discovery_result_map.results)
41
42     network discovery result map =
```

```
1 drills:
2   discover_hosts:
3     name: Discover Hosts
4     author: vonal3@bfh.ch
5     repository:
6       ↪ https://gitlab.ti.bfh.ch/gausf1-vonal3/swiss-army-knife-network-sniffe
7
8     dependencies:
9       system:
10        - arp-scan
11        - nmap
12       python: [ ]
13
14     interface:
15       arguments:
16         interface:
17           type: str
18         subnet:
19           type: str
20           default: null
21       return_type: NetworkDiscoveryResultMap
```

---

Listing 63: Discover Hosts Drill: Resource Configuration



```
1 import logging
2 import dataclasses
3 import subprocess
4 from ipaddress import IPv4Address, IPv6Address
5
6 from nsak.core import IPAddress
7 from nsak.core import config
8 from nsak.core.network import NetworkDiscoveryResultMap,
9     ↳ NetworkDiscoveryResult, NetworkService,
10     ↳ NetworkServiceEndpoint
11 from nsak.core.network.configuration import
12     ↳ EthernetConfiguration
13
14 logger = logging.getLogger(__name__)
15
16 @dataclasses.dataclass(frozen=True, kw_only=True)
17 class ARPScanResult:
18     ip: IPAddress
19     mac: str
20     vendor: str
21
22 def parse_arp_scan_result(arp_scan_result_output: str) ->
23     ↳ list[ARPScanResult]:
24     """
25     :param arp_scan_result_output:
26     :return:
27     """
28
29     entries = []
30     for line in arp_scan_result_output.splitlines():
31         parts = line.split(None, 2)
32         if len(parts) < 2:
33             continue
34
35         raw_ip, mac = parts[:2]
36         try:
37             ip = IPv4Address(raw_ip)
38         except ValueError:
39             ip = IPv6Address(raw_ip)
40         vendor = parts[2] if len(parts) == 3 else ""
41         entries.append(ARPScanResult(ip=ip, mac=mac,
42     ↳ vendor=vendor)) 183
43
44     return entries
45
46 def parse_nmap_result(nmap_output: str) -> list[ARPScanResult]:
47     """
48     Parses the output of `nmap -sn` into a list of
49     ↳ ARPScanResults.
```

---

```

1 import logging
2 import re
3 import subprocess
4 from typing import Literal
5
6 from nsak.core.network import NetworkDiscoveryResultMap,
  ↳ NetworkService, NetworkServiceEndpoint
7
8 logger = logging.getLogger(__name__)
9
10 # matches: "22/tcp open  ssh      OpenSSH 8.9p1"
11 _PORT_LINE =
  ↳ re.compile(r"^(\d+)/(\tcp|udp)\s+open\s+(\S+)\s*(.*)")
  ↳ #compile regex once in obj (p,prot,name,version,
12
13
14 def parse_nmap(stdout: str) -> list[tuple[int, Literal["tcp",
  ↳ "udp"], str, str]]:
15     """Extract open ports from nmap stdout as (port, protocol,
  ↳ name, version) tuples.
16
17     :param stdout: Raw nmap output text.
18     :return: List of (port, protocol, service name, version)
  ↳ tuples for open ports.
19     """
20     results = []
21     # compare the output with the regex and skip if output does
  ↳ not match | group(0)=whole line, group(1)=port...
22     for line in stdout.splitlines():
23         match = _PORT_LINE.match(line.strip())
24         if match:
25             port = int(match.group(1))
26             protocol = match.group(2)
27             name = match.group(3)
28             version = match.group(4).strip()
29             results.append((port, protocol, name, version))
30     return results
31
32
33 def run(discovery_result: NetworkDiscoveryResultMap) ->
  ↳ NetworkDiscoveryResultMap:
34     """Run nmap service scan on all target IPs and append open
  ↳ ports to the discovery result.
35
36     :param discovery_result: Existing host discovery result
  ↳ containing target IPs per interface.
37     :return: The same result map, with open port/service
  ↳ information.
38     """
39     for iface_name, result in discovery_result.results.items():
40         mac_by_ip = {
41             endpoint.ip: endpoint.mac
42             for service in result.network_services
43             for endpoint in service.endpoints

```

---

```

1  import logging
2  import re
3  import subprocess
4
5  from nsak.core.network import NetworkDiscoveryResultMap
6  from nsak.core.network.enumerate_services_result import
   ↳ EnumerateServicesResult, EnumerateServicesResultEntry
7
8  logger = logging.getLogger(__name__)
9
10 _NSE_LINE = re.compile(r"^\\[_ ]?\\s*(.*)") # nmap NSE output:
   ↳ "| text" or "|_ text"
11
12 # Known services get targeted scripts; everything else falls
   ↳ back to "banner"
13 # (banner grabs the initial TCP response – gives software name +
   ↳ version for any unknown service)
14 _SERVICE_SCRIPTS: dict[str, list[str]] = {
15     "http": ["http-title", "http-headers",
   ↳ "http-robots.txt"],
16     "https": ["http-title", "http-headers",
   ↳ "http-robots.txt"],
17     "http-alt": ["http-title", "http-headers",
   ↳ "http-robots.txt"],
18     "domain": ["dns-zone-transfer", "dns-brute"],
19     "smtp": ["smtp-commands", "smtp-enum-users"],
20     "smtps": ["smtp-commands"],
21     "ftp": ["ftp-anon", "ftp-ls"],
22     "netbios-ssn": ["smb-security-mode", "smb2-security-mode"],
23     "microsoft-ds": ["smb-security-mode", "smb2-security-mode"],
24     "ldap": ["ldap-rootdse"],
25     "ldaps": ["ldap-rootdse"],
26 }
27
28
29 def _parse_nse_output(stdout: str) -> list[str]:
30     """
31     Extract script result lines from nmap stdout.
32
33     :param stdout: Raw nmap output.
34     :return: Parsed output lines with pipe prefixes stripped.
35     """
36     findings = []
37     for line in stdout.splitlines():
38         m = _NSE_LINE.match(line.strip())
39         if m:
40             text = m.group(1).strip()
41             if text:
42                 findings.append(text)
43     return findings
44
45
46 def run(discovery_result: NetworkDiscoveryResultMap) ->
   EnumerateServicesResult:

```